

Shift Happens: Building Robust AI Models with Domain Adaptation

Aleksandra Ciprijanovic

Associate Scientist, Fermilab
aleksand@fnal.gov



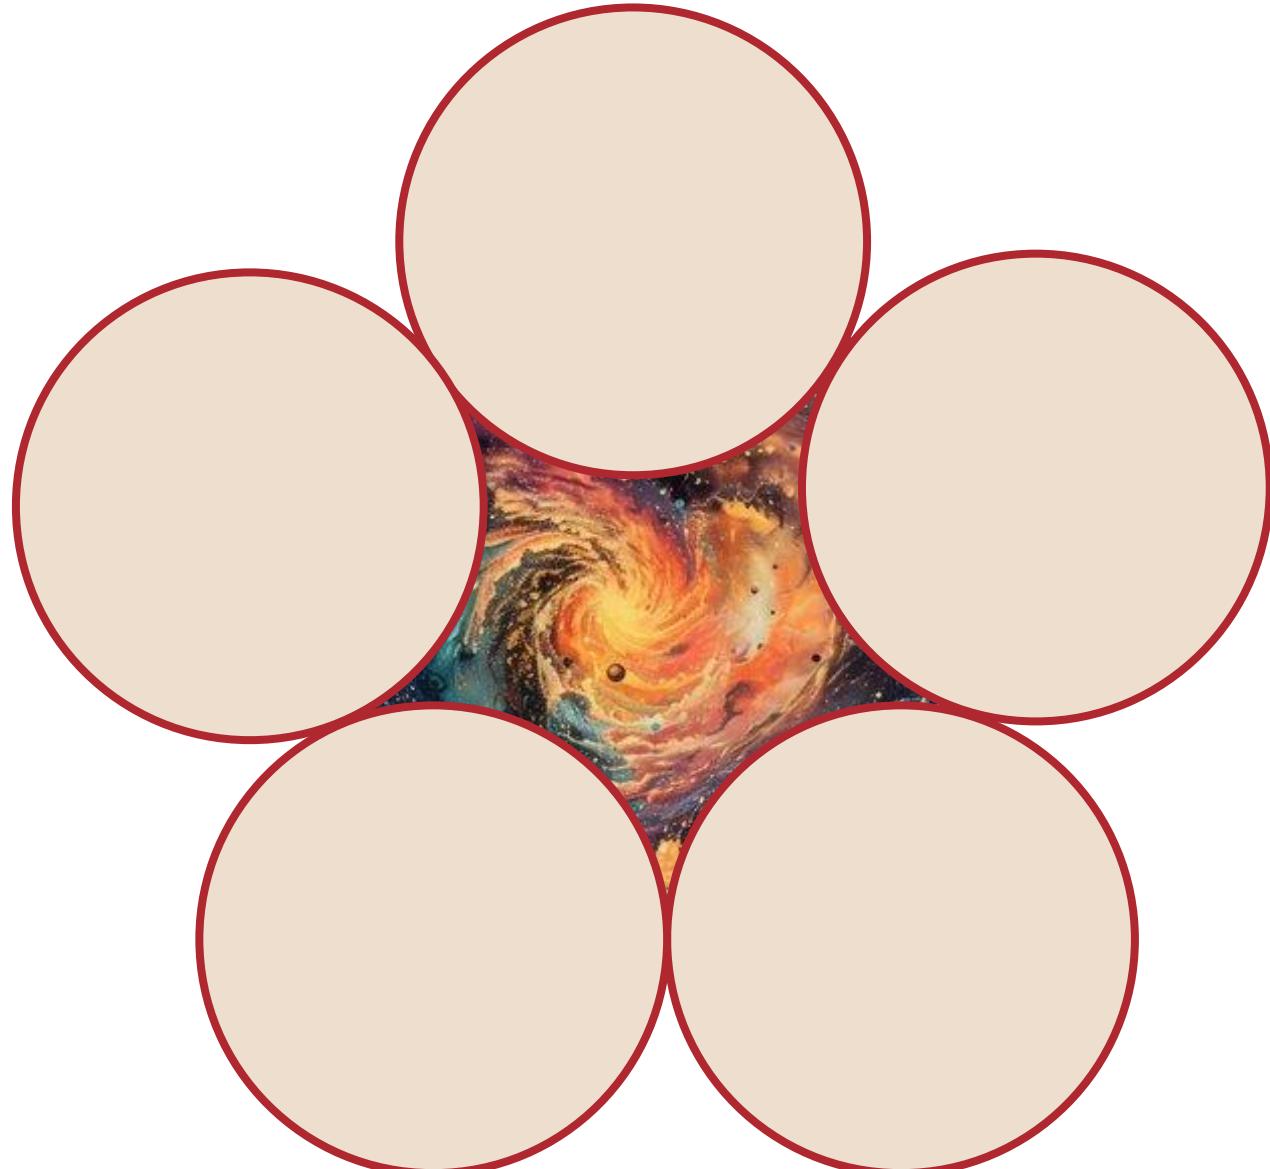
**U.S. DEPARTMENT
of ENERGY**

Fermi National Accelerator Laboratory is managed by
FermiForward for the U.S. Department of Energy Office of Science



Major challenges for AI applications in Science

As listed in The Dawes Review 10: The impact of
deep learning for the analysis of galaxy surveys
[Huertas-Company & Lanusse 2022](#)



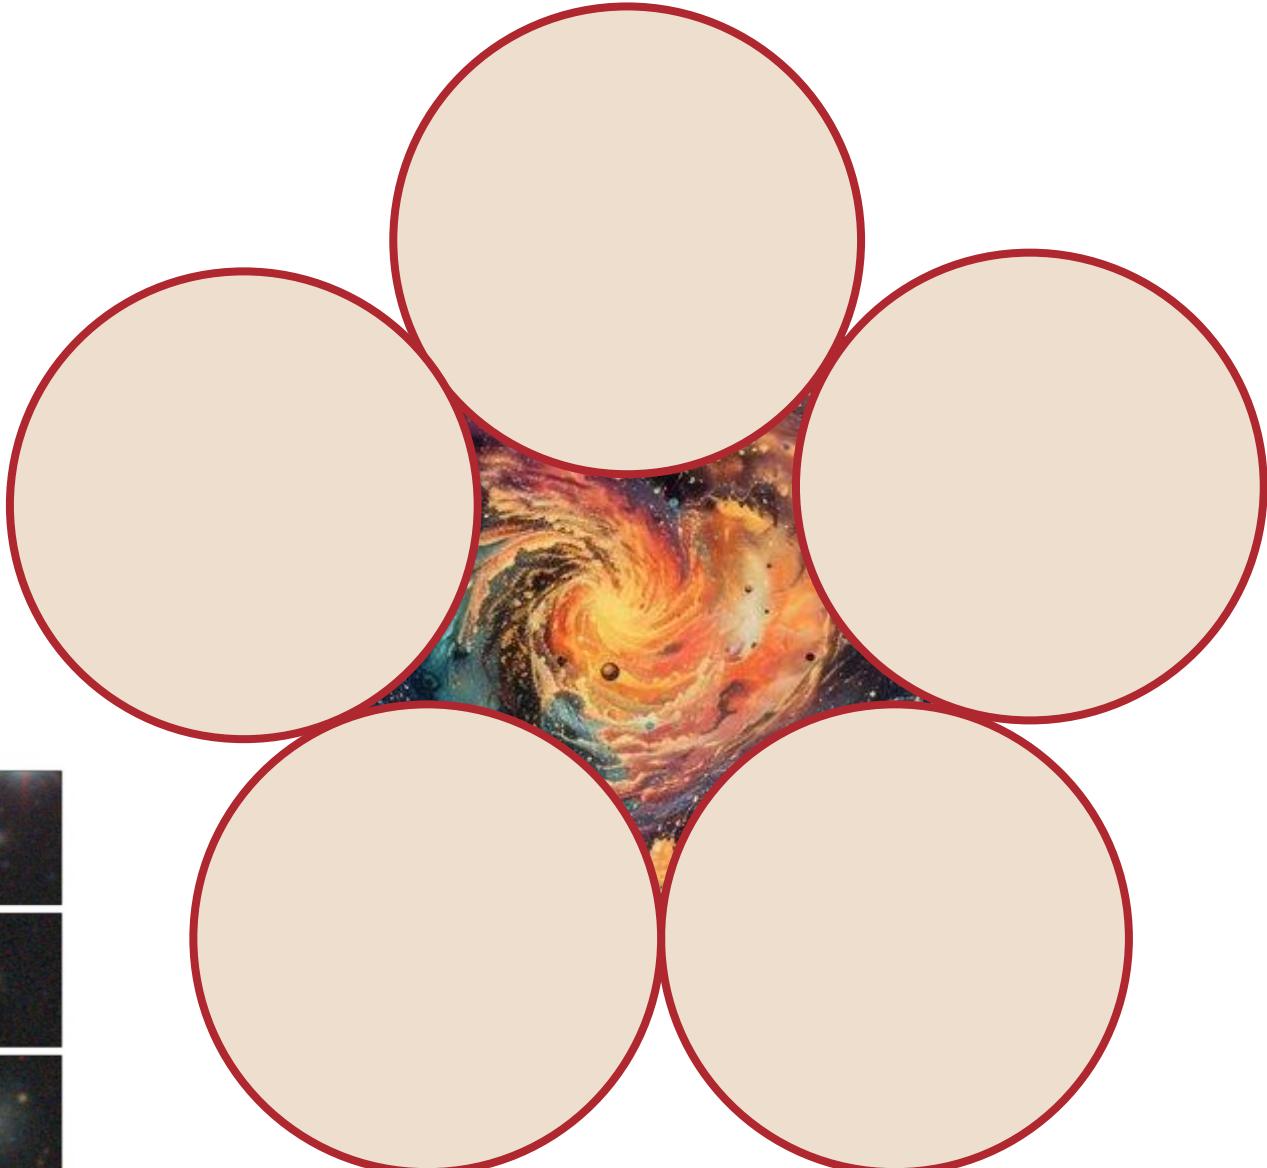
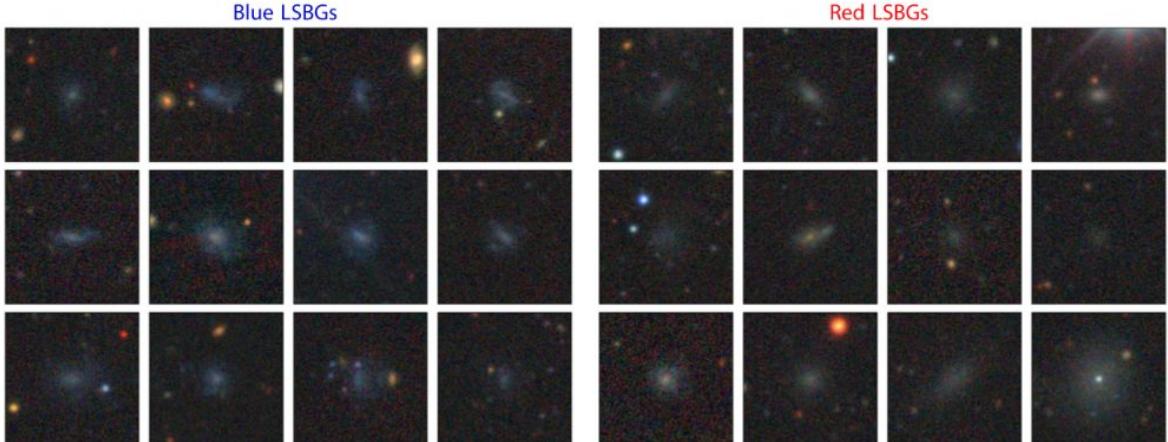


Major challenges for AI applications in Science

As listed in The Dawes Review 10: The impact of
deep learning for the analysis of galaxy surveys
[Huertas-Company & Lanusse 2022](#)

Low Surface Brightness Galaxies (LSBGs)

Tanoglidis et al. 2021.





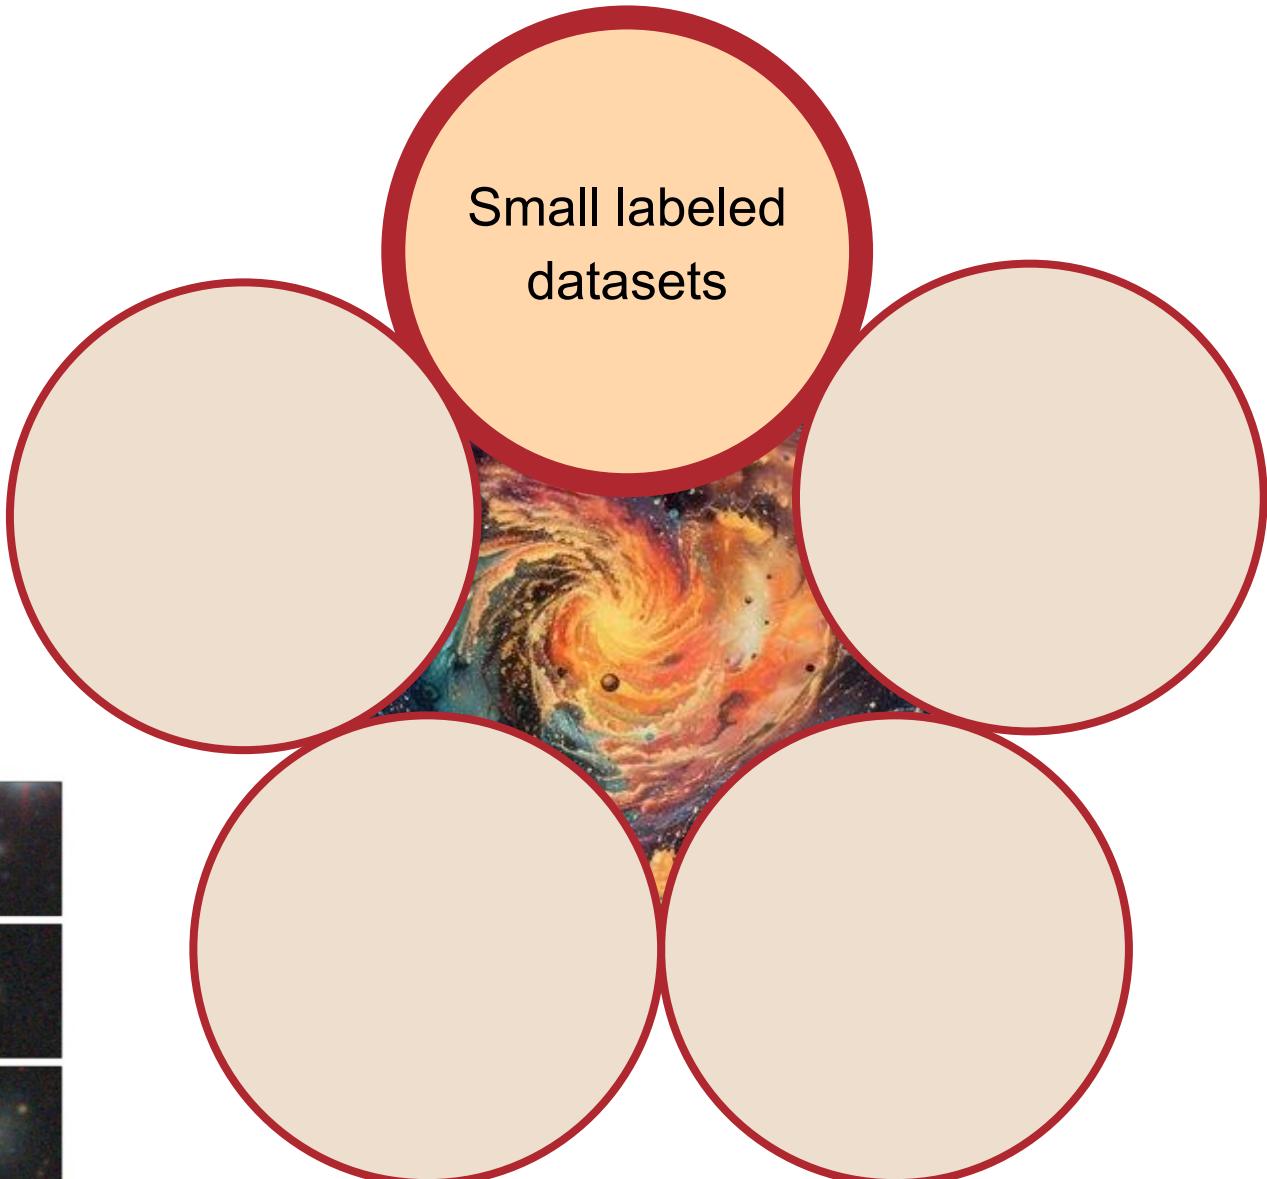
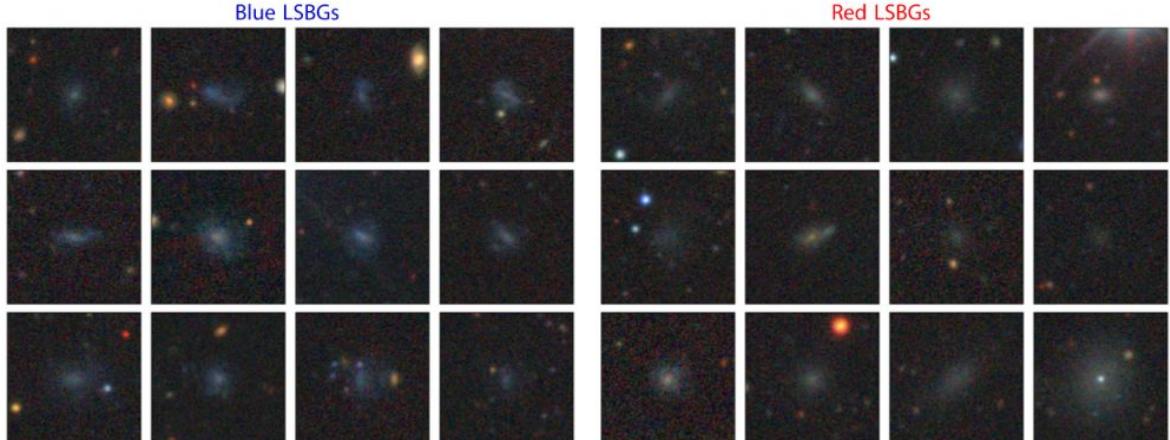
Major challenges for AI applications in Science

As listed in The Dawes Review 10: The impact of
deep learning for the analysis of galaxy surveys
[Huertas-Company & Lanusse 2022](#)

I want to train my AI model by I don't
have enough trustworthy labels?

Low Surface Brightness Galaxies (LSBGs)

Tanoglidis et al. 2021.



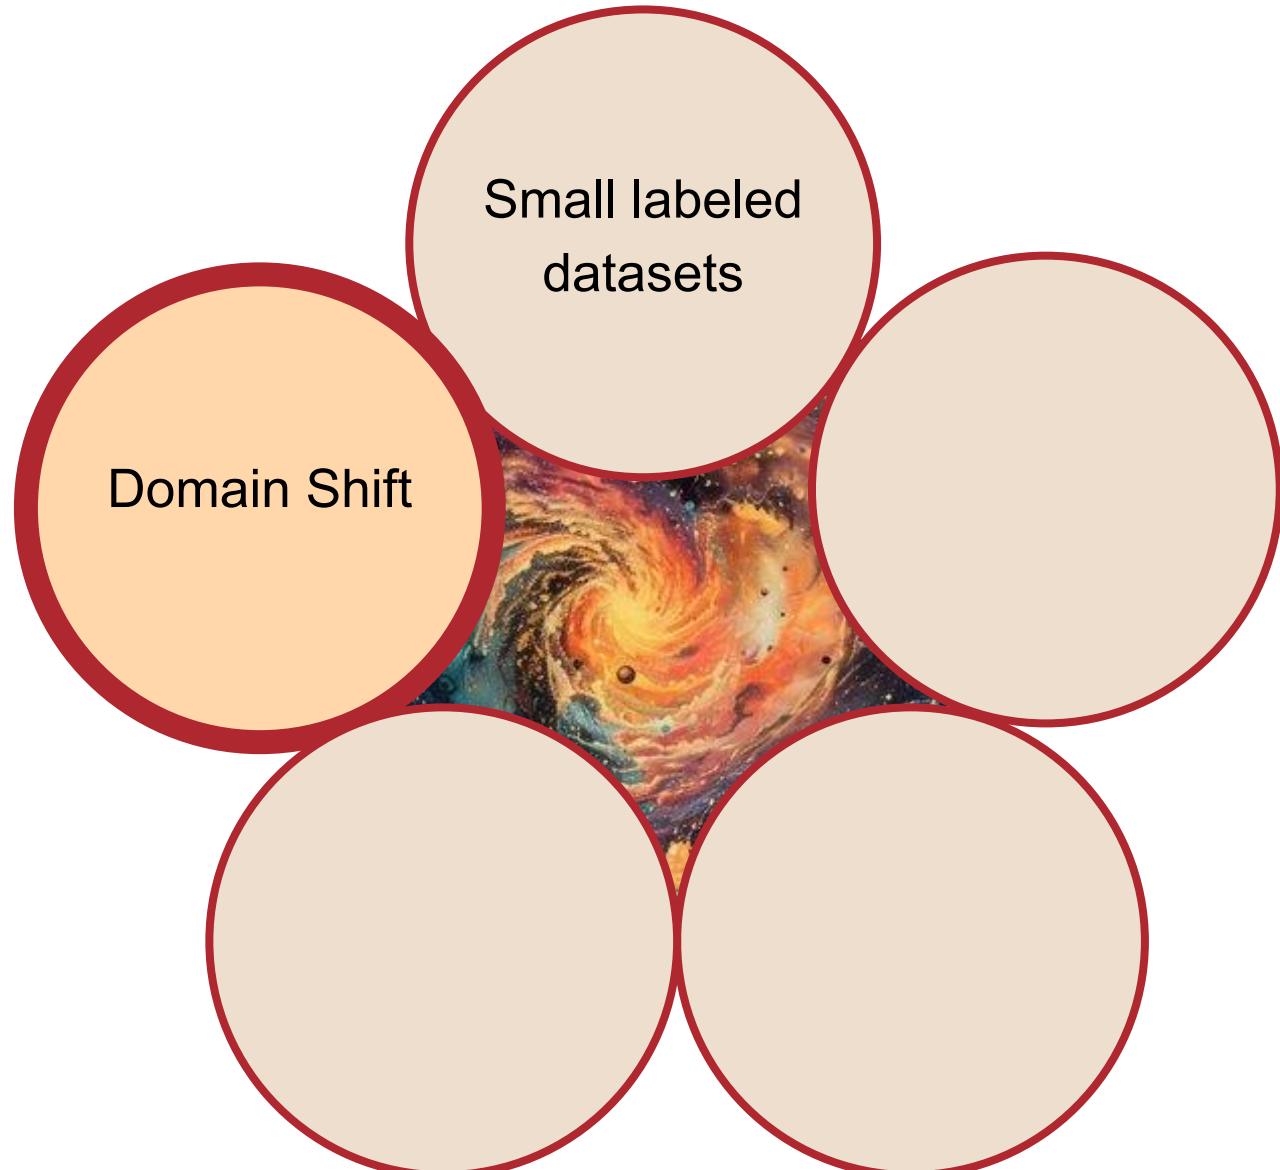


Major challenges for AI applications in Science

As listed in The Dawes Review 10: The impact of
deep learning for the analysis of galaxy surveys
[Huertas-Company & Lanusse 2022](#)

I want to train my AI model by I don't
have enough trustworthy labels?

**Can I train on simulations and apply
the model to my real data?**





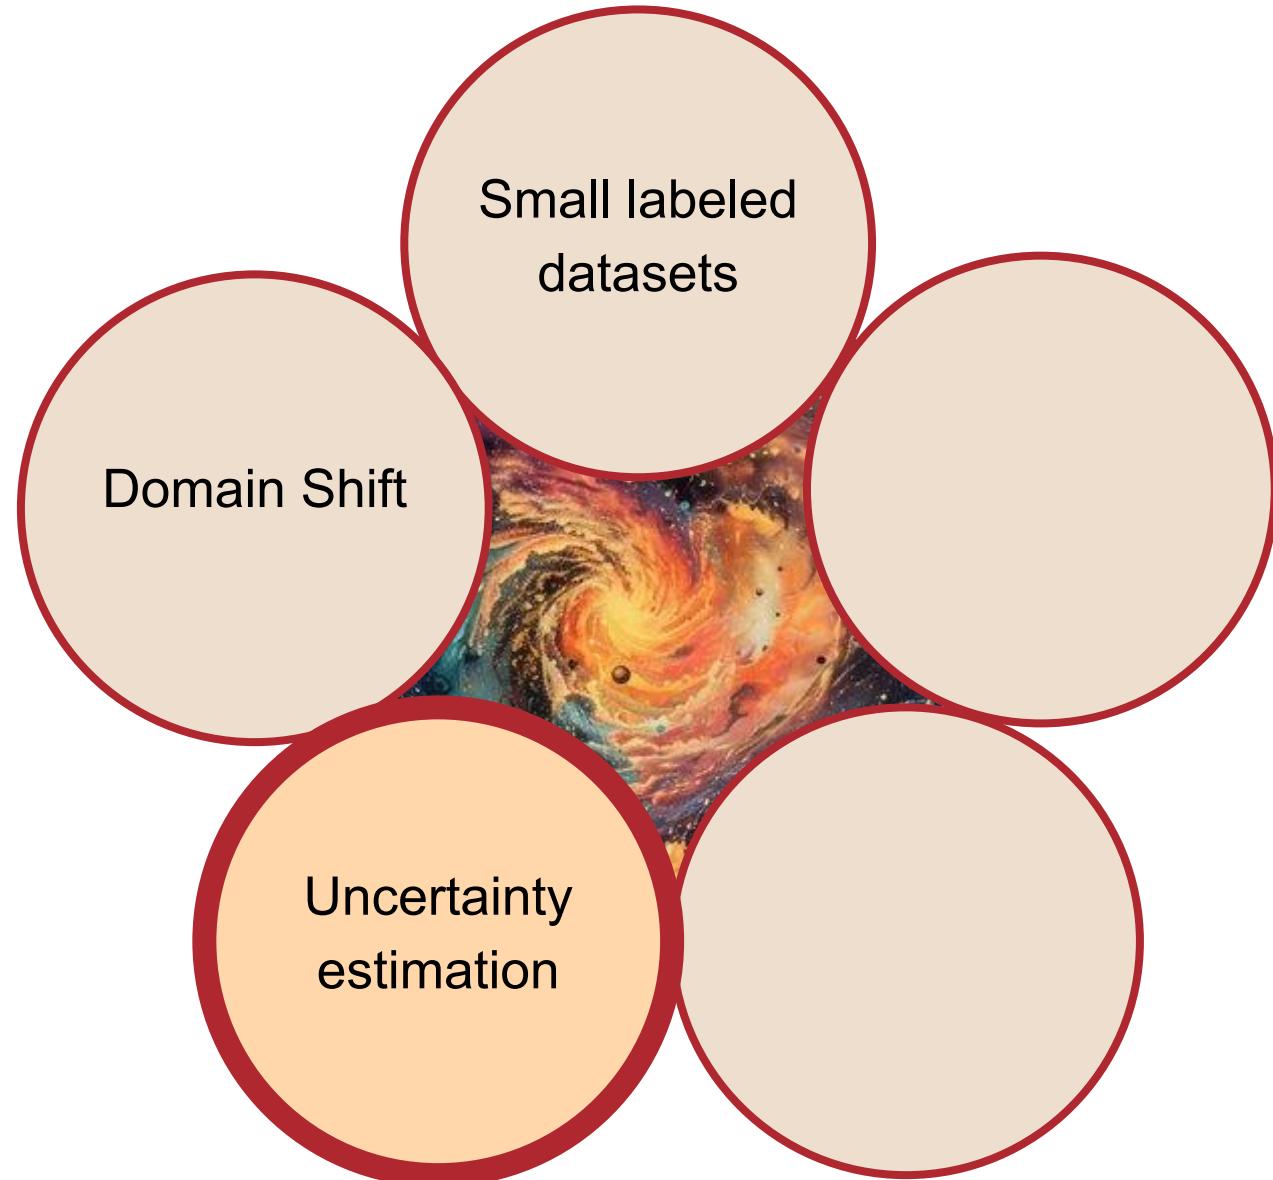
Major challenges for AI applications in Science

As listed in The Dawes Review 10: The impact of
deep learning for the analysis of galaxy surveys
[Huertas-Company & Lanusse 2022](#)

I want to train my AI model by I don't
have enough trustworthy labels?

Can I train on simulations and apply
the model to my real data?

**How well is my model going to
perform?**





Major challenges for AI applications in Science

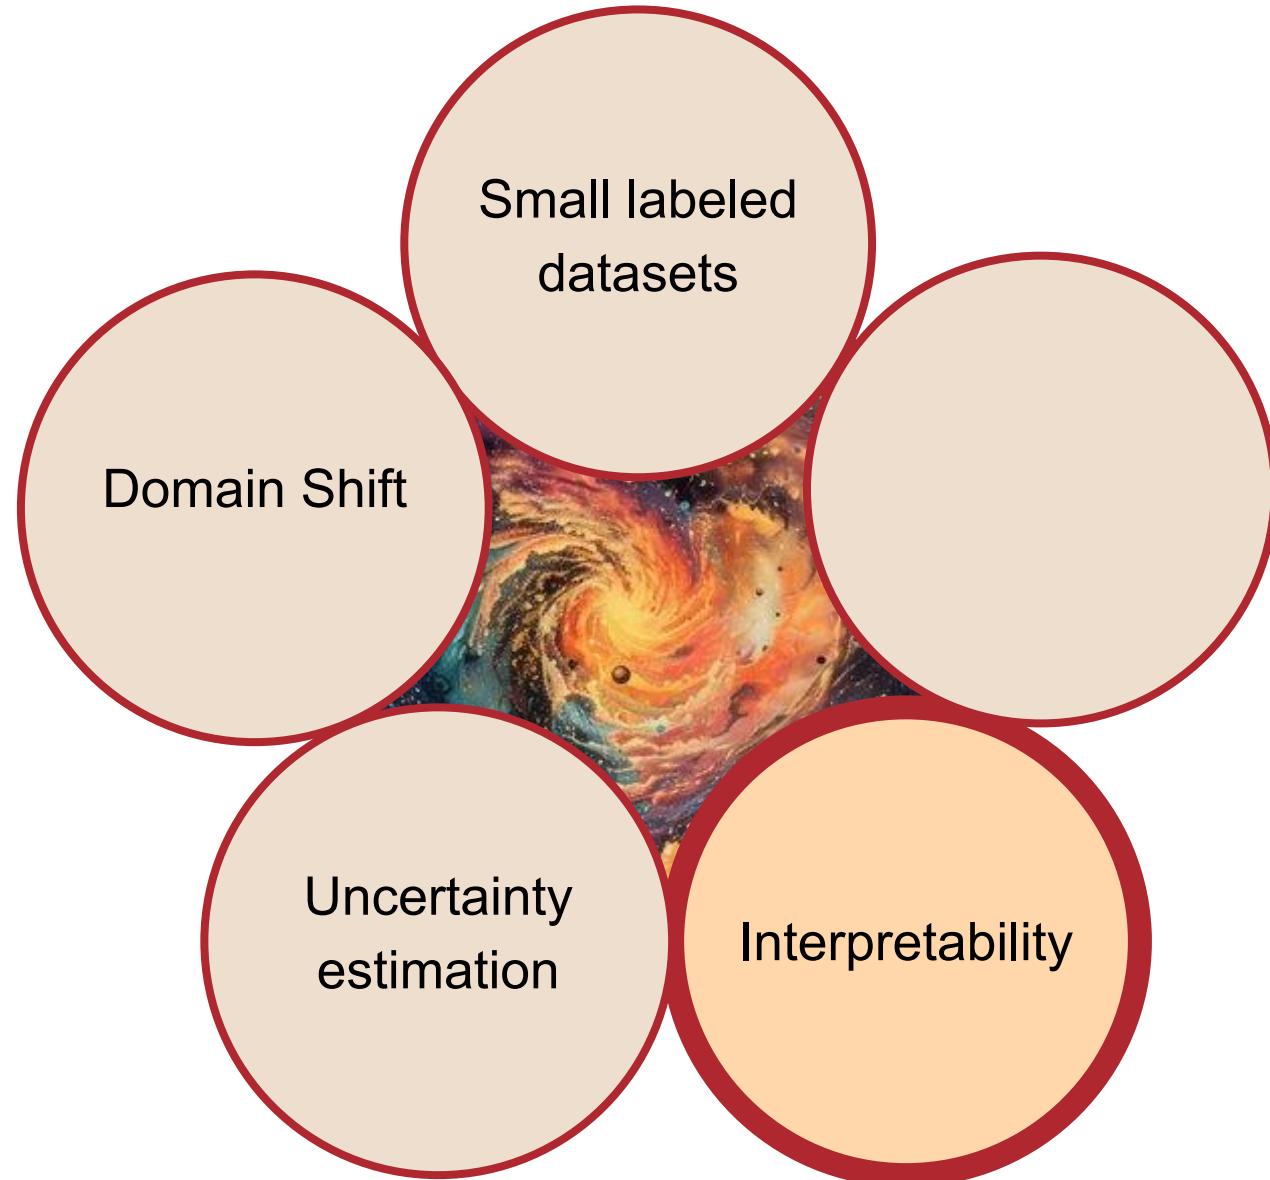
As listed in The Dawes Review 10: The impact of
deep learning for the analysis of galaxy surveys
[Huertas-Company & Lanusse 2022](#)

I want to train my AI model by I don't
have enough trustworthy labels?

Can I train on simulations and apply
the model to my real data?

How well is my model going to
perform?

Can I understand its performance?





Major challenges for AI applications in Science

As listed in The Dawes Review 10: The impact of deep learning for the analysis of galaxy surveys
[Huertas-Company & Lanusse 2022](#)

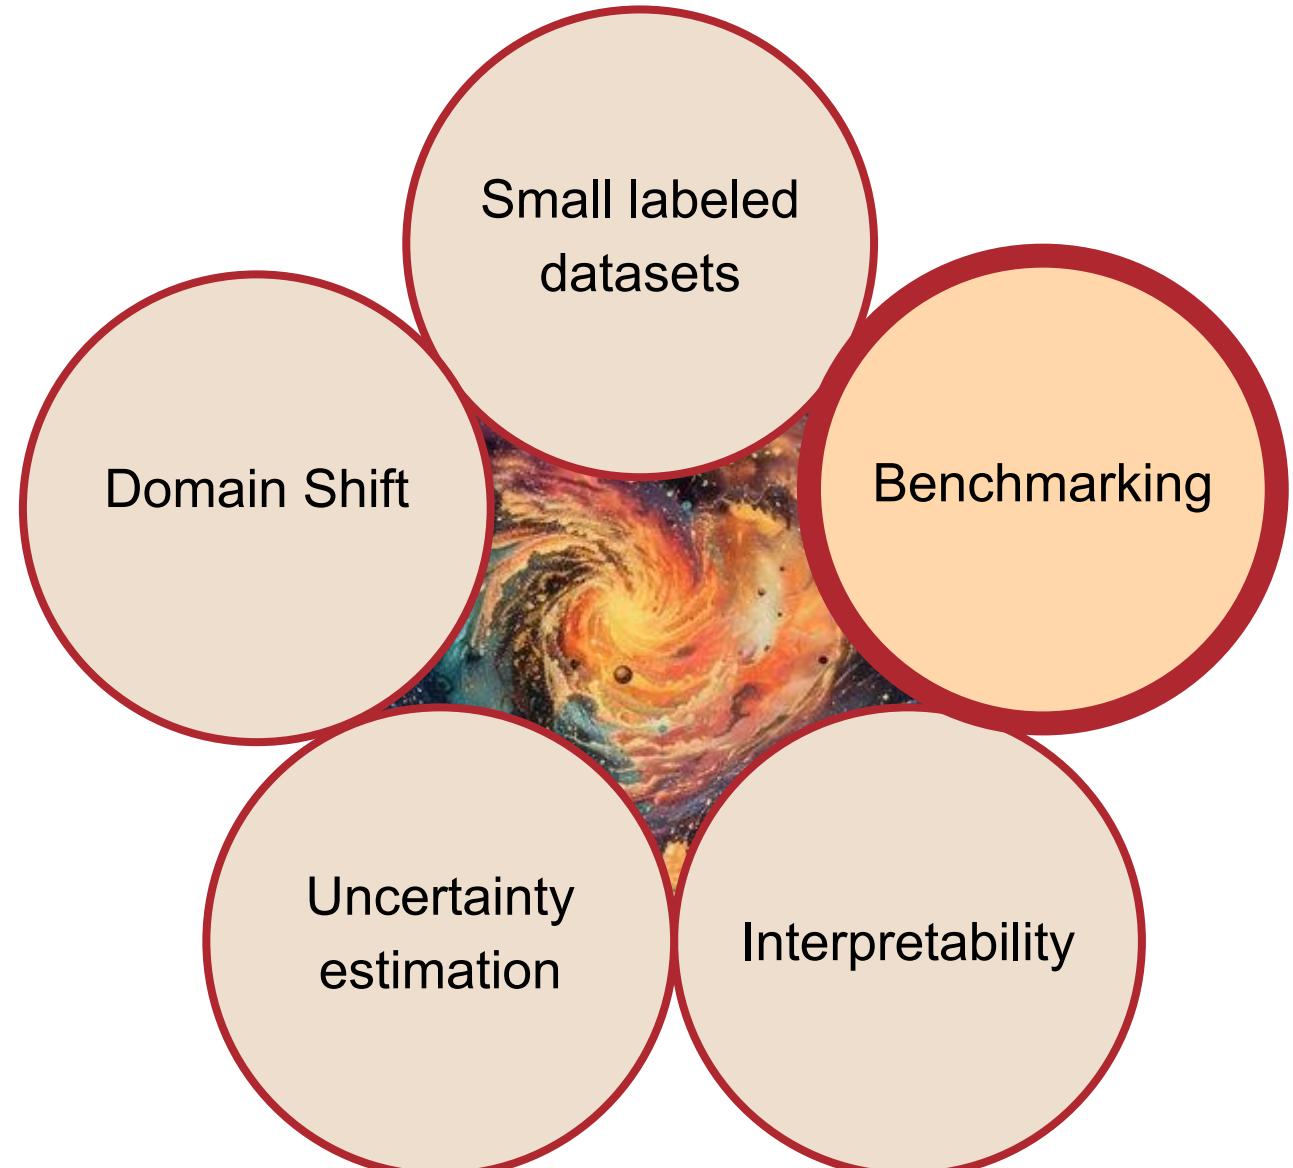
I want to train my AI model by I don't have enough trustworthy labels?

Can I train on simulations and apply the model to my real data?

How well is my model going to perform?

Can I understand its performance?

Can I compare my model with other researchers?





Major challenges for AI applications in Science

As listed in The Dawes Review 10: The impact of deep learning for the analysis of galaxy surveys
[Huertas-Company & Lanusse 2022](#)

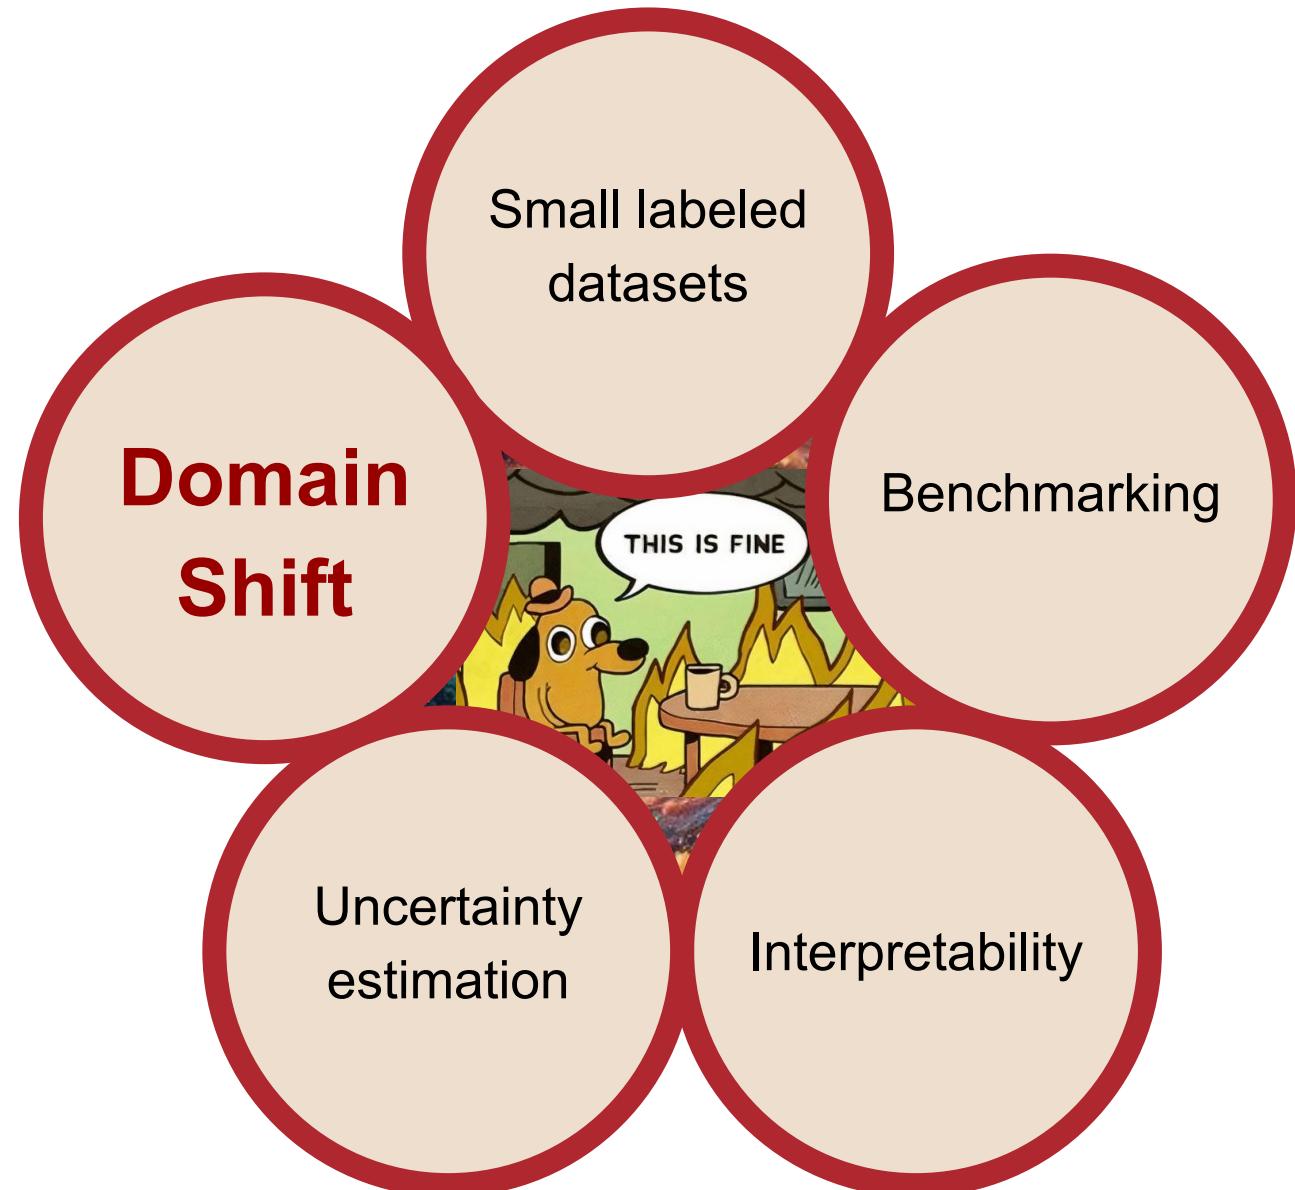
I want to train my AI model by I don't have enough trustworthy labels?

Can I train on simulations and apply the model to my real data?

How well is my model going to perform?

Can I understand its performance?

Can I compare my model with other researchers?





Talk Outline

Part 1: Domain Shift Problem

Part 2: Domain Adaptation Techniques

Part 3: Examples

Part 4: Distance-based Methods

Part 5: Uncertainty Quantification *

Part 6: Open Questions & Research Trends

A perspective view looking down a long aisle between two rows of tall, dark server racks in a data center. The racks are filled with various components and cables. The perspective creates a strong sense of depth.

01

Domain Shift Problem



More Data - More problems

All areas of astro(physics) often need to create
**model trained on simulated data, that also work
on real data!**

DOMAIN SHIFT

Missing and unknown
physics, wrong geometry,
background levels

Computational constraints
for simulations

Detector problems,
transients, errors, data
compression

Imperfect addition of
observational effects

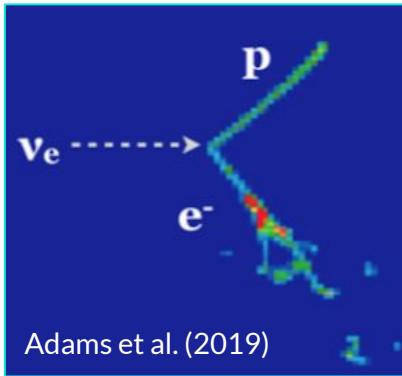
Different detectors or
telescopes

MicroBooNE
(neutrinos)

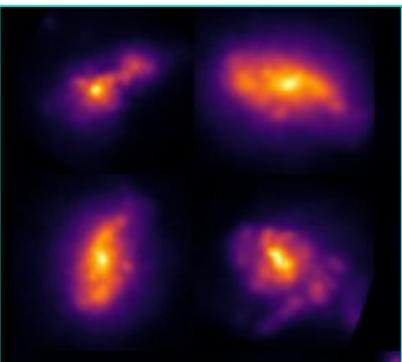
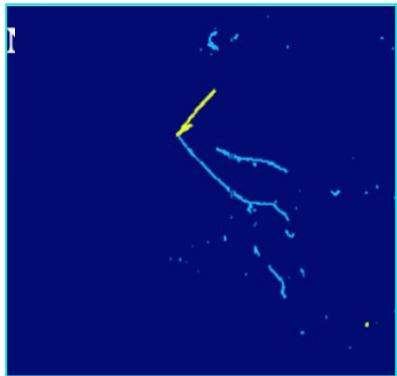
Illustris / Hubble
(merging galaxies)

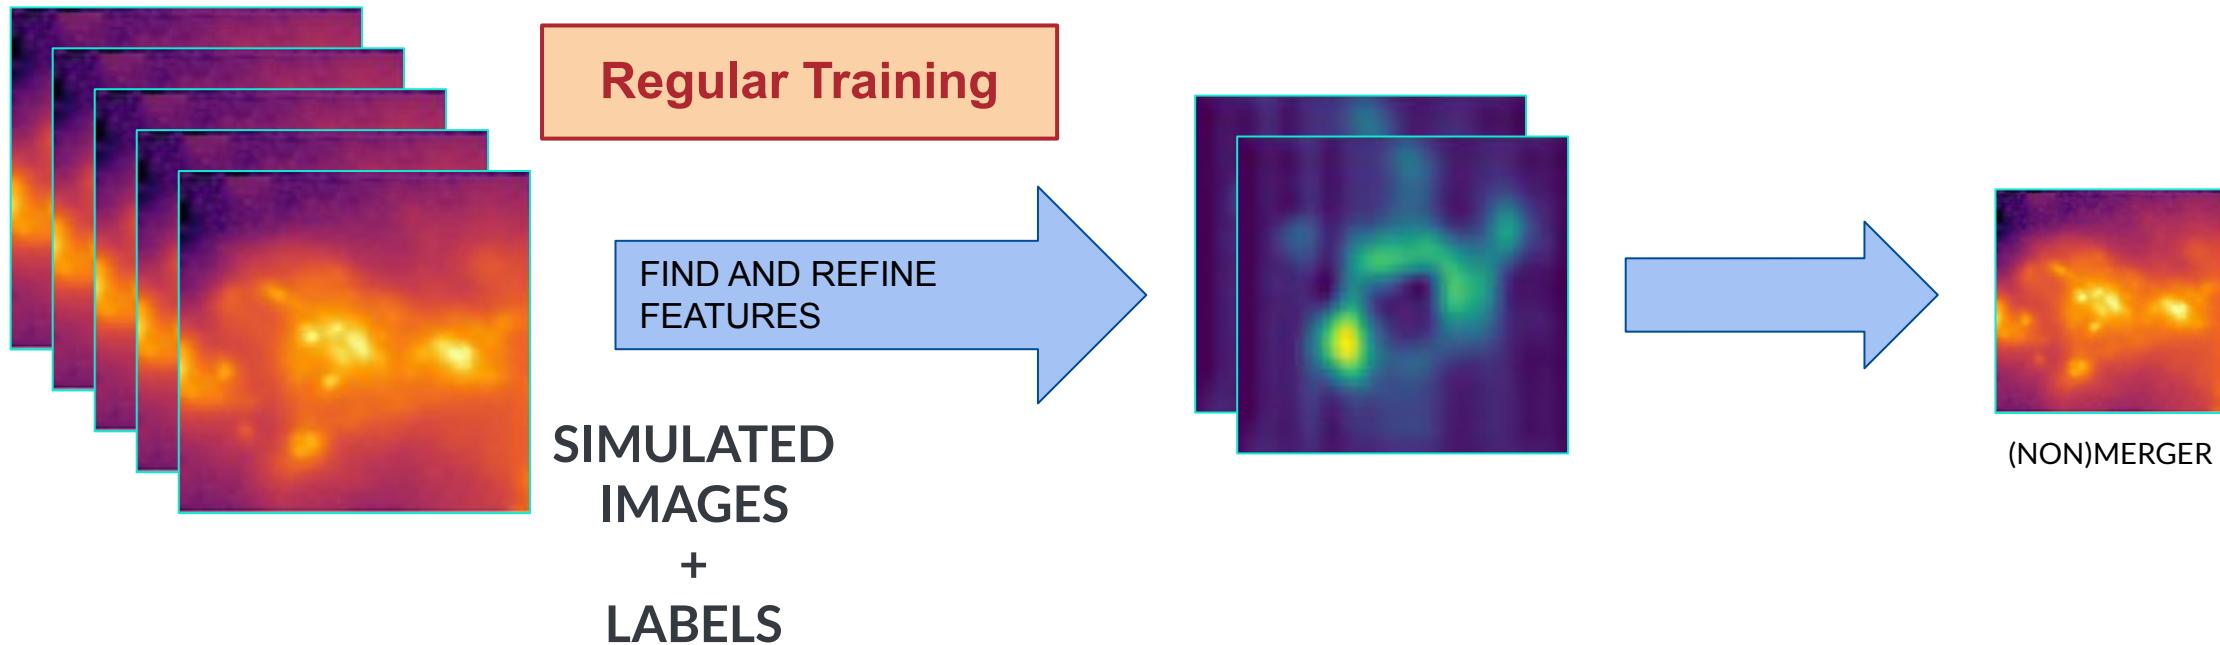
SIMULATED

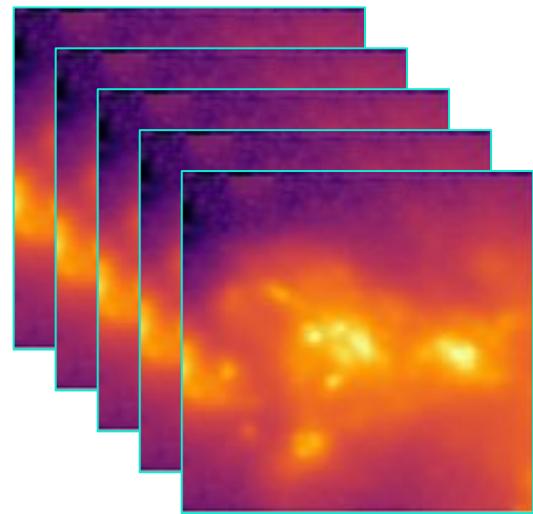
REAL



Adams et al. (2019)

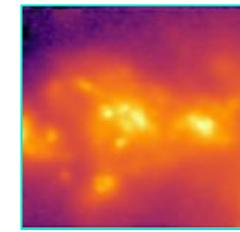
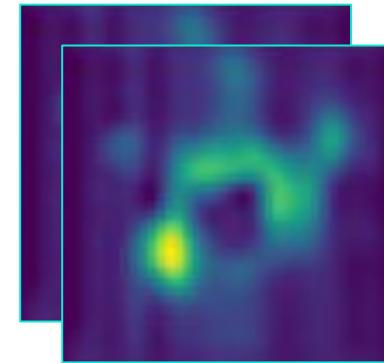


 Traditional ML

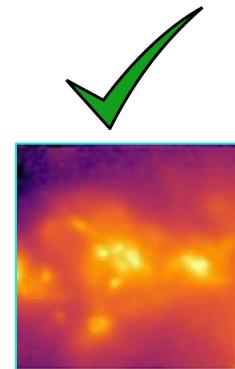
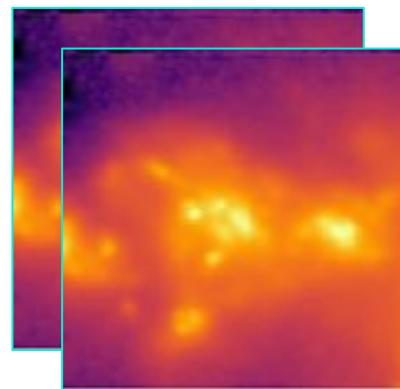
 Traditional ML

SIMULATED
IMAGES
+
LABELS

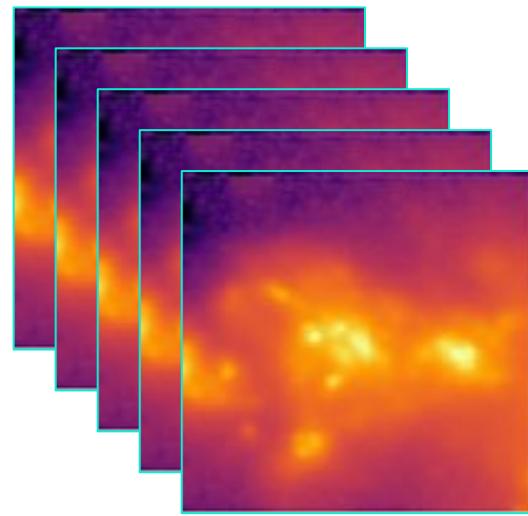
FIND AND REFINE
FEATURES



(NON)MERGER

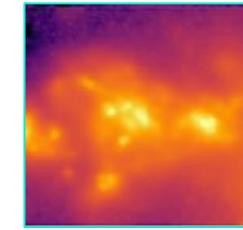
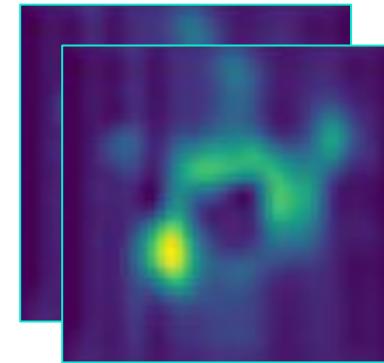


Testing the model

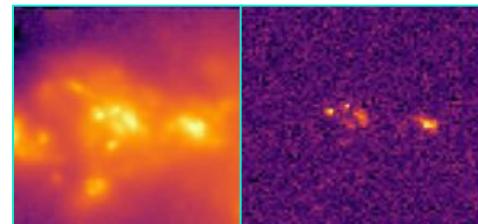
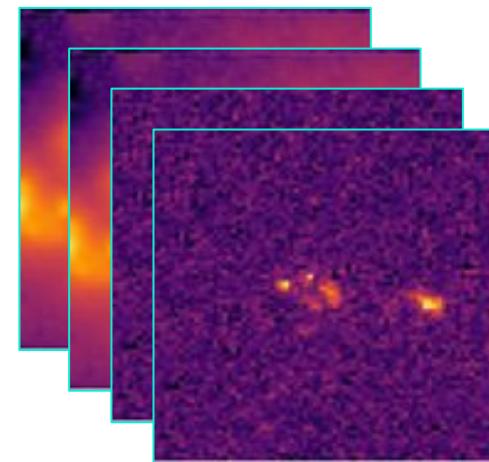
 Traditional ML

SIMULATED
IMAGES
+
LABELS

FIND AND REFINE
FEATURES



(NON)MERGER



Simulated Observed

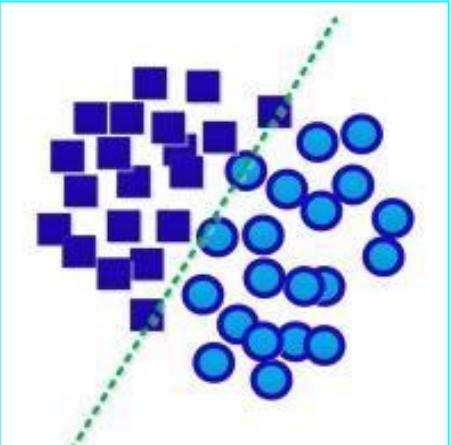
Testing the model



Why do models fail?

Train the model
on source
dataset and find
the decision
boundary.

Source Domain

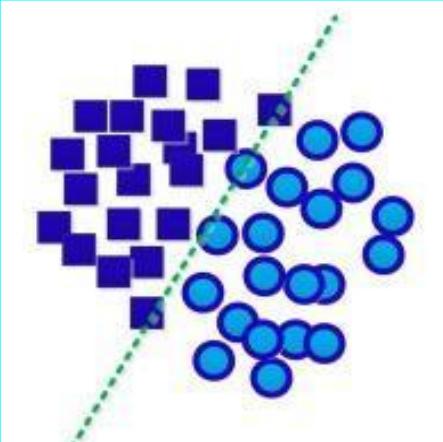




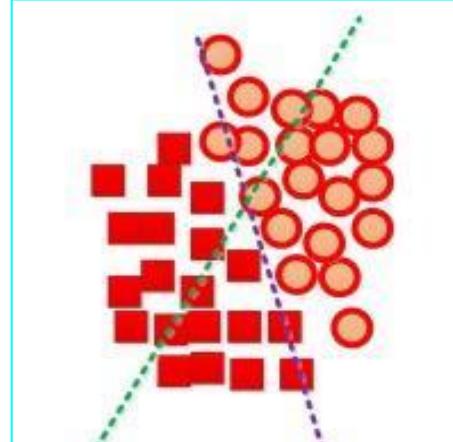
Why do models fail?

New domain is shifted,
learned decision boundary doesn't work.

Source Domain



Target Domain

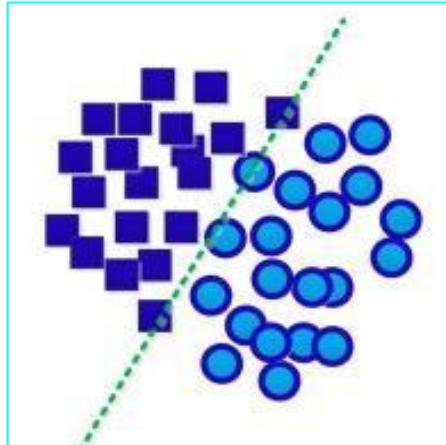




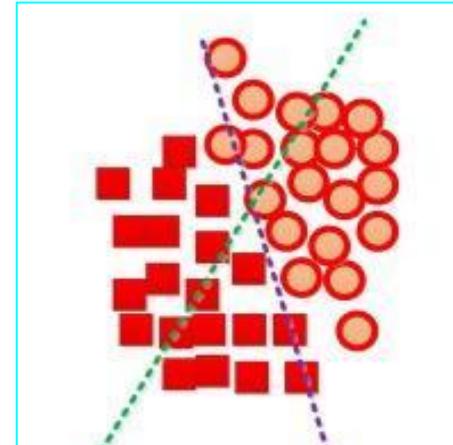
Why do models fail?

We need to align
the data during
training!

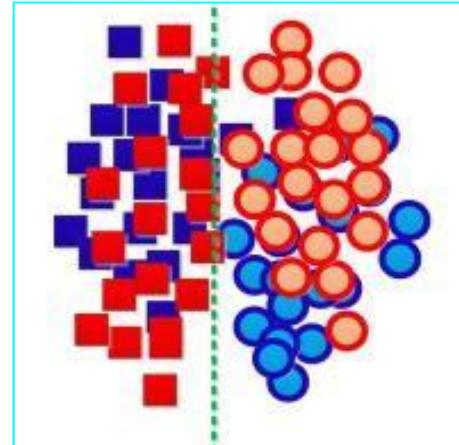
Source Domain



Target Domain

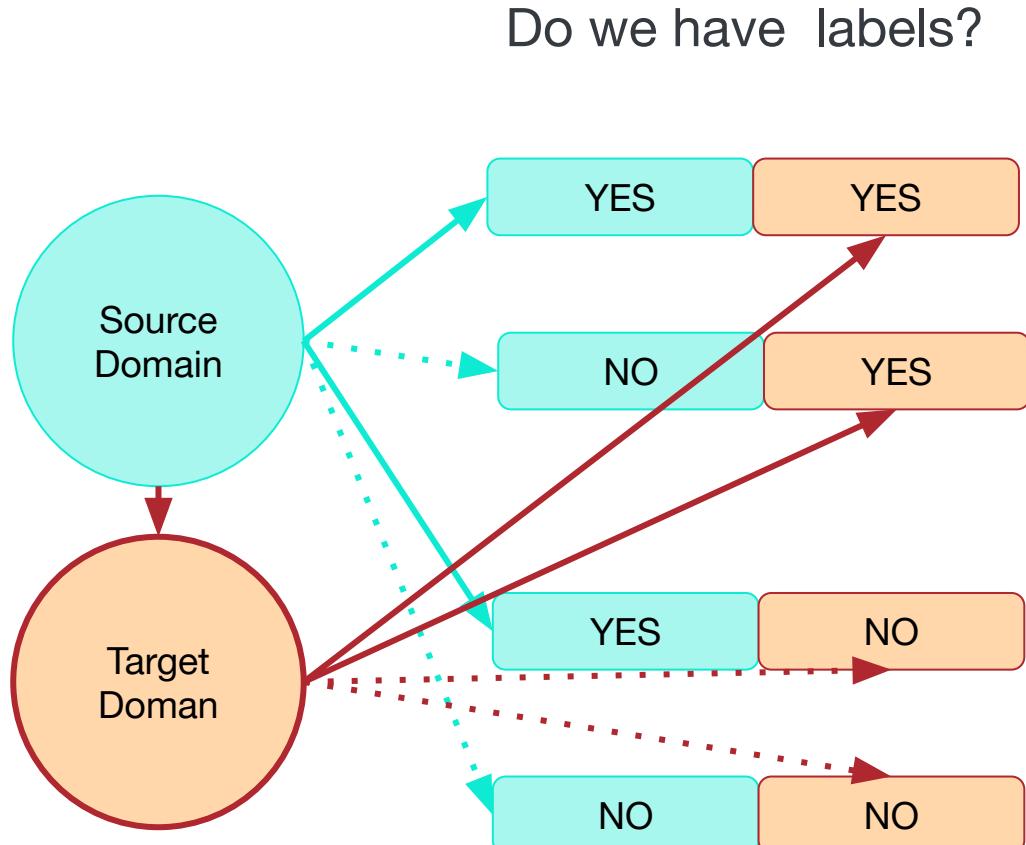


Domain Alignment





The Landscape of Multi-Dataset Learning



Inductive transfer learning. Multi-task learning

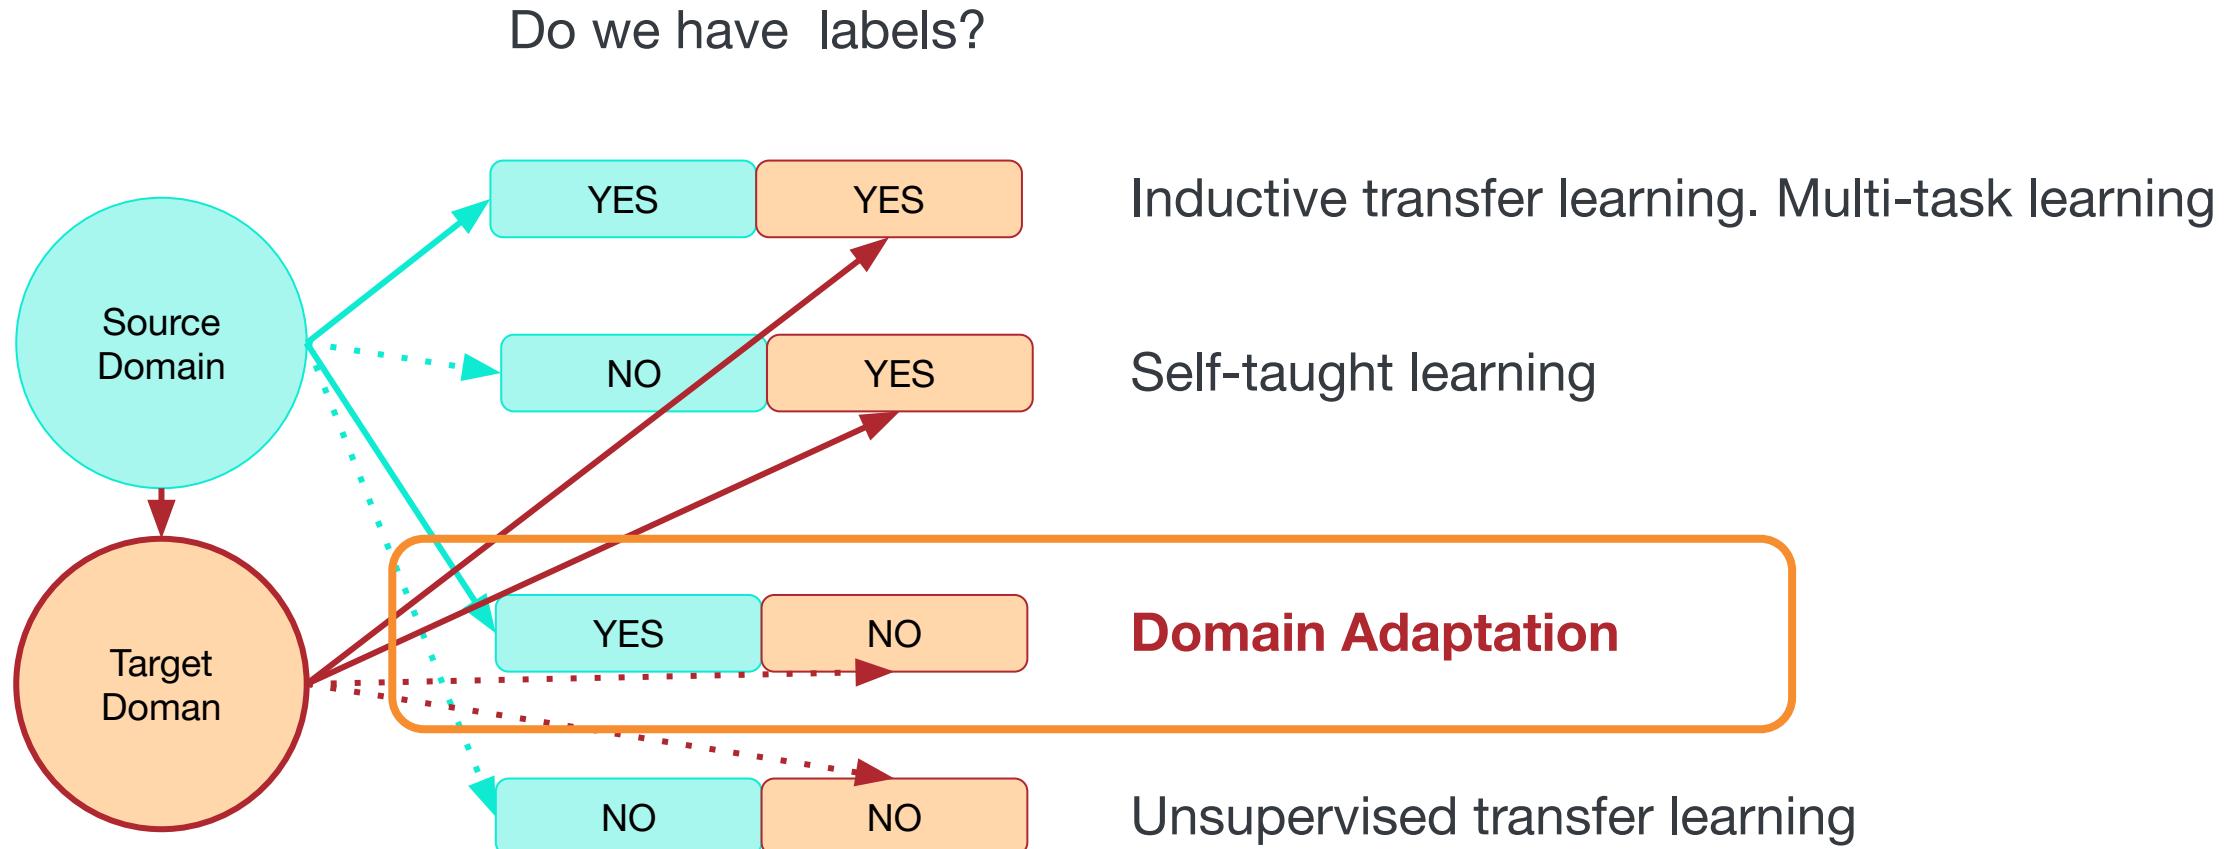
Self-taught learning

Domain Adaptation

Unsupervised transfer learning

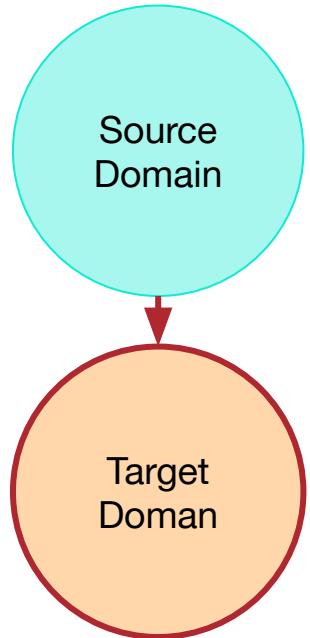


The Landscape of Multi-Dataset Learning





Datasets and Types of Shifts



$$\mathcal{D}_S = \{(x_i^S, y_i^S)\}_{i=1}^{n_S}$$

with n_S labelled samples, from joint distribution $P_S(x, y)$

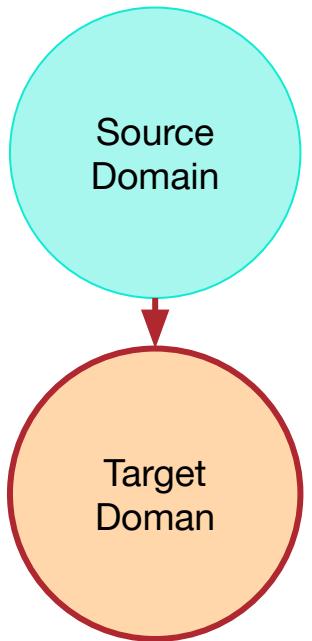
$$\mathcal{D}_T = \{x_i^T\}_{i=1}^{n_T}$$

with n_T unlabelled samples, from joint distribution $P_T(x)$

$$P_S(x, y) \neq P_T(x, y)$$

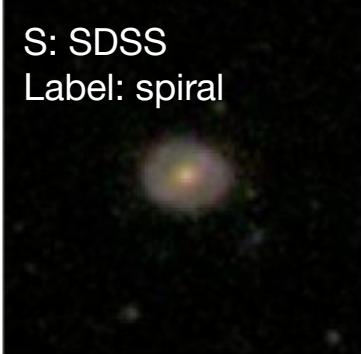


Datasets and Types of Shifts



$$\mathcal{D}_S = \{(x_i^S, y_i^S)\}_{i=1}^{n_S}$$

with n_S labelled samples, from joint distribution $P_S(x, y)$



$$\mathcal{D}_T = \{x_i^T\}_{i=1}^{n_T}$$

with n_T unlabelled samples, from joint distribution $P_T(x)$

$$P_S(x, y) \neq P_T(x, y)$$



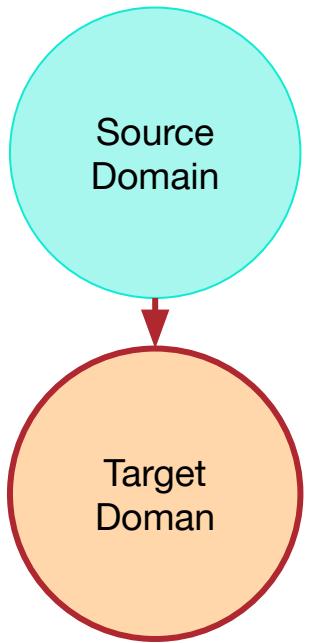
Label Shift:

$$P_s(Y) \neq P_t(Y) \quad \text{while} \quad P_s(X | Y) = P_t(X | Y)$$

Conditional distribution of data given a label stays the same, but the label distribution changes i.e., the frequency of classes differs across domains (e.g., class imbalance), but each class "looks" the same in both domains.

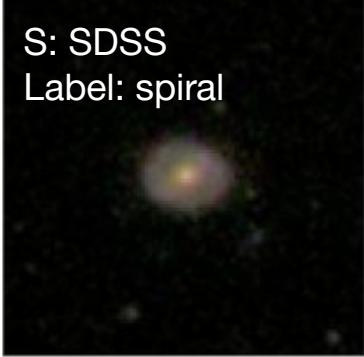


Datasets and Types of Shifts



$$\mathcal{D}_S = \{(x_i^S, y_i^S)\}_{i=1}^{n_S}$$

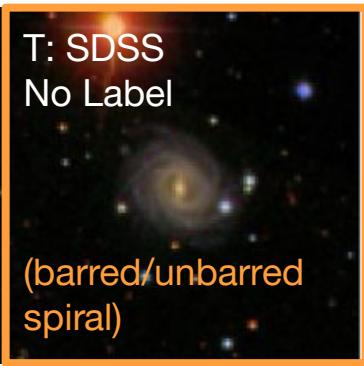
with n_S labelled samples, from joint distribution $P_S(x, y)$



$$\mathcal{D}_T = \{x_i^T\}_{i=1}^{n_T}$$

with n_T unlabelled samples, from joint distribution $P_T(x)$

$$P_S(x, y) \neq P_T(x, y)$$

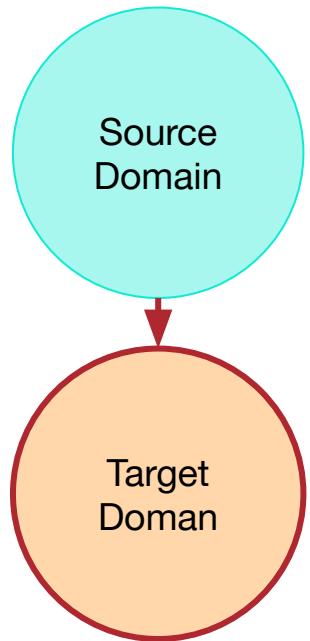


Concept Shift: $P_s(X) = P_t(X)$ but $P_s(Y | X) \neq P_t(Y | X)$

The input distribution of data stays the same, but the mapping from input to label changes i.e., the same input might be labeled differently across domains (e.g., due to different labeling policies or evolving concepts).

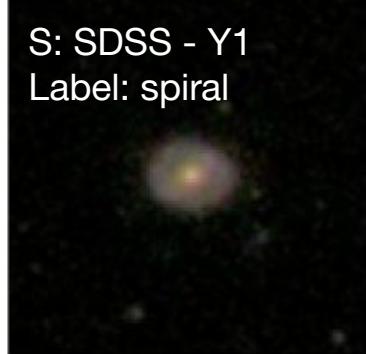


Datasets and Types of Shifts



$$\mathcal{D}_S = \{(x_i^S, y_i^S)\}_{i=1}^{n_S}$$

with n_S labelled samples, from joint distribution $P_S(x, y)$



$$\mathcal{D}_T = \{x_i^T\}_{i=1}^{n_T}$$

with n_T unlabelled samples, from joint distribution $P_T(x)$

$$P_S(x, y) \neq P_T(x, y)$$

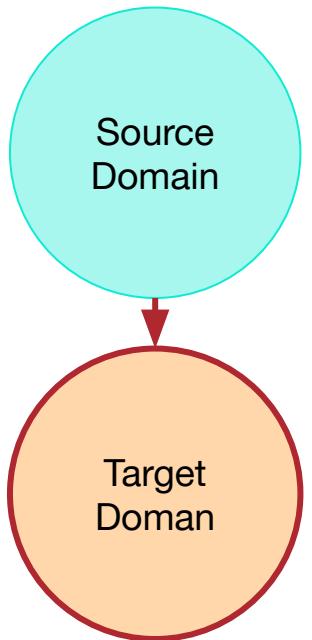


Conditional Shift: $P_S(Y) = P_T(Y)$ but $P_S(X | Y) \neq P_T(X | Y)$

The label distribution the same, but the mapping from input to label changes i.e., the same labeled corresponds to different data due to measurement differences.



Datasets and Types of Shifts



$$\mathcal{D}_S = \{(x_i^S, y_i^S)\}_{i=1}^{n_S}$$

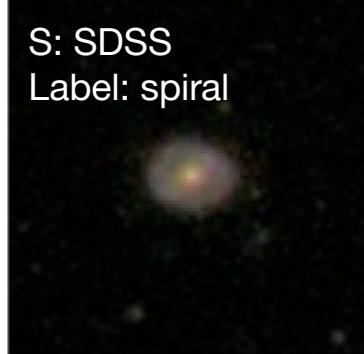
with n_S labelled samples, from joint distribution $P_S(x, y)$

$$\mathcal{D}_T = \{x_i^T\}_{i=1}^{n_T}$$

with n_T unlabelled samples, from joint distribution $P_T(x)$

$$P_S(x, y) \neq P_T(x, y)$$

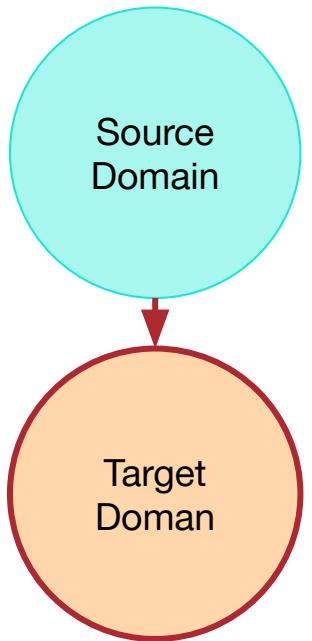
Covariate Shift: $P_s(X) \neq P_t(X)$ but $P_s(Y | X) = P_t(Y | X)$



Data distribution changes (often not just image properties but class distribution as well for example e.g., different sensors, sampling bias), but the conditional distribution of labels given the data remains the same.

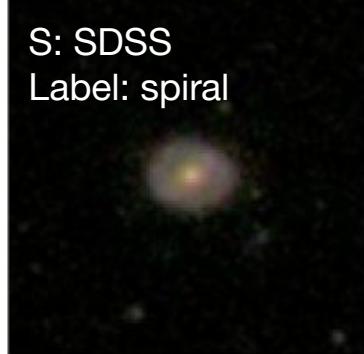


Datasets and Types of Shifts



$$\mathcal{D}_S = \{(x_i^S, y_i^S)\}_{i=1}^{n_S}$$

with n_S labelled samples, from joint distribution $P_S(x, y)$



$$\mathcal{D}_T = \{x_i^T\}_{i=1}^{n_T}$$

with n_T unlabelled samples, from joint distribution $P_T(x)$

$$P_S(x, y) \neq P_T(x, y)$$

Compound Shift:

Any potential combination of two or more shift!





Dataset Shift Quiz !

Scenario 1:

You trained a model to detect a rare disease in a hospital with a low prevalence rate (5%). It's now deployed in a different region where the prevalence is 30%. Patient features (symptoms, test values) behave similarly for each diagnosis.

Covariate

Concept / Conditional

Label

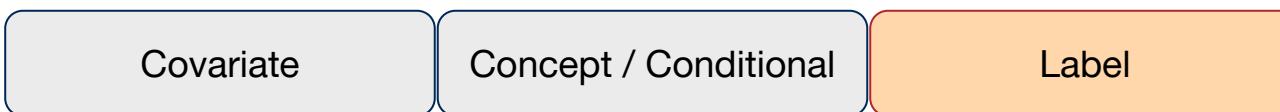
Answer:



Dataset Shift Quiz !

Scenario 1:

You trained a model to detect a rare disease in a hospital with a low prevalence rate (5%). It's now deployed in a different region where the prevalence is 30%. Patient features (symptoms, test values) behave similarly for each diagnosis.



Answer:

Label shift — $P(y)$ has changed (disease rate), but $P(x | y)$ is stable.



Dataset Shift Quiz !

Scenario 2:

Your spam classifier was trained on English-language emails. Now it's deployed to classify emails in both English and French. The spam-to-nonspam ratio is about the same, and spam emails still tend to have similar structures (e.g., lots of links, certain keywords).

Covariate

Concept / Conditional

Label

Answer:



Dataset Shift Quiz !

Scenario 2:

Your spam classifier was trained on English-language emails. Now it's deployed to classify emails in both English and French. The spam-to-nonspam ratio is about the same, and spam emails still tend to have similar structures (e.g., lots of links, certain keywords).

Covariate

Concept / Conditional

Label



Answer:

Covariate shift — the input distribution $P(x)$ changed (language), but the conditional $P(y | x)$ stayed relatively consistent.

Dataset Shift Quiz !

Scenario 3:

Your credit risk model was trained 5 years ago. Now, due to economic changes and inflation, people with low income are more likely to default than they used to. Income distribution remains unchanged.



Answer:



Dataset Shift Quiz !

Scenario 3:

Your credit risk model was trained 5 years ago. Now, due to economic changes and inflation, people with low income are more likely to default than they used to. Income distribution remains unchanged.



Answer:

Concept shift — the relationship $P(y | x)$ changed due to evolving economic factors even though $P(x)$ remained the same.



Dataset Shift Quiz !

Scenario 4:

You trained a model on a balanced dataset of merging and non-merging galaxies. At test time, your realistic catalog contains 90% of the non-merging galaxy images, and the remaining 10% are merging. The images for each galaxy still look the same.

Covariate

Concept / Conditional

Label

Answer:



Dataset Shift Quiz !

Scenario 4:

You trained a model on a balanced dataset of merging and non-merging galaxies. At test time, your realistic catalog contains 90% of the non-merging galaxy images, and the remaining 10% are merging. The images for each galaxy still look the same.

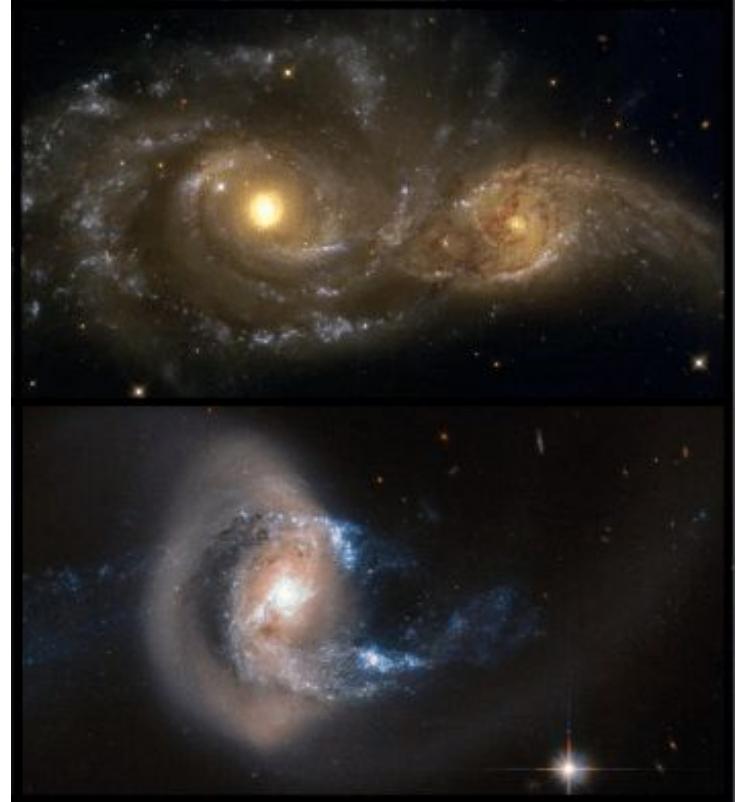
Covariate

Concept / Conditional

Label

Answer:

Label shift — the label distribution $P(y)$ changed (mergers are rare), but image features ($P(x | y)$) stayed the same.



Dataset Shift Quiz !

Scenario 5:

A remote sensing model is trained to classify crop types (e.g., wheat, corn, rice) using satellite imagery from **Region A** during a specific season. Later, the model is applied to **Region B**, where crop types and labels are the same, but **soil composition, weather patterns, and sunlight levels differ**, affecting how the crops look in satellite images

Covariate

Concept / Conditional

Label

Answer:



Dataset Shift Quiz !

Scenario 5:

A remote sensing model is trained to classify crop types (e.g., wheat, corn, rice) using satellite imagery from **Region A** during a specific season. Later, the model is applied to **Region B**, where crop types and labels are the same, but **soil composition, weather patterns, and sunlight levels differ**, affecting how the crops look in satellite images .

Covariate

Concept / Conditional

Label

Answer:

Covariate shift — the input distribution $\mathbf{P}(\mathbf{x})$ (satellite pixel features) has changed due to environmental factors, but $\mathbf{P}(\mathbf{y} | \mathbf{x})$ (how features map to crop types) is assumed to remain the same.



02

Domain Adaptation Techniques

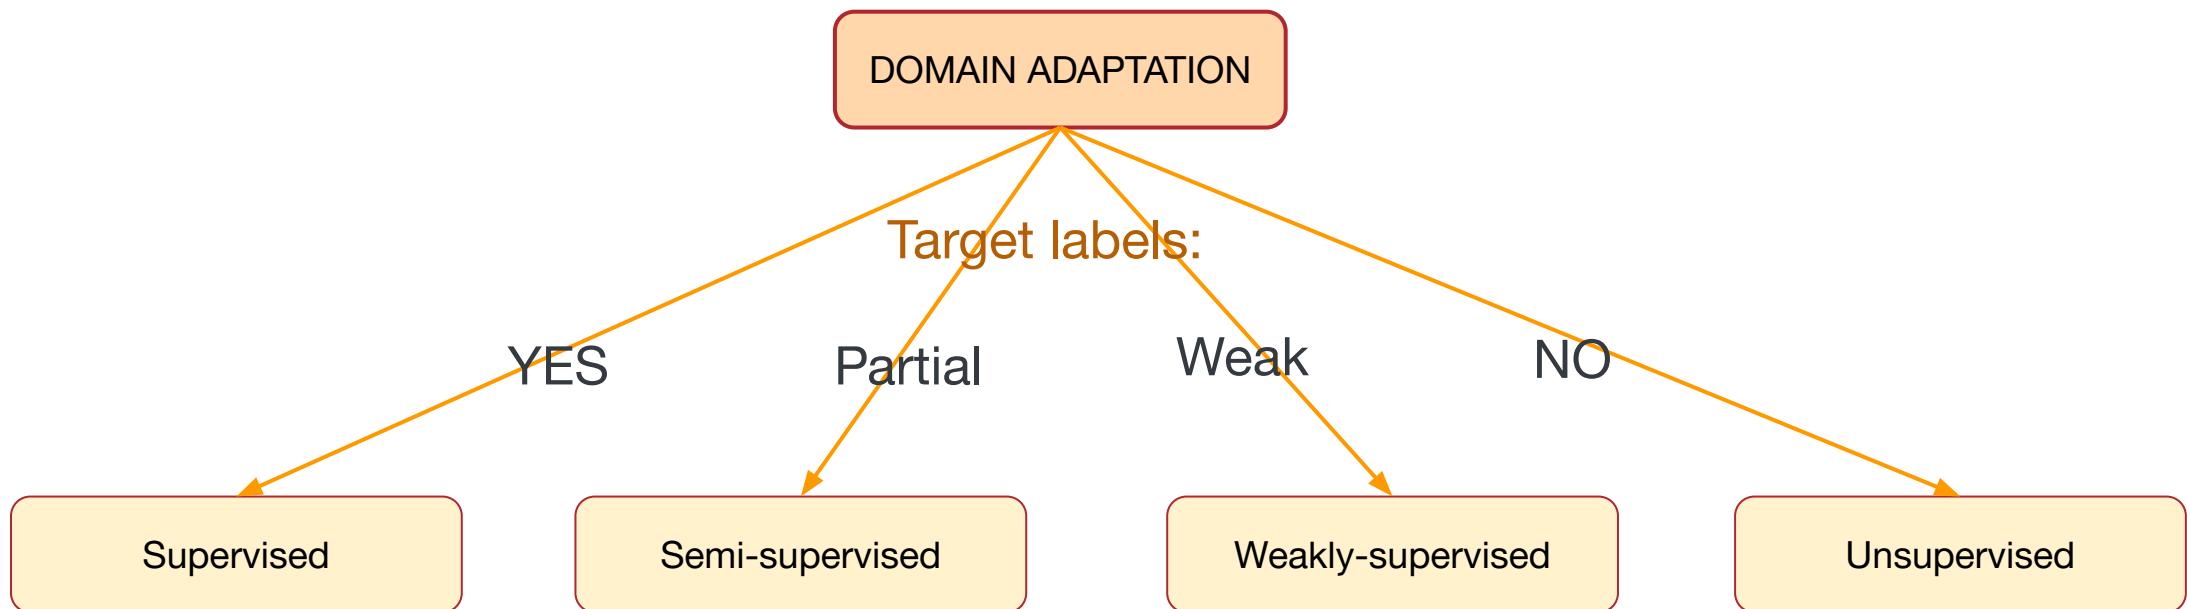
Useful Resources:

- ["A Brief Review of Domain Adaptation"](#) Farahini et al. 2020.
- ["Deep Visual Domain Adaptation: A Survey"](#) Wand & Deng 2018.
- ["Domain Adaptation for Visual Applications: A Comprehensive Survey"](#) Csurka 2017.



Types of Domain Adaptation Approaches

There is a huge number of domain adaptation methods to suit different needs.

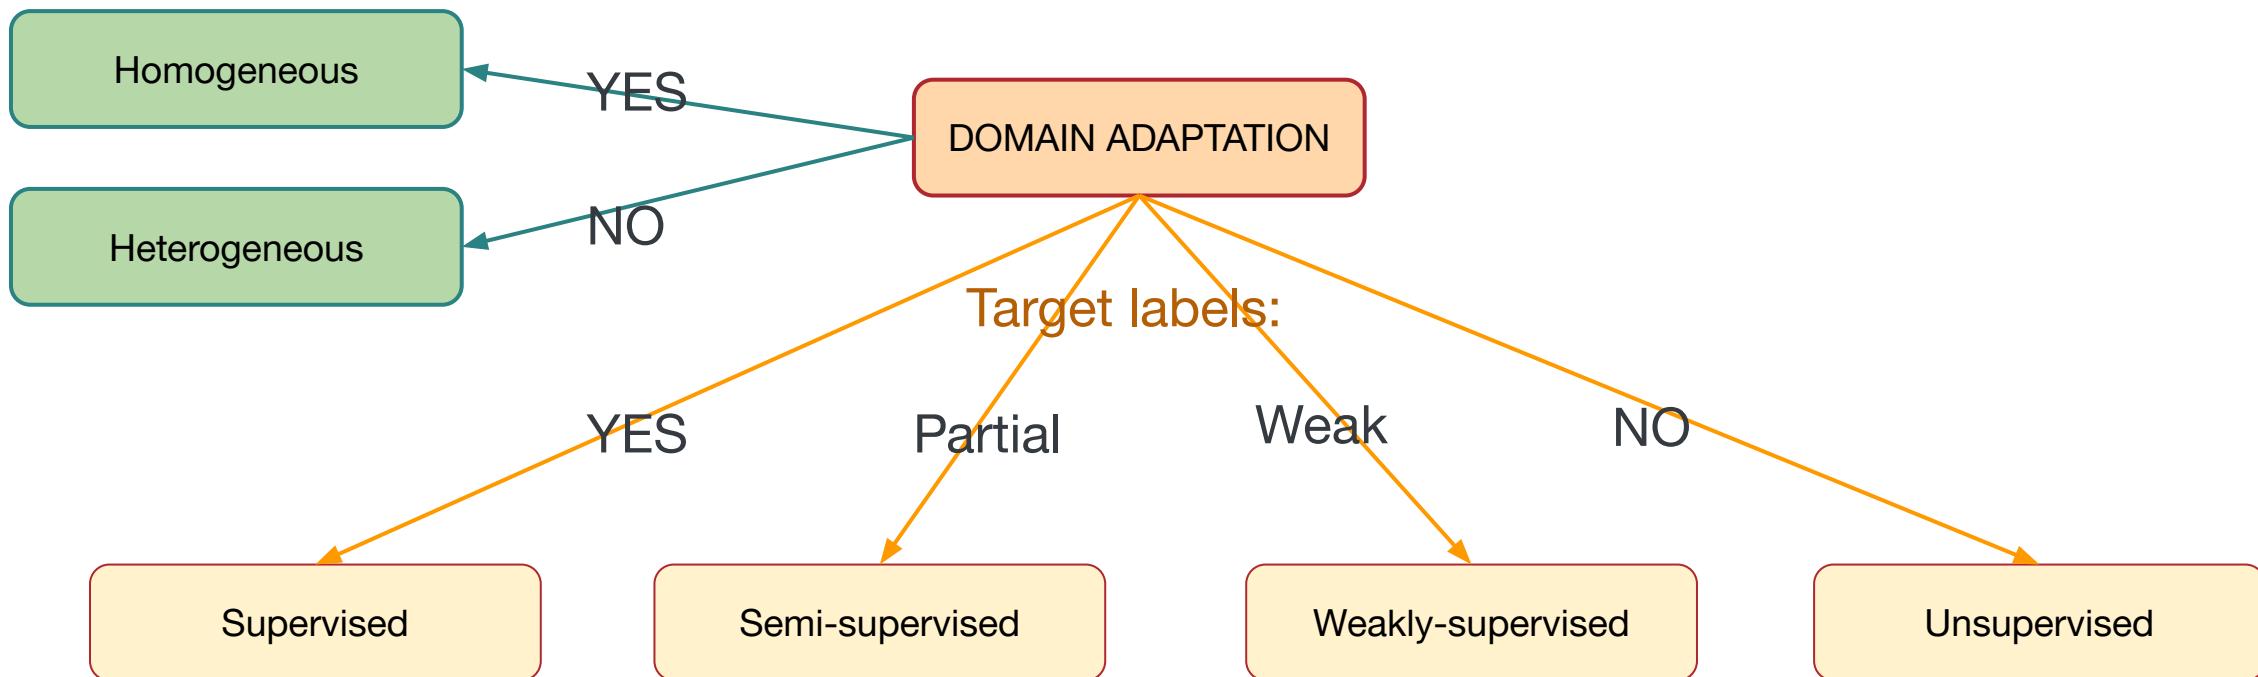




Types of Domain Adaptation Approaches

There is a huge number of domain adaptation methods to suit different needs.

Similarity of the feature space:





Types of Domain Adaptation Approaches

There is a huge number of domain adaptation methods to suit different needs.

Similarity of the feature space:

Homogeneous

YES

Heterogeneous

NO

DOMAIN ADAPTATION

Knowledge directly transferred:

One-step

YES

Multi-step

NO

Target labels:

Partial

Weak

NO

YES

Supervised

Semi-supervised

Weakly-supervised

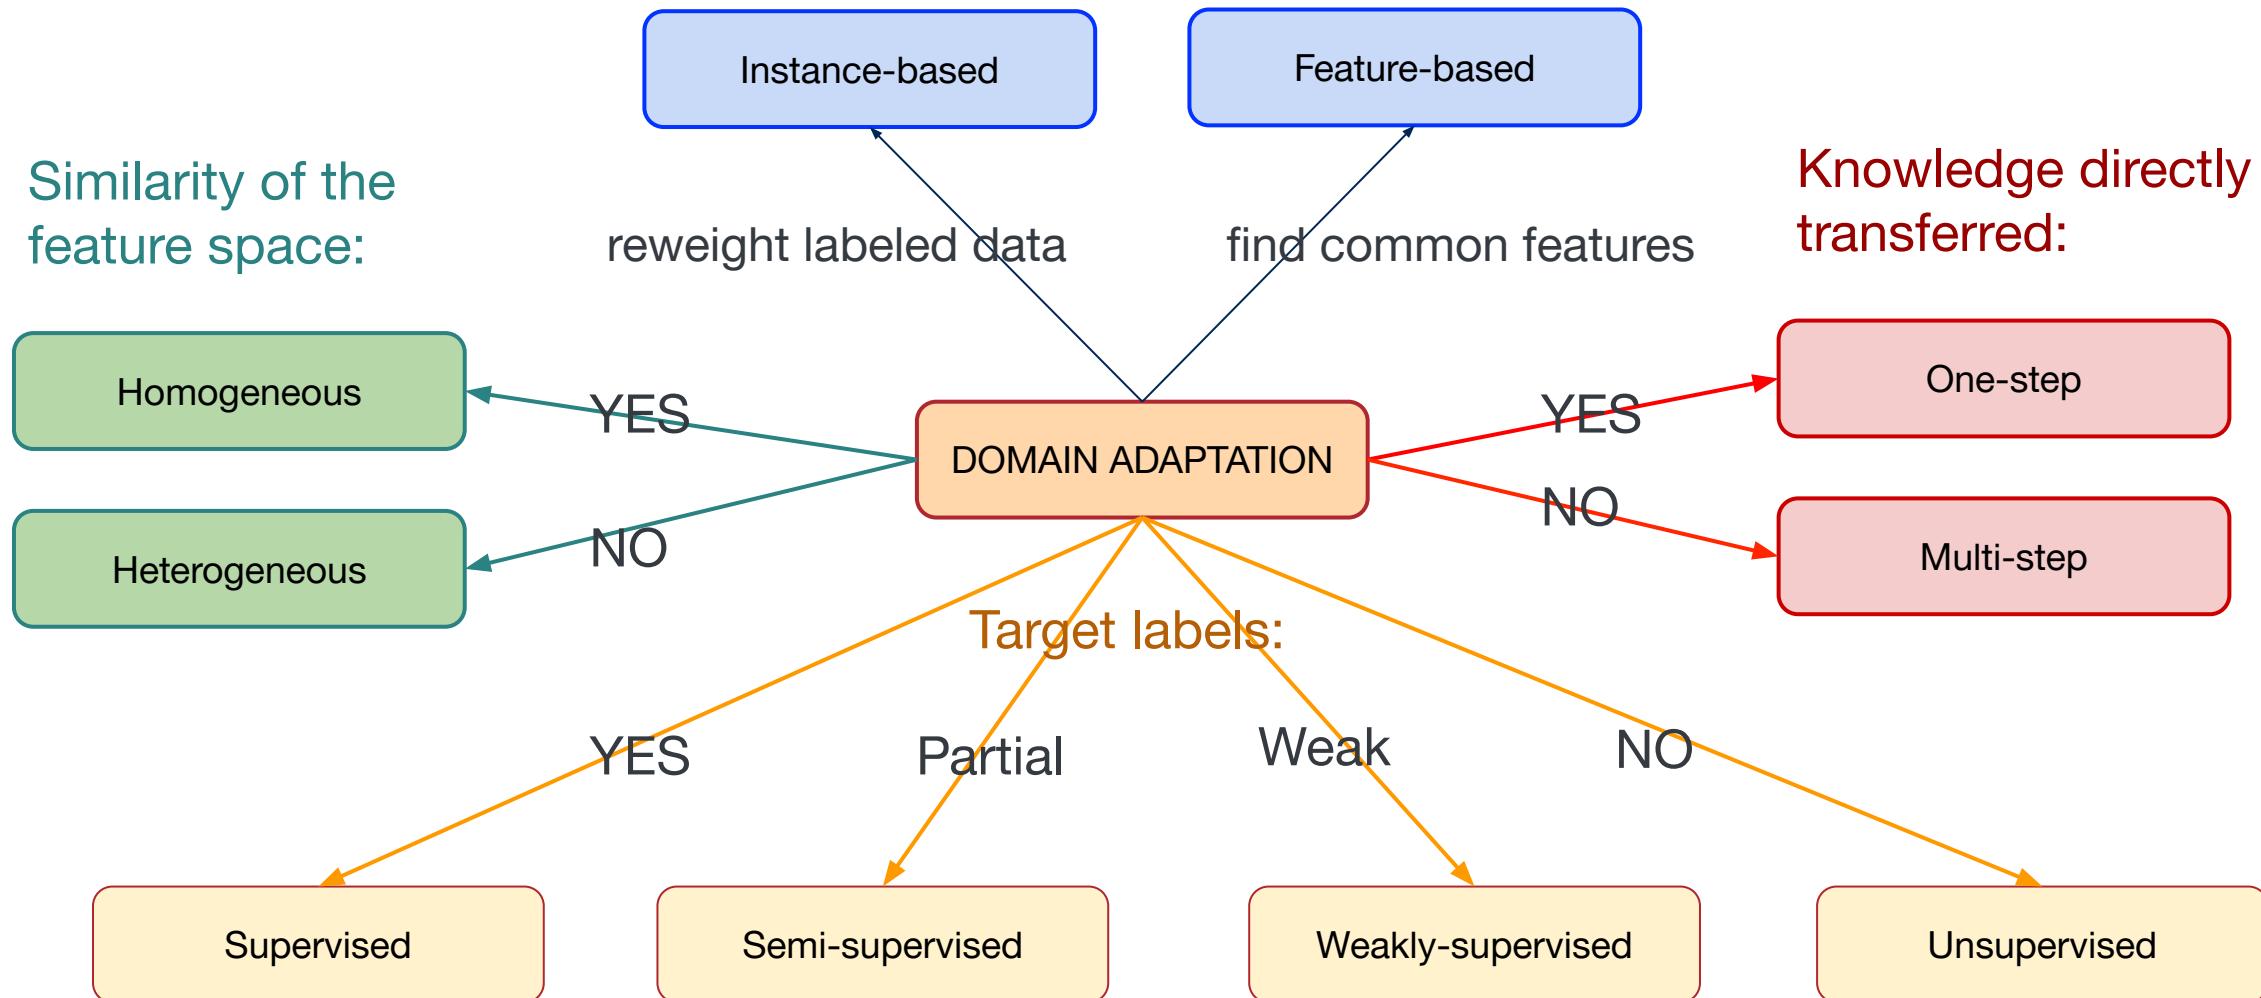
Unsupervised



Types of Domain Adaptation Approaches

There is a huge number of domain adaptation methods to suit different needs.

Alignment approach:





Domain Adaptation Quiz !

Scenario 1:

You trained a cosmic ray classifier using labeled simulated data. You also have a small set of real cosmic ray events with noisy labels from an old detector. What is this setup best described as?

Supervised

Weakly supervised

Unsupervised



Answer:



Domain Adaptation Quiz !

Scenario 1:

You trained a cosmic ray classifier using labeled simulated data. You also have a small set of real cosmic ray events with noisy labels from an old detector. What is this setup best described as?

Supervised

Weakly supervised

Unsupervised



Answer:

You have real data labels, but they are **noisy or imprecise** — so you need to use **weak supervision**. Supervised DA assumes accurate labels, while unsupervised DA assumes no labels. Weak supervision methods handle uncertainty, label correction, or probabilistic label modeling.



Domain Adaptation Quiz !

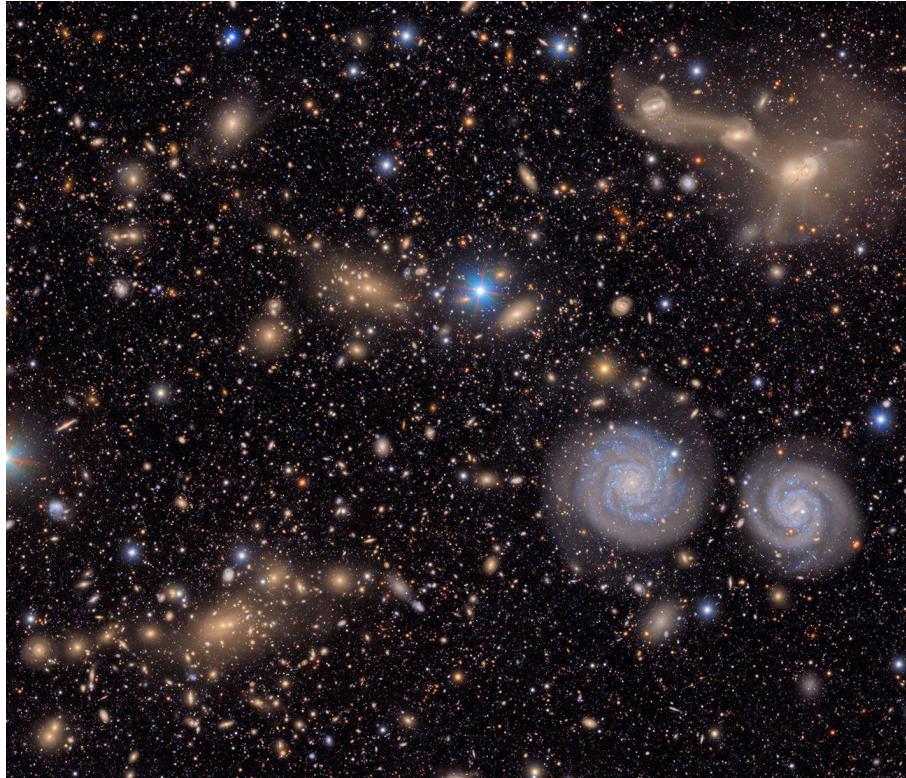
Scenario 2:

You align galaxy image features across domains by first performing PCA, then minimizing some domain adaptation loss. You then fine-tune a regressor on the aligned target features. What type of method is this?

One-step, instance based

Multi-step, feature based

Multi-step, instance based



Answer:



Domain Adaptation Quiz !

Scenario 2:

You align galaxy image features across domains by first performing PCA, then minimizing some domain adaptation loss. You then fine-tune a regressor on the aligned target features. What type of method is this?

One-step, instance based

Multi-step, feature based

Multi-step, instance based



Answer:

This approach first **transforms features** (PCA), then **aligns distributions** using a DA method, and finally adapts the model — making it **multi-step and feature-based**. Instance-based methods directly reweight or select examples data points.



Domain Adaptation Quiz !

Scenario 3:

You trained a redshift estimation model using spectroscopic data (rich, high-dimensional spectra).

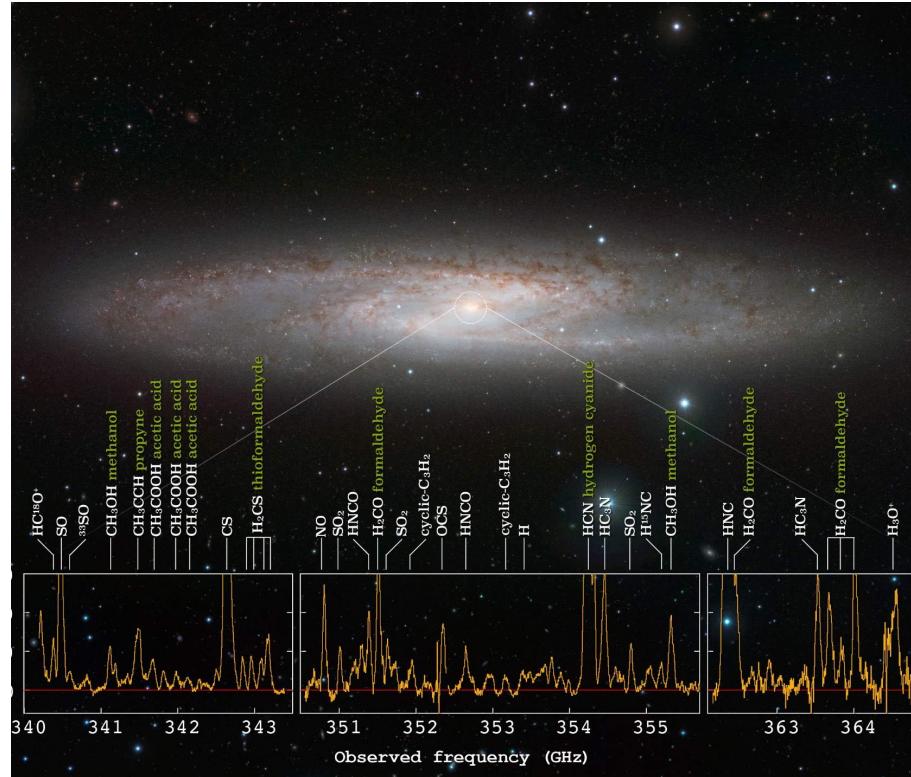
You now want to apply it to photometric data (broadband magnitudes). What kind of domain adaptation challenge is this?

Unsupervised

Homogeneous

Heterogeneous

Answer:





Domain Adaptation Quiz !

Scenario 3:

You trained a redshift estimation model using spectroscopic data (rich, high-dimensional spectra).

You now want to apply it to photometric data (broadband magnitudes). What kind of domain adaptation challenge is this?

Unsupervised

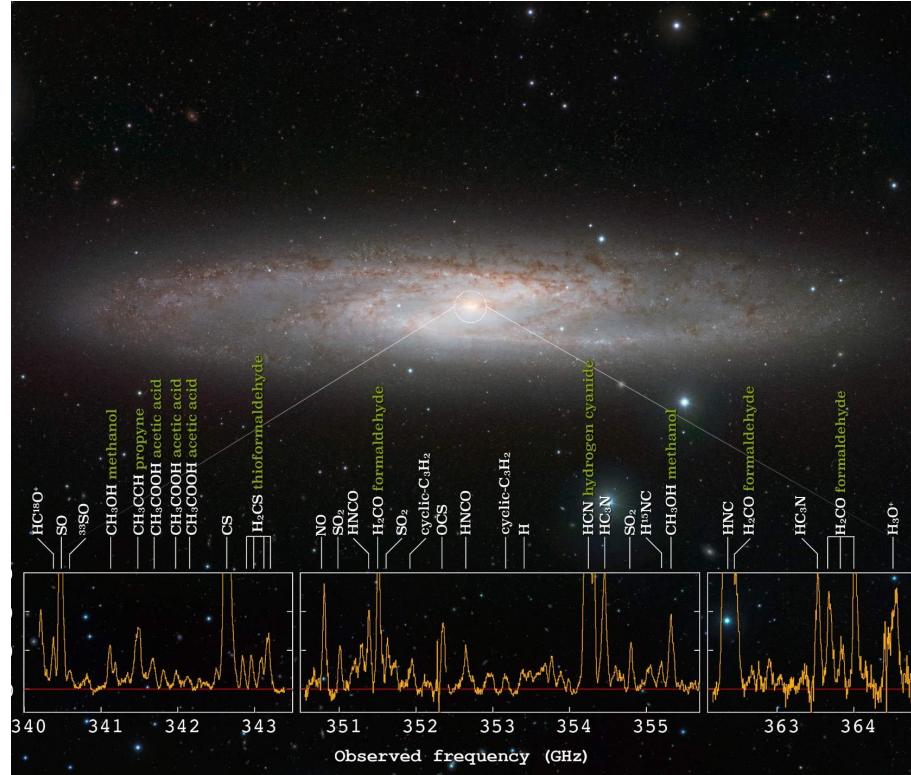
Homogeneous

Heterogeneous

Answer:

The **input features differ in type and dimension** — spectra vs. photometry — so this is a **heterogeneous DA** problem.

Homogeneous DA assumes the same input space. The key challenge here lies in mapping between incompatible feature spaces (e.g., using shared latent representations or cross-modal alignment).

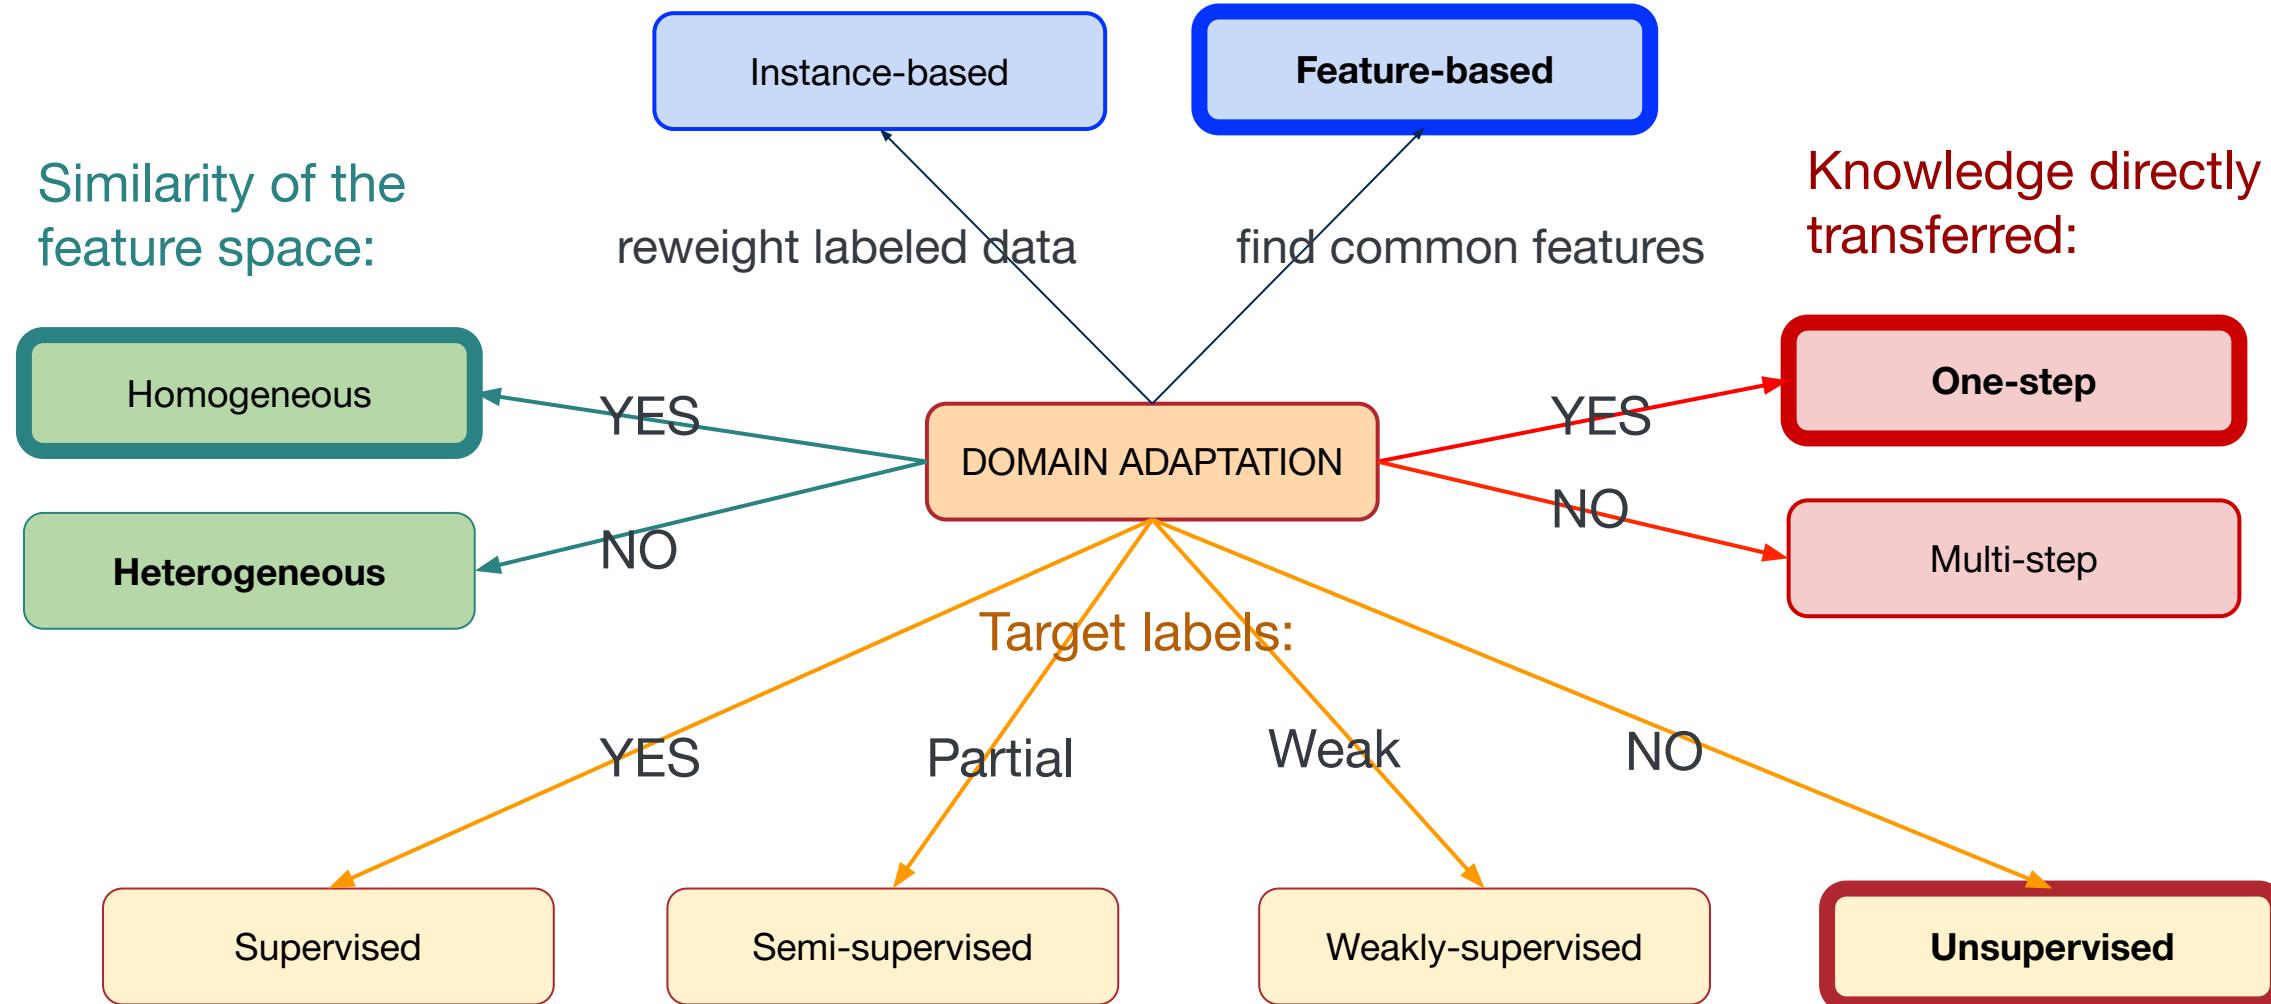




Types of Domain Adaptation Approaches

There is a huge number of domain adaptation methods to suit different needs.

Alignment approach:





Combining Datasets

DOMAIN ADAPTATION

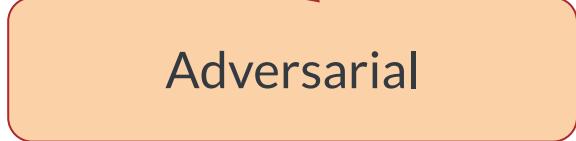
Align data distributions in the latent space of the network by forcing the network
to **find more robust domain-invariant features**.



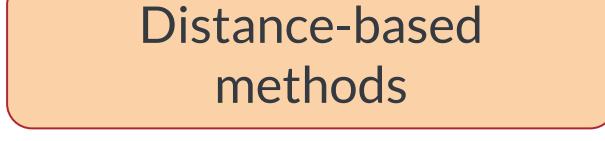
Combining Datasets

DOMAIN ADAPTATION

Align data distributions in the latent space of the network by forcing the network to **find more robust domain-invariant features**.



Ganin et al. (2016)



Gretton et al. (2012)

Minimize a distance metric between source and target data.



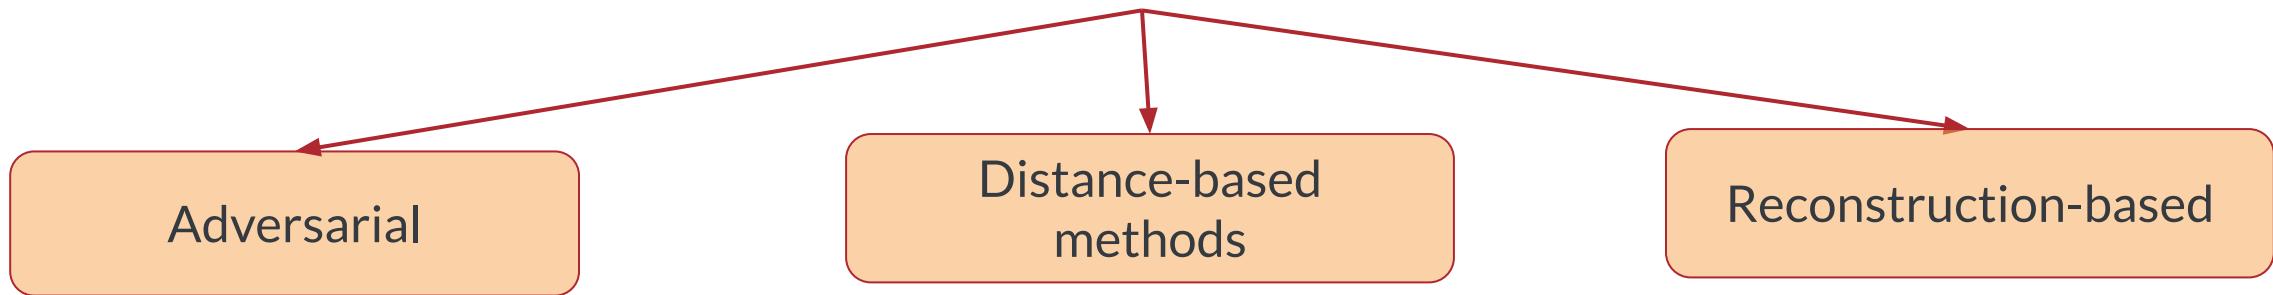
Ghifary et al. (2016)

Data reconstruction as an auxiliary task to ensure feature invariance.

Using domain discriminators to encourage domain confusion through an adversarial objective.

DOMAIN ADAPTATION

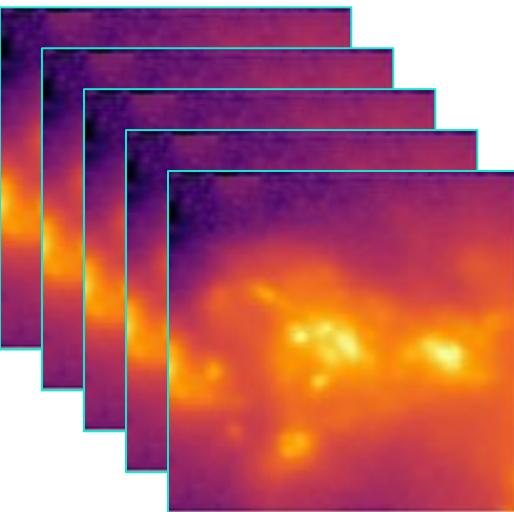
Align data distributions in the latent space of the network by forcing the network to **find more robust domain-invariant features**.



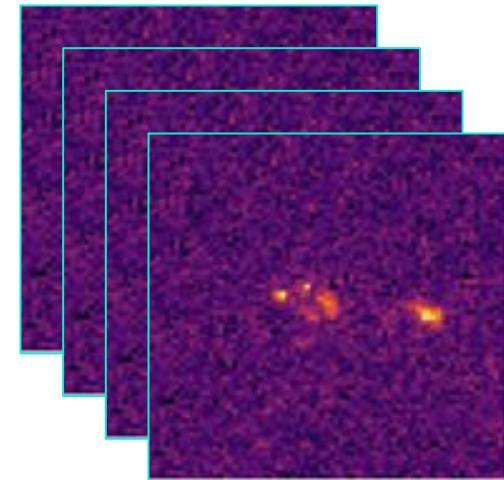
Works on **unlabeled target domain!**
Can be applied to **new data**, no need for scientists to label anything.



Training with Domain Adaptation



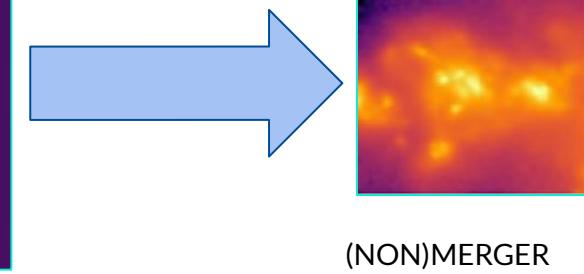
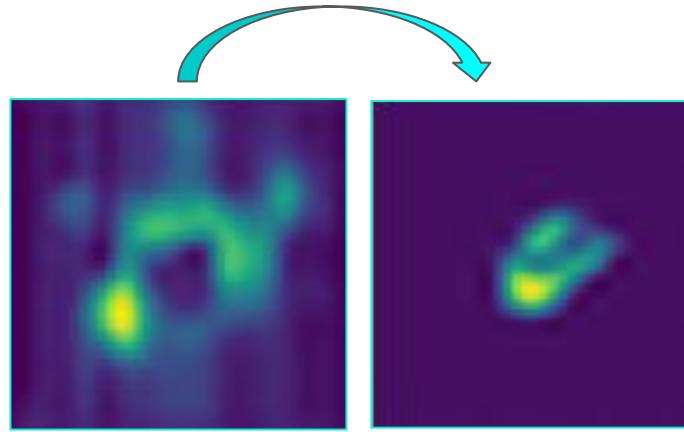
SIMULATED
IMAGES
+
LABELS



OBSERVED
IMAGES

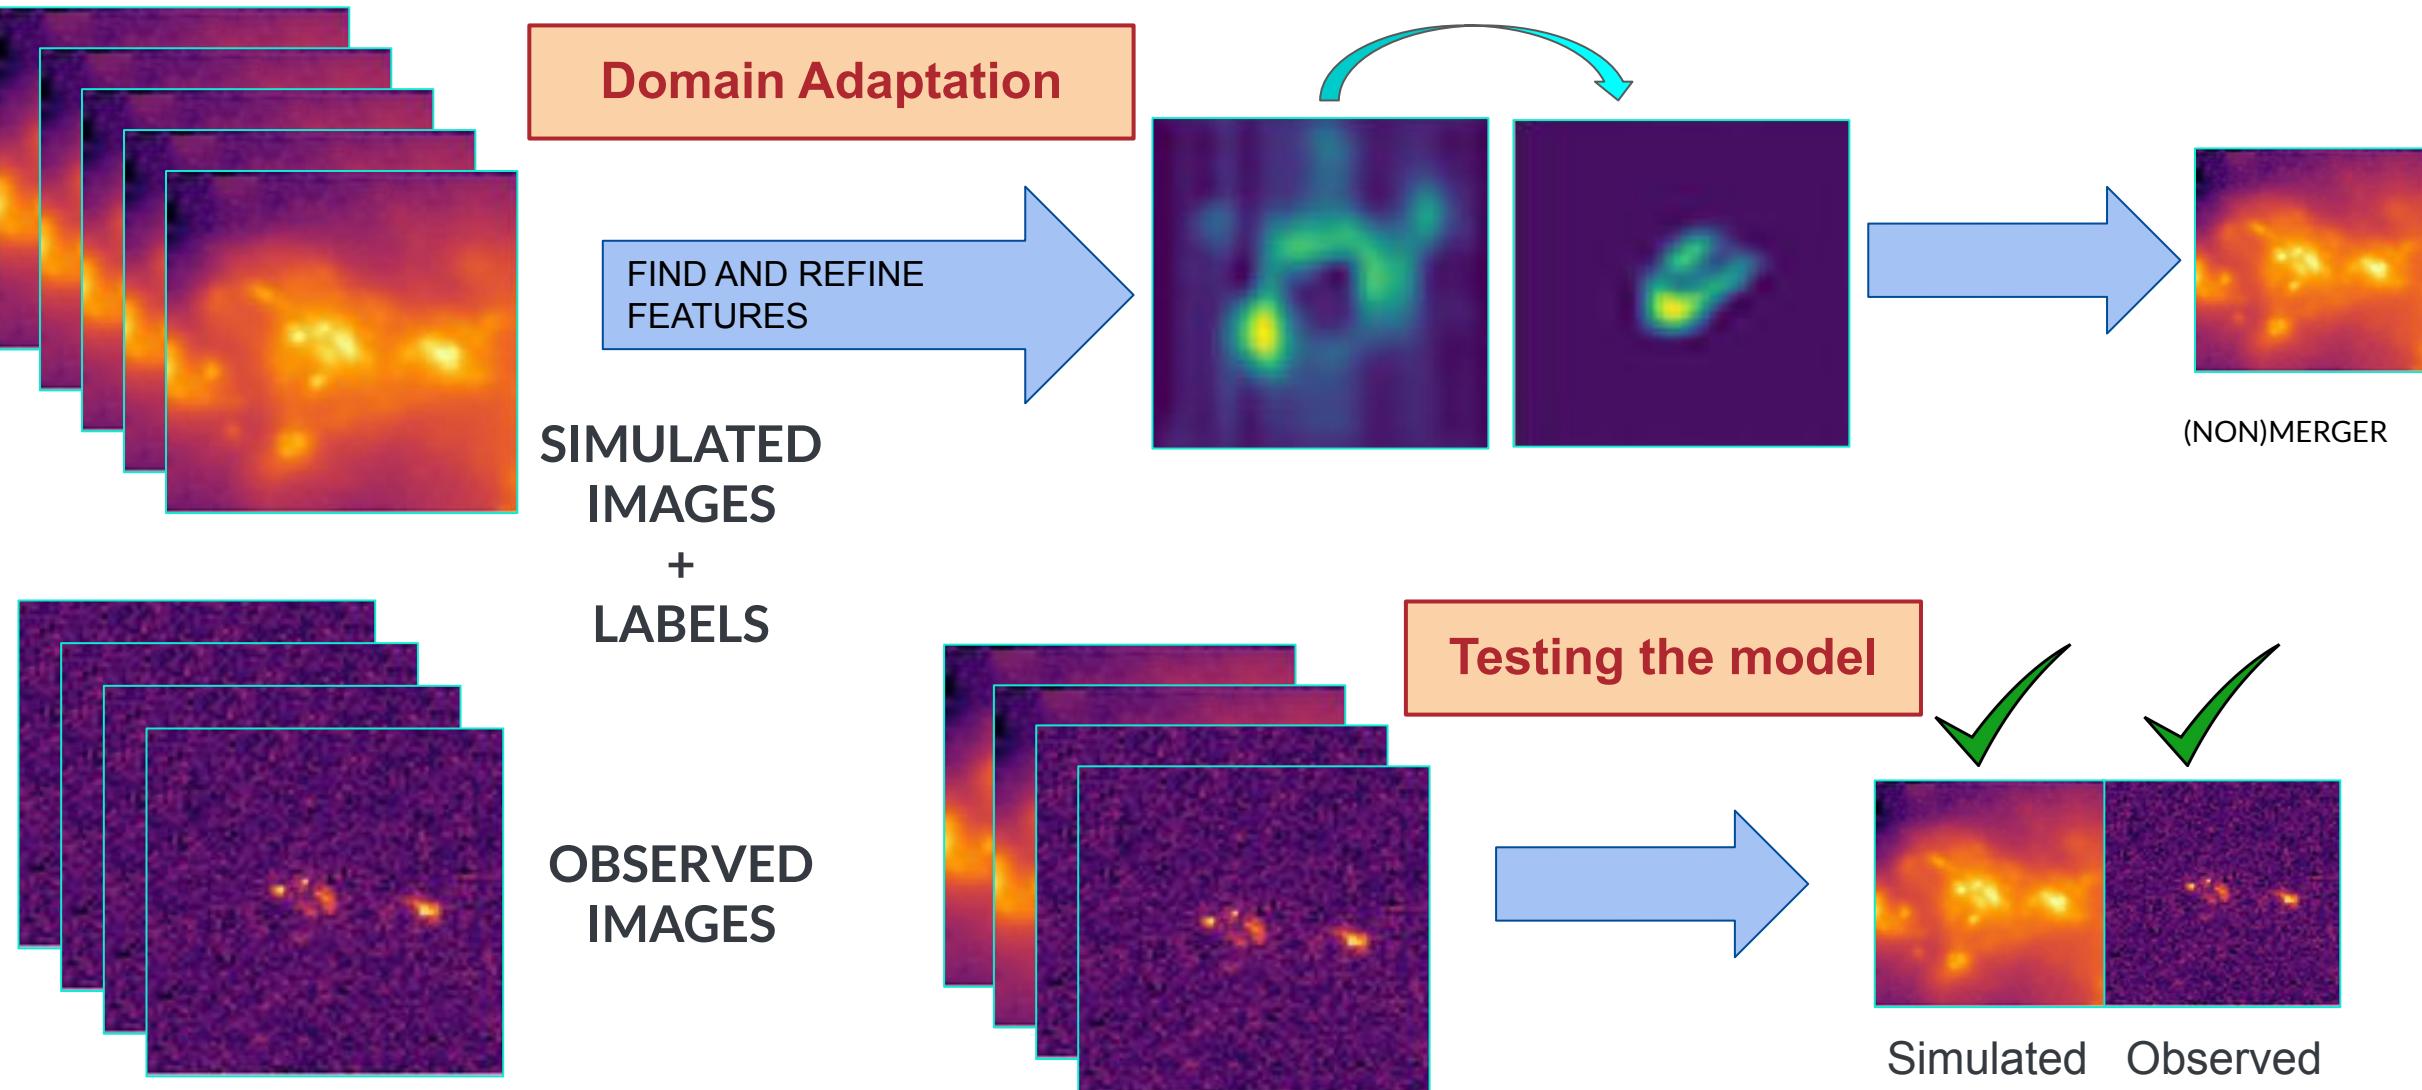
Domain Adaptation

FIND AND REFINE
FEATURES





Training with Domain Adaptation

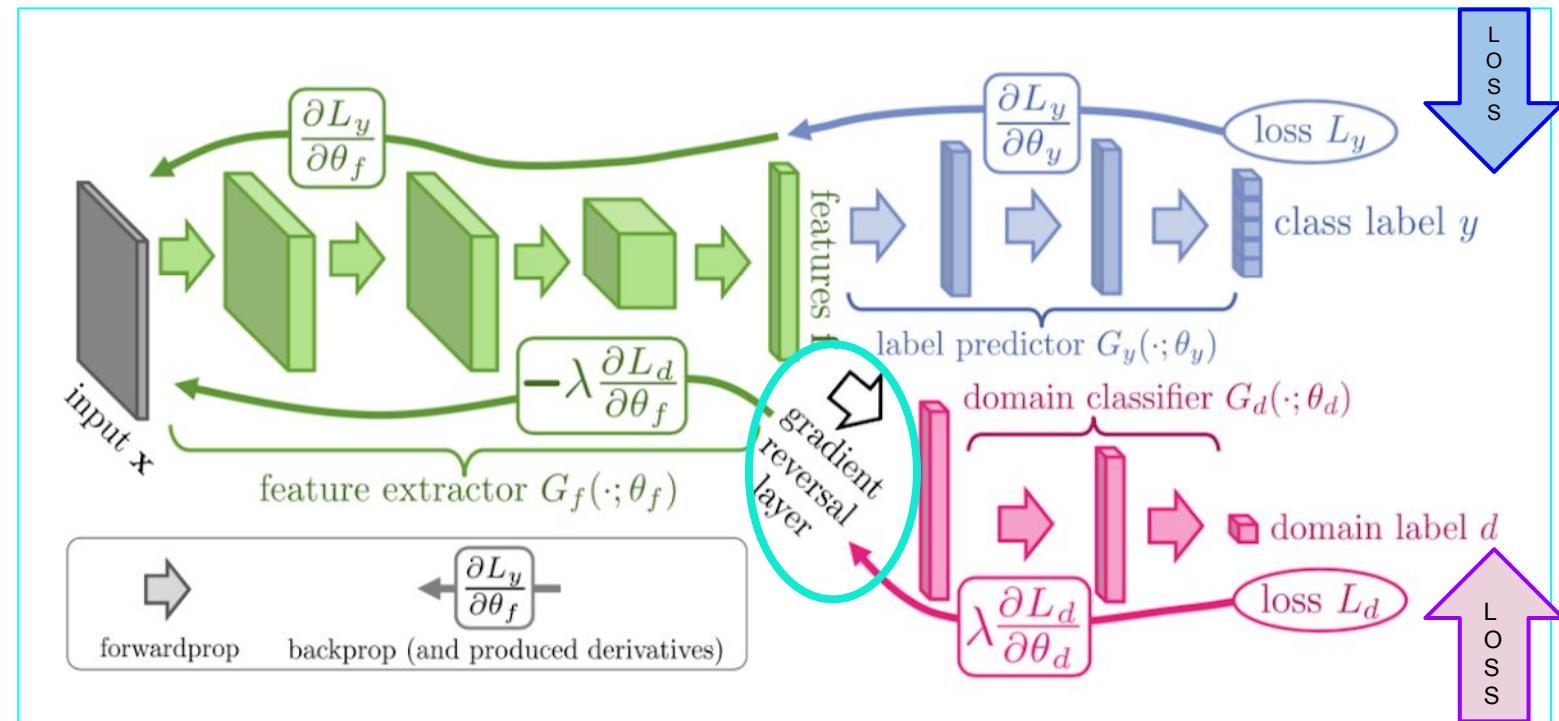


Adversarial Training

Domain Adversarial Neural Networks - DANNs

DANN - **feature extractor** + **label predictor** + **domain classifier**

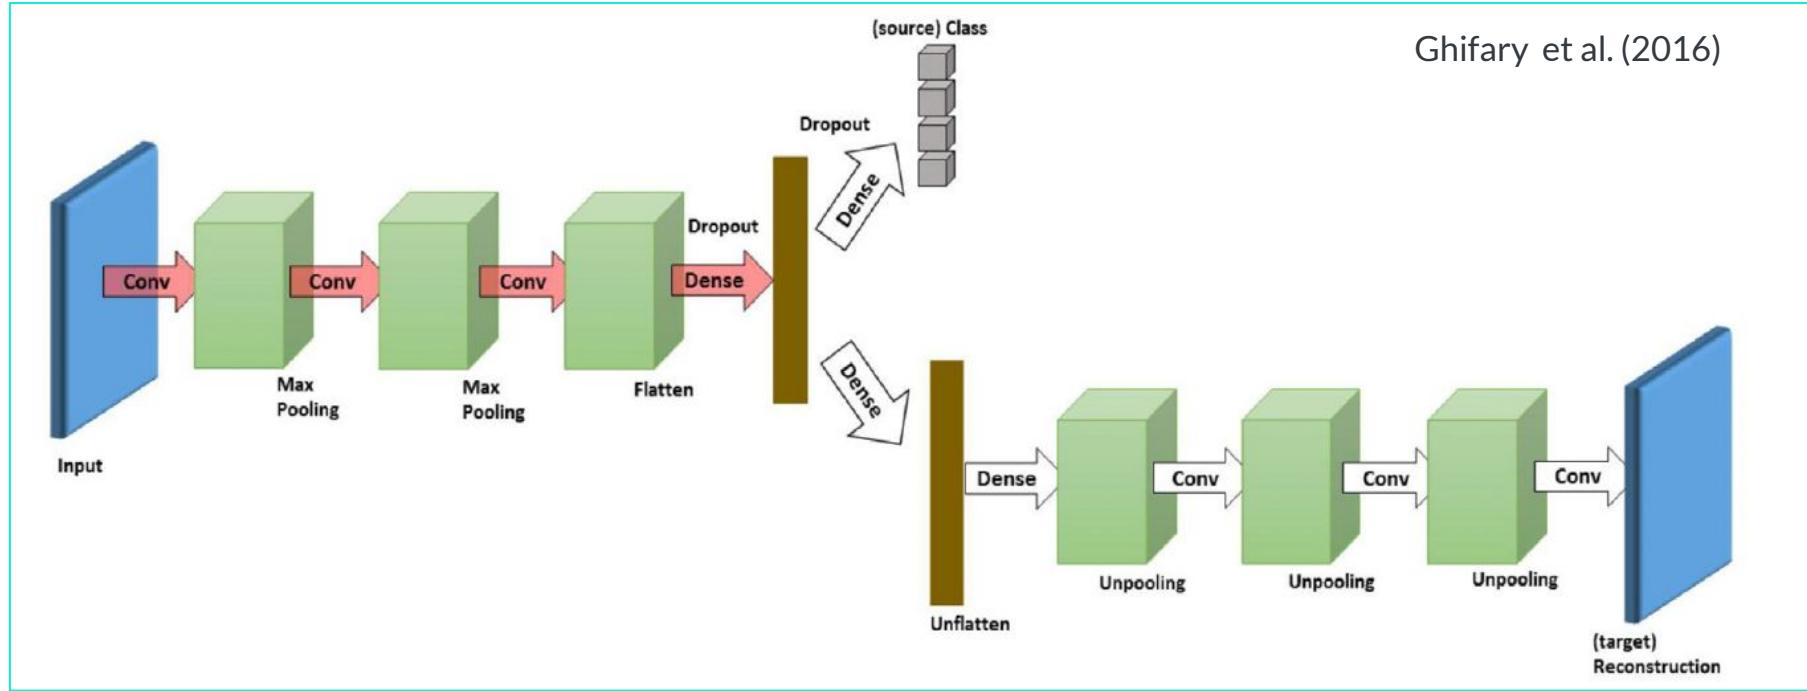
- **Gradient reversal layer** - multiplies the gradient by a negative constant during the backpropagation.
- Results in the extraction of **domain-invariant features**.
- Only source domain images are labeled during training.



Ganin et al. (2016)

Reconstruction-Based Training

Deep Reconstruction-Classification Networks - DRCN



Model jointly learns a shared encoding representation for two tasks:

- **supervised classification** of labeled source data
- **unsupervised reconstruction** of unlabeled target data.

The learnt representation not only preserves discriminability, but also encodes useful information from the target domain.



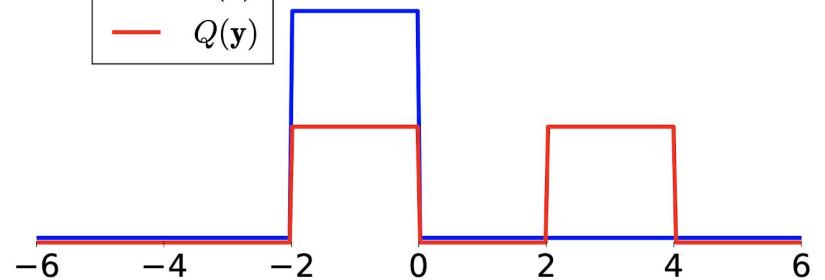
Distance-Based Training

Maximum Mean Discrepancy - MMD

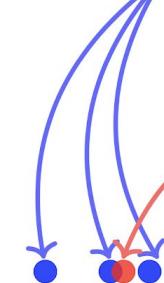
Smola et al. (2007)
Gretton et al. (2012)

Are P and Q different?

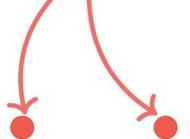
$$\begin{array}{l} \text{--- } P(x) \\ \text{--- } Q(y) \end{array}$$



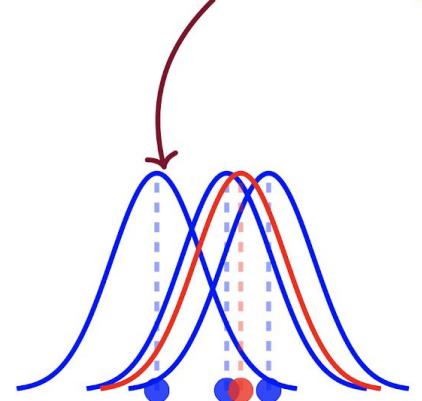
Observe $X = \{x_1, \dots, x_n\} \sim P$



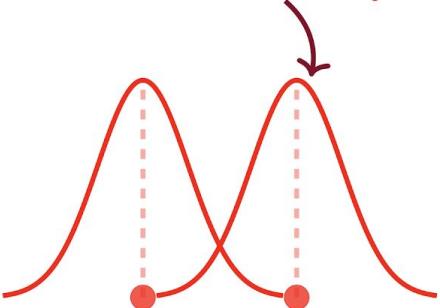
Observe $Y = \{y_1, \dots, y_n\} \sim Q$



Gaussian kernel on x_i



Gaussian kernel on y_i



$$\hat{\mu}_P(v) := \frac{1}{m} \sum_{i=1}^m k(x_i, v)$$

$\hat{\mu}_Q(v)$: mean embedding of Q

$$\underbrace{\text{witness}(v) = \hat{\mu}_P(v) - \hat{\mu}_Q(v)}_{\text{witness function}}$$

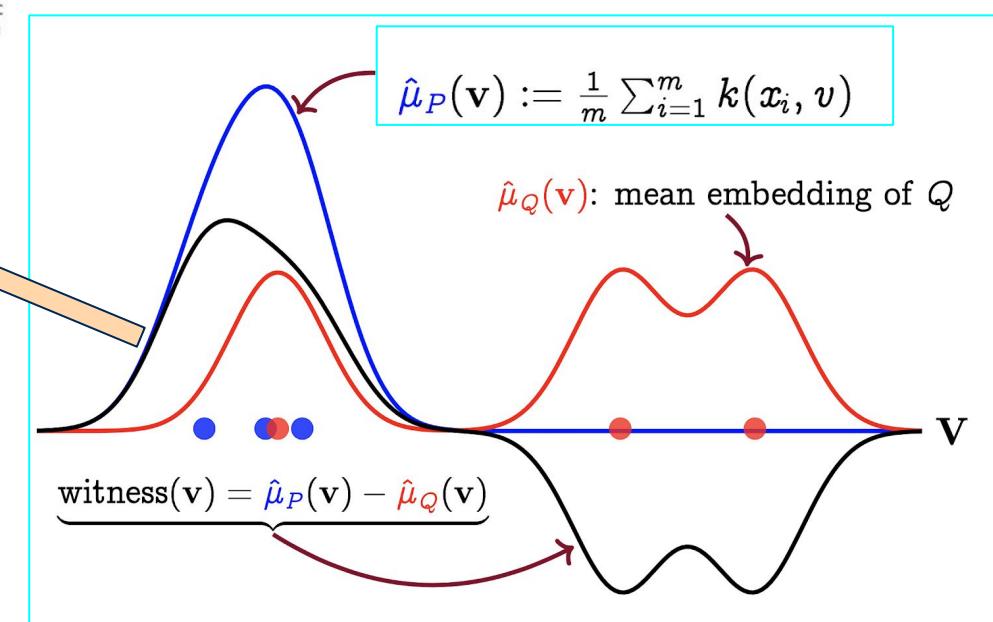
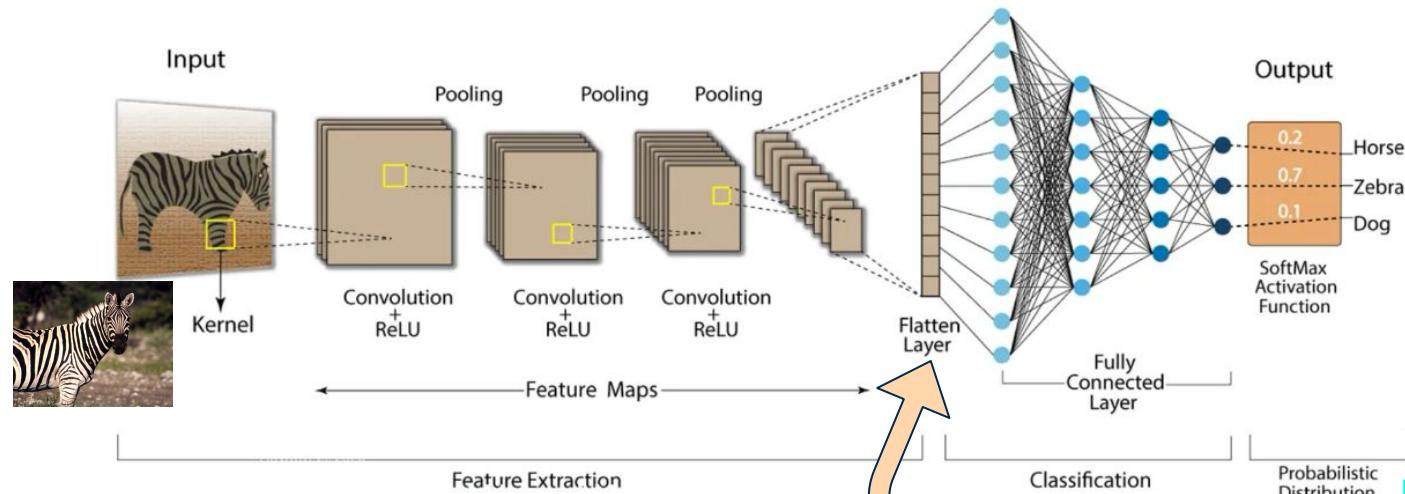
v



Distance-Based Training

Maximum Mean Discrepancy - MMD

Smola et al. (2007)
Gretton et al. (2012)



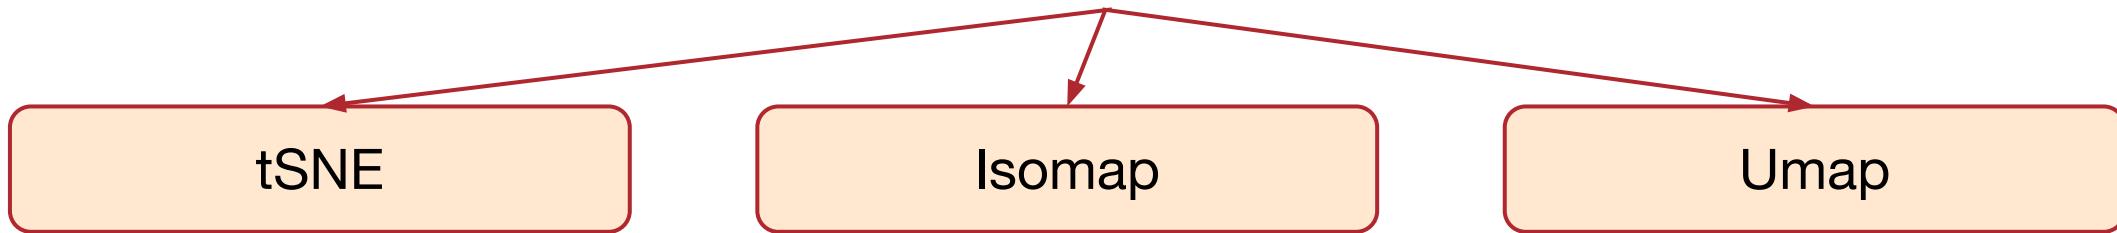
BREAK



How do visualise what DA is doing?

Dimensionality reduction techniques

We need to be able to examine the
multi-dimensional latent (embedding) space
i.e. go from ND to 2D or 3D.





How do visualise what DA is doing?

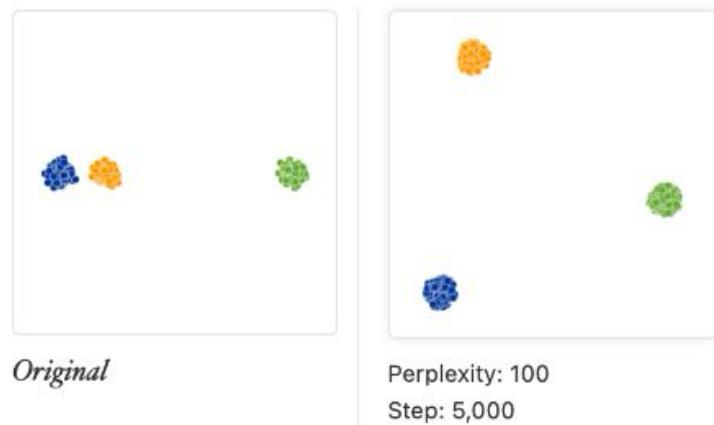
Dimensionality reduction techniques

We need to be able to examine the
multi-dimensional latent (embedding) space
i.e. go from ND to 2D or 3D.

tSNE

Isomap

Umap



t-distributed Stochastic Neighbor Embedding

A nonlinear dimensionality reduction technique that preserves local similarities by modeling high-dimensional data into low-dimensional space using probability distributions. It excels at visualizing clusters but can distort global structure and is sensitive to parameters like perplexity. van der Maaten et al. 2008.



How do visualise what DA is doing?

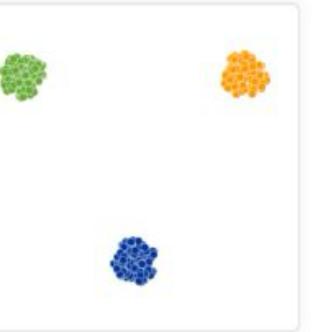
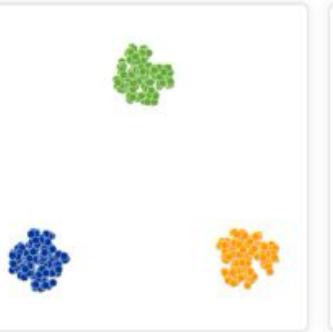
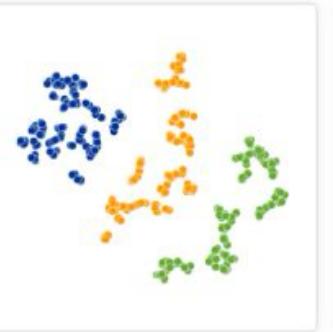
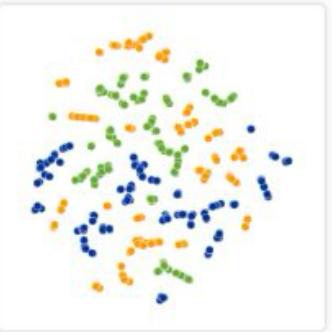
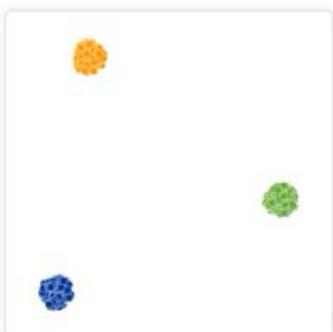
Dimensionality reduction techniques

We need to be able to examine the
multi-dimensional latent (embedding) space
i.e. go from ND to 2D or 3D.

tSNE

Isomap

Umap





How do visualise what DA is doing?

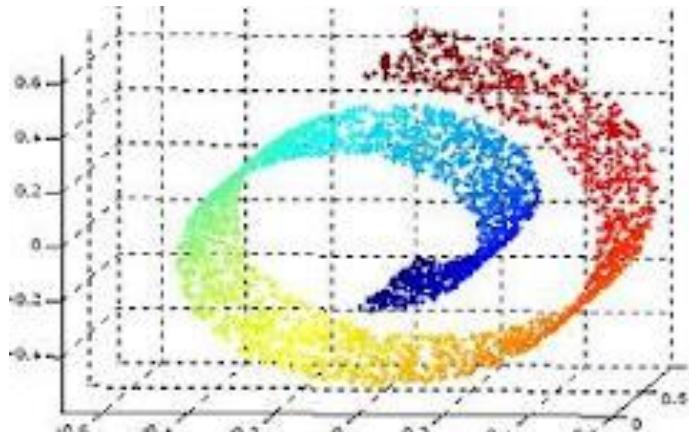
Dimensionality reduction techniques

We need to be able to examine the
multi-dimensional latent (embedding) space
i.e. go from ND to 2D or 3D.

tSNE

Isomap

Umap



Isometric Mapping

Isomap is a manifold learning method that preserves global geometry by computing geodesic distances along a graph connecting nearest neighbors. It is useful for capturing the underlying structure of data that lies on a nonlinear manifold. Tennenbaum et al. 2000.



How do visualise what DA is doing?

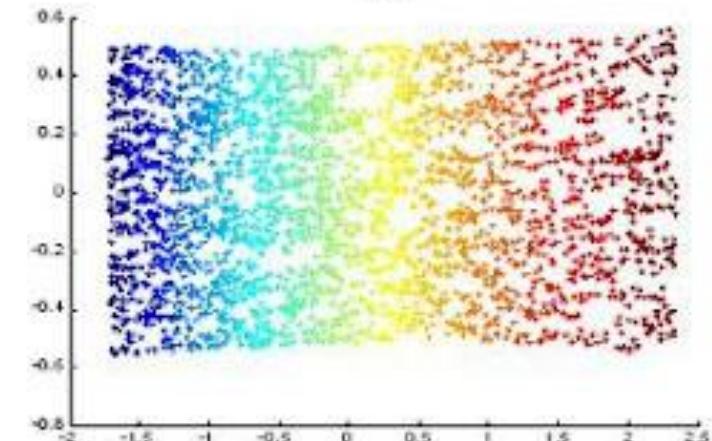
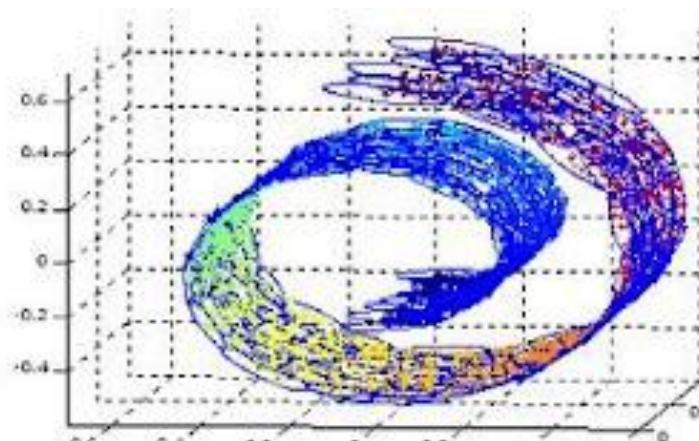
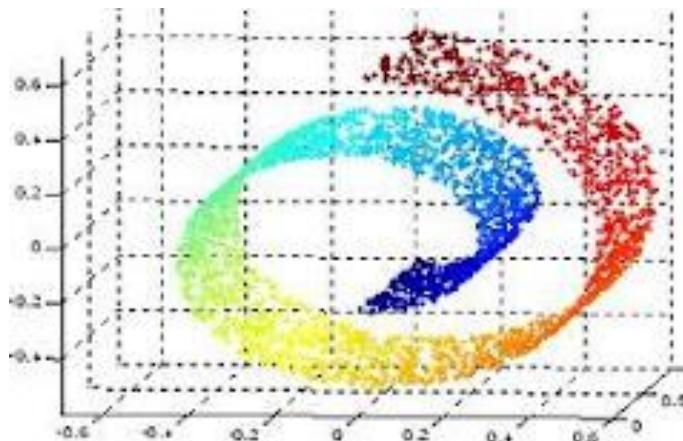
Dimensionality reduction techniques

We need to be able to examine the
multi-dimensional latent (embedding) space
i.e. go from ND to 2D or 3D.

tSNE

Isomap

Umap





How do visualise what DA is doing?

Dimensionality reduction techniques

We need to be able to examine the
multi-dimensional latent (embedding) space
i.e. go from ND to 2D or 3D.

tSNE

Isomap

Umap

~4 million single-cell transcriptomes from adult mouse brain labeled by source brain region represented by a UMAP (Yao Z. et al. 2023, bioRxiv).

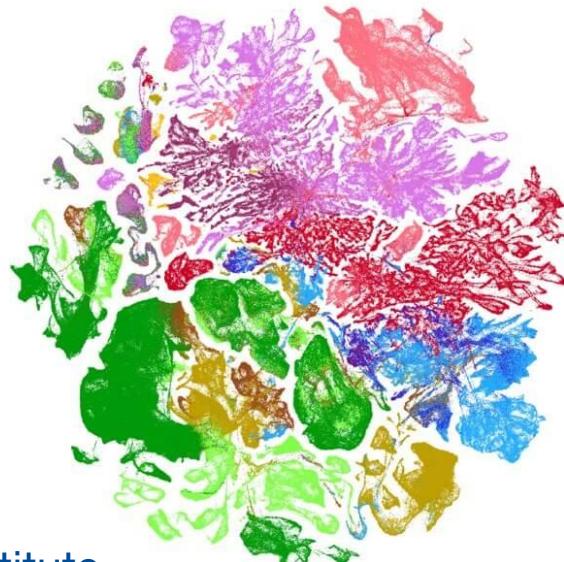


Figure from: [Allen Institute](#)

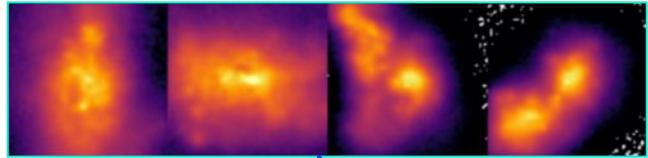
Uniform Manifold Approximation and Projection

UMAP is a fast, scalable nonlinear dimensionality reduction technique that preserves both local and some global structure using fuzzy topological representations. It often outperforms t-SNE in preserving global relationships and is more suitable for large datasets. McInnes et al. 2020.

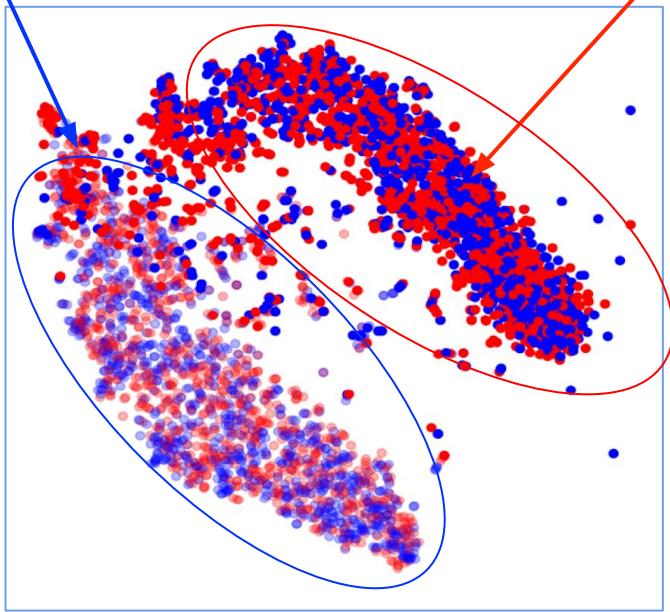
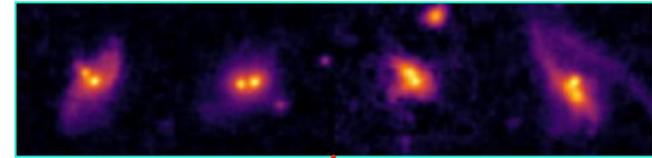


Latent Space Alignment

Source - Illustris



Target - SDSS observations

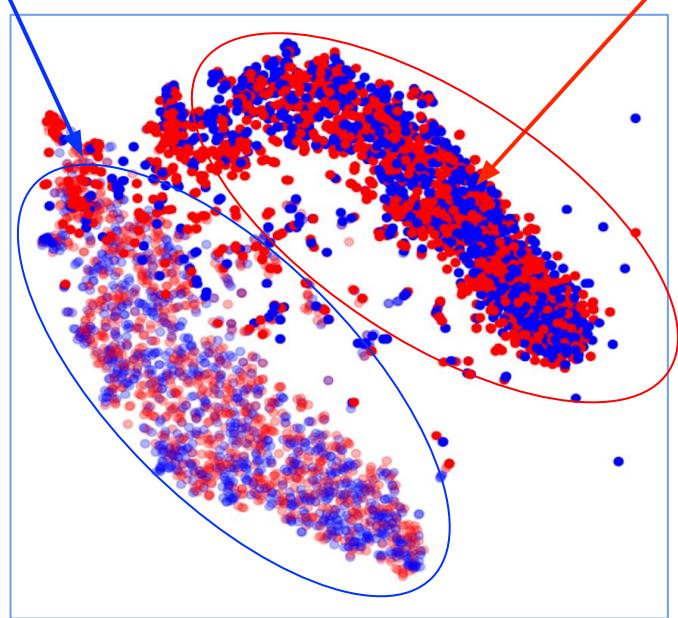
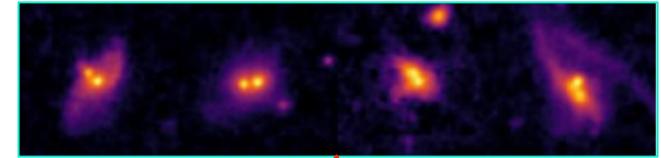
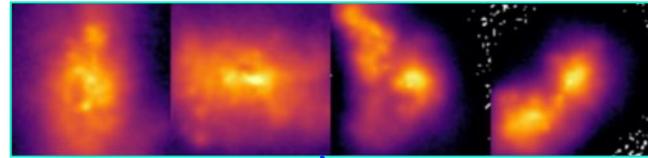


This is how the network sees the data.
2D representation of network's latent space.

Ćiprijanović et al. 2020.
Ćiprijanović et al. 2021.

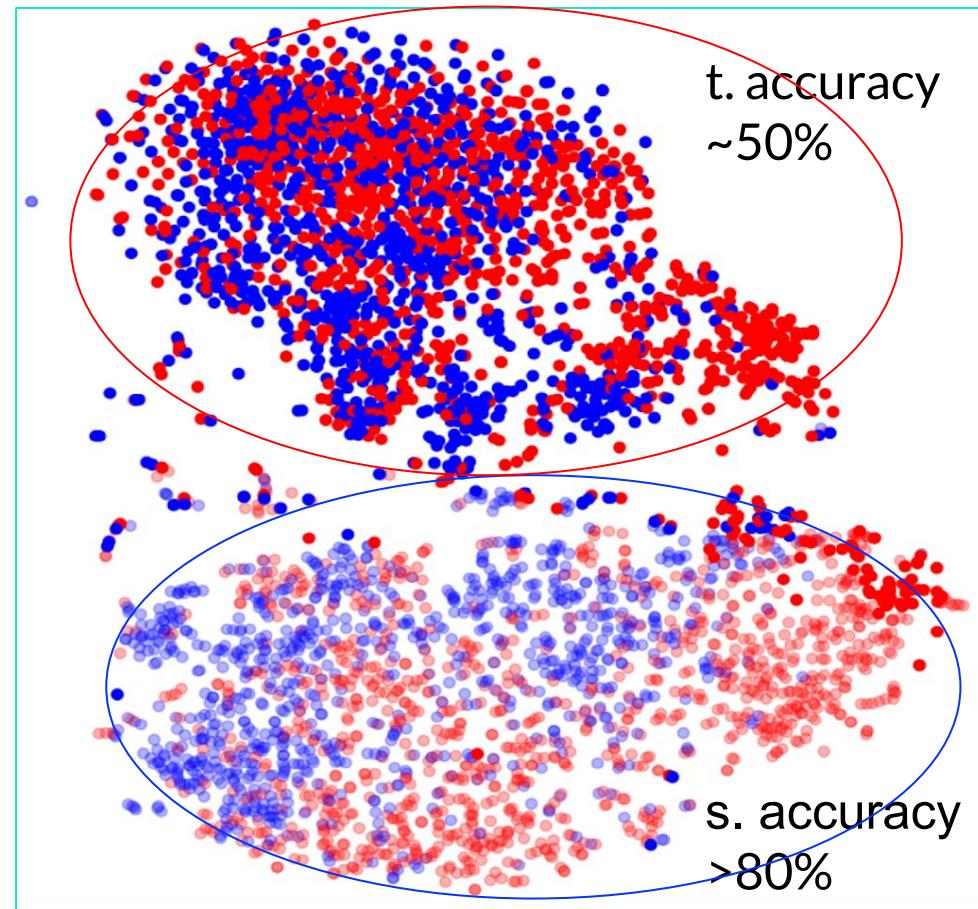
Latent Space Alignment

Source - Illustris Target - SDSS observations



Ćiprijanović et al. 2020.
Ćiprijanović et al. 2021.

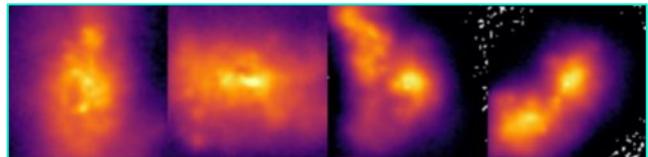
Regular Training



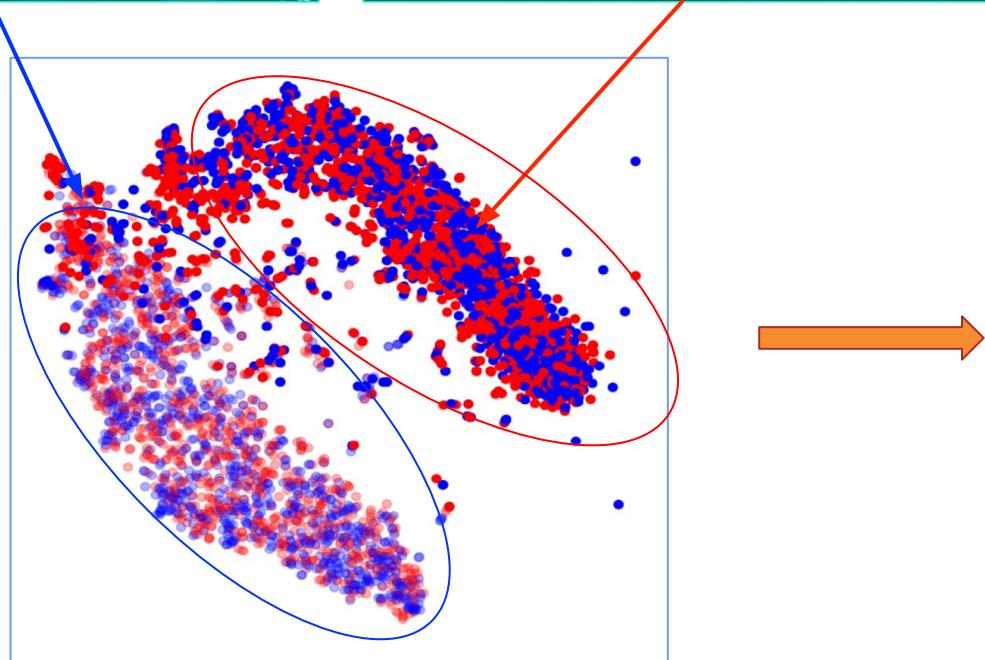
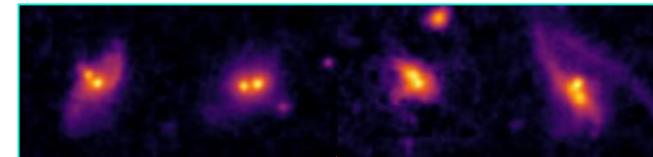


Latent Space Alignment

Source - Illustris



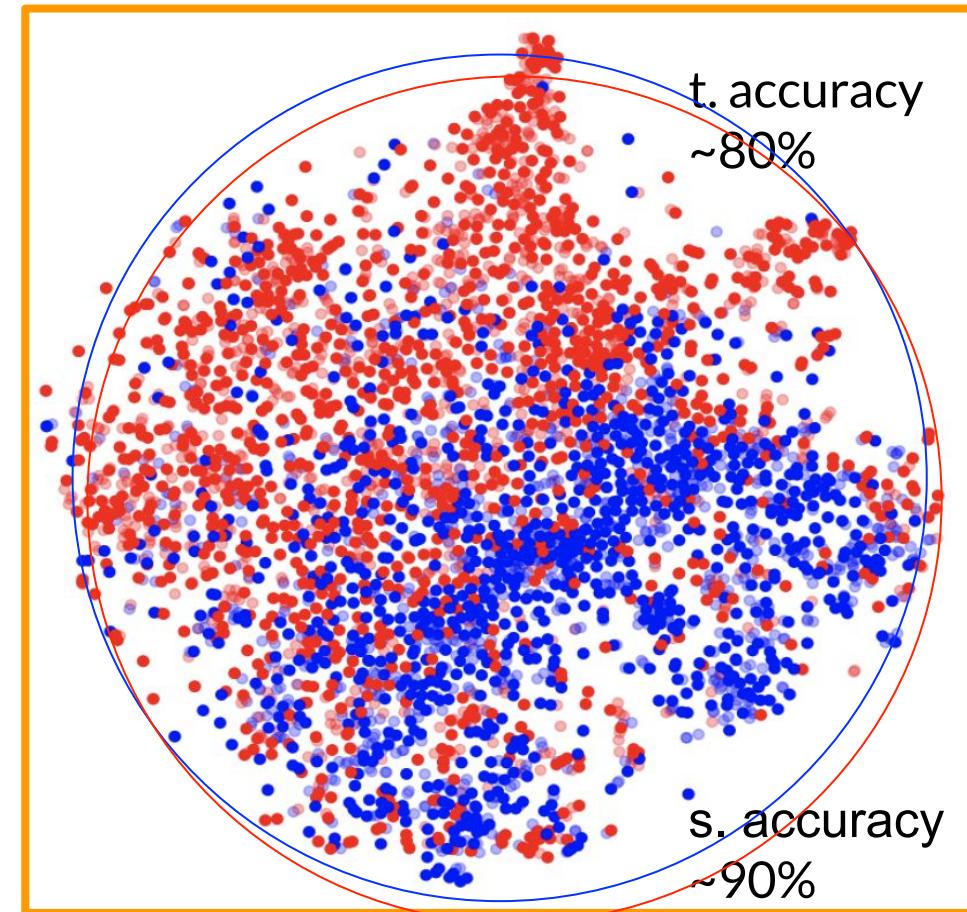
Target - SDSS observations



Ćiprijanović et al. 2020.
Ćiprijanović et al. 2021.

Domain Adaptation

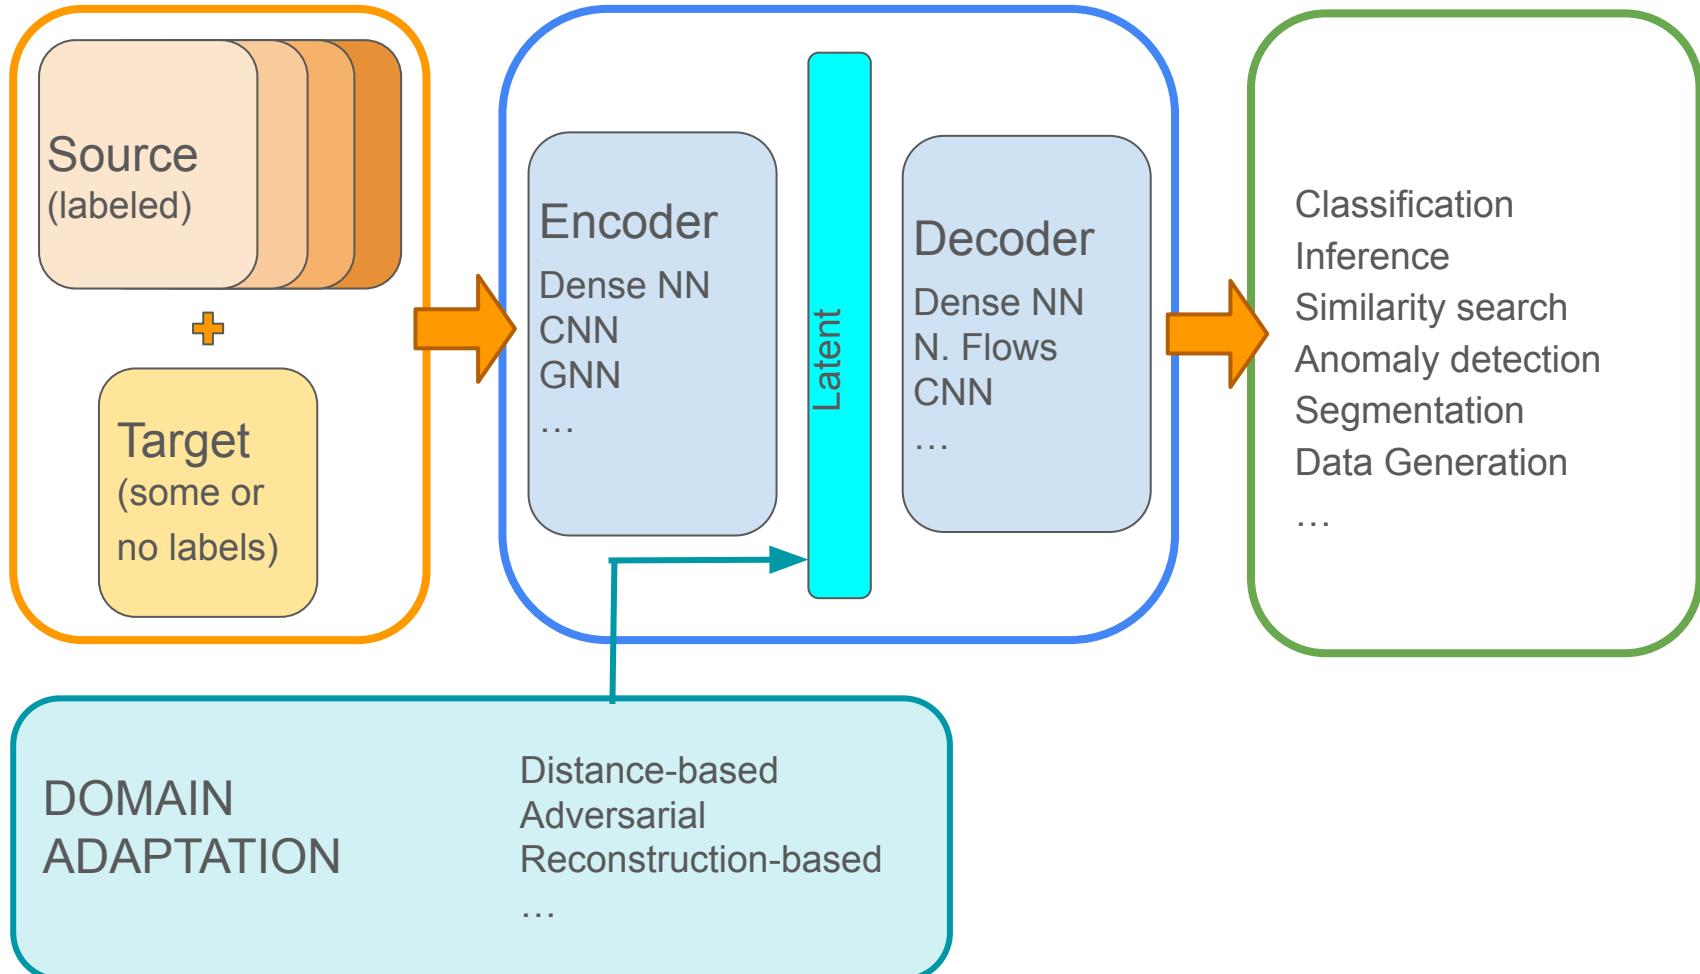
Up to 30% increase!





Where can we use it?

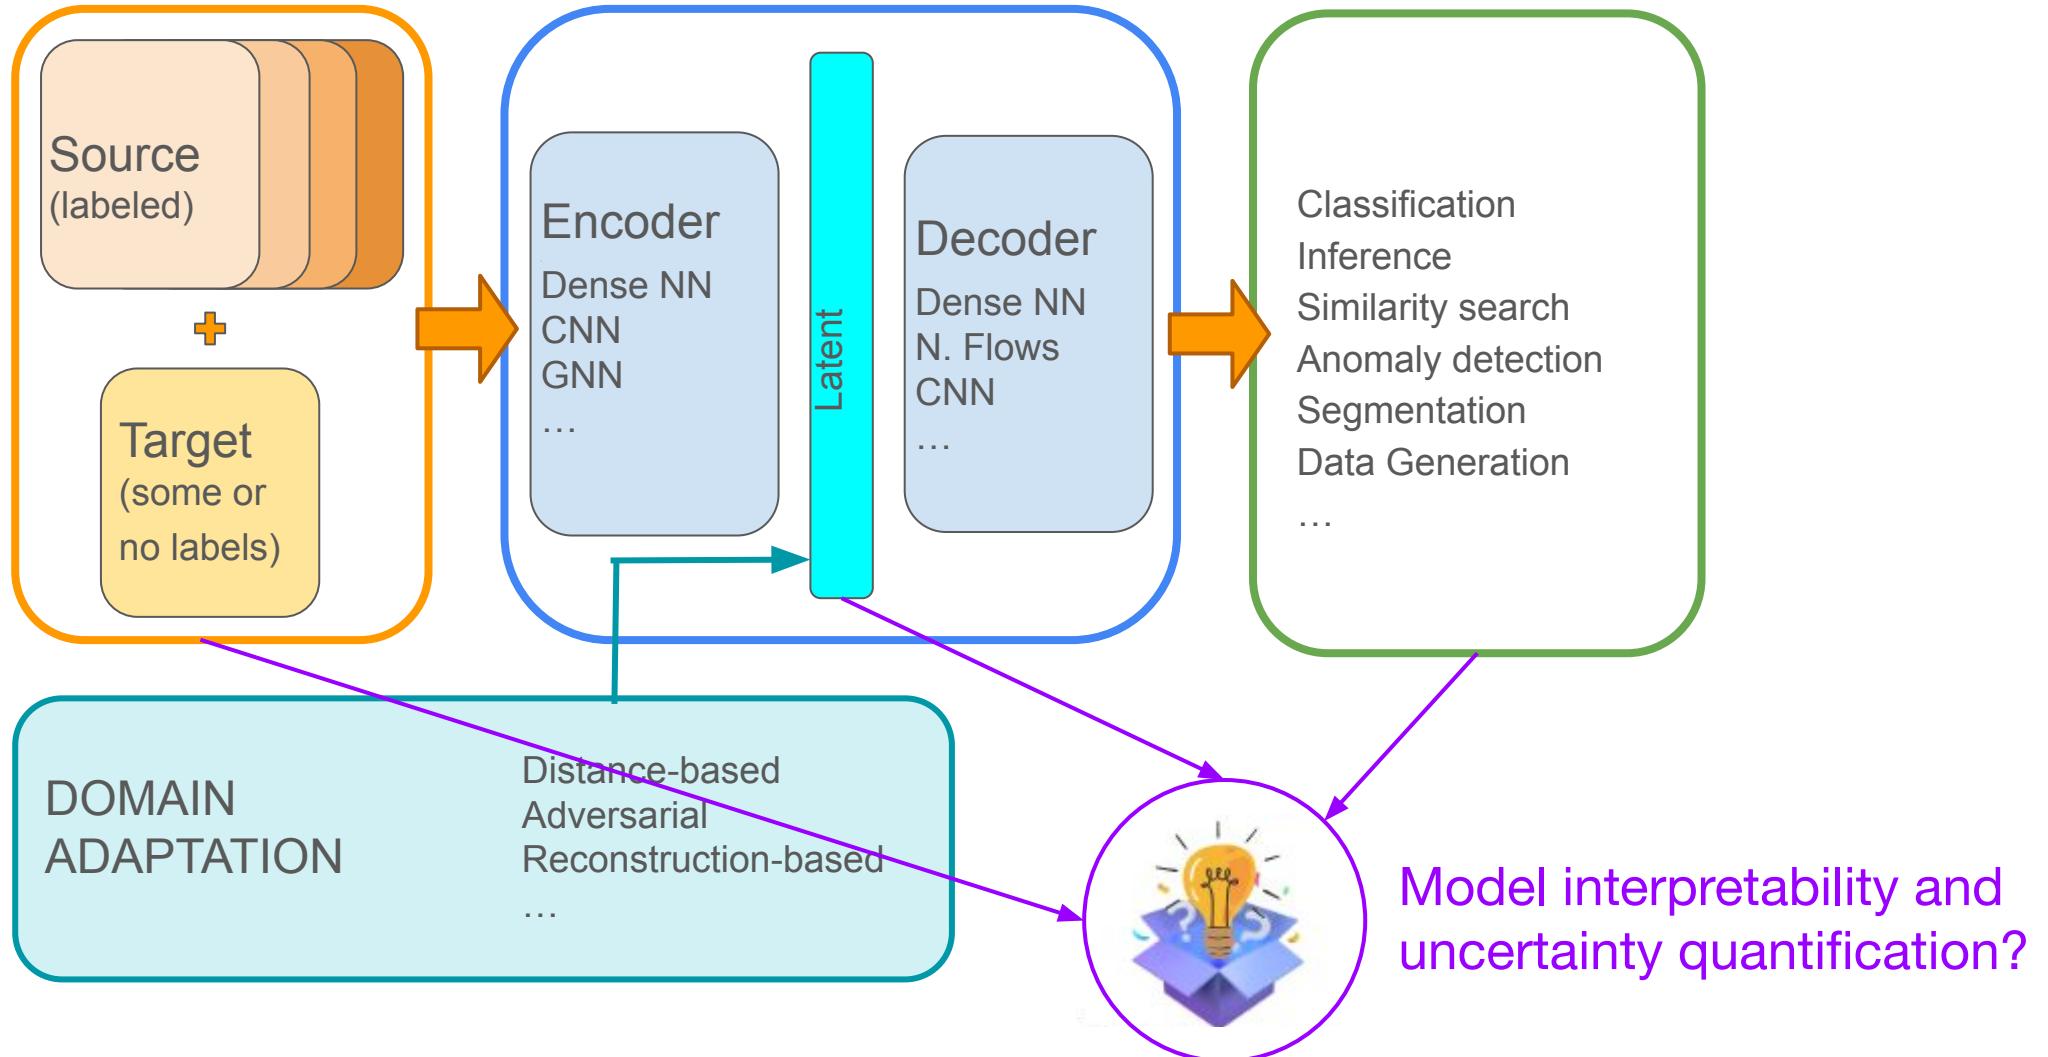
Any data, any task, any problem (within reason!)





Where can we use it?

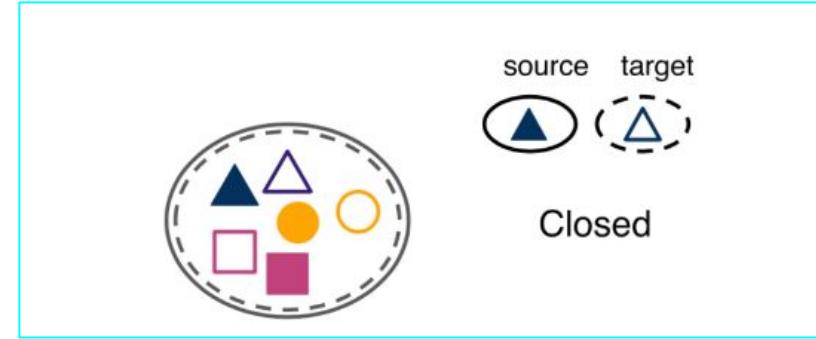
Any data, any task, any problem (within reason!)





Limitations and things to be careful about

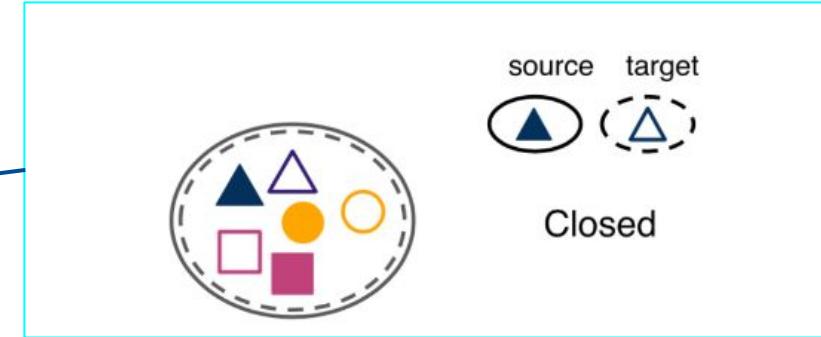
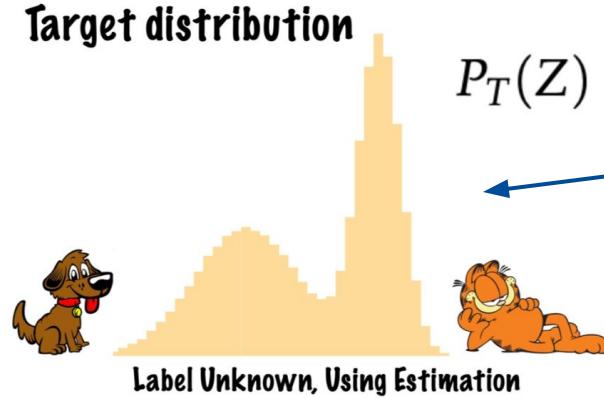
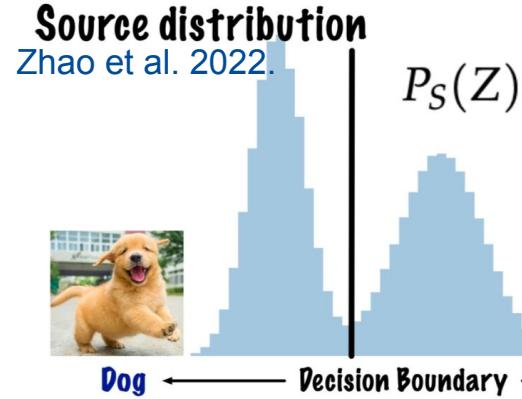
Types of domain overlap





Limitations and things to be careful about

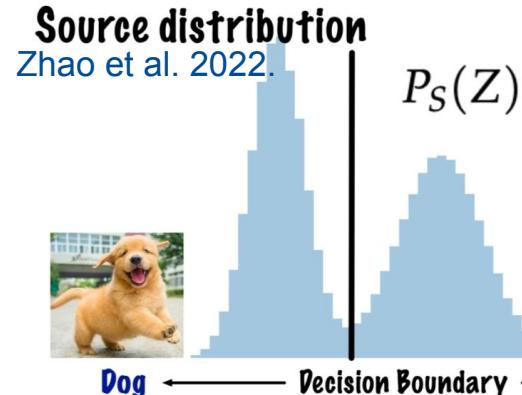
Types of domain overlap



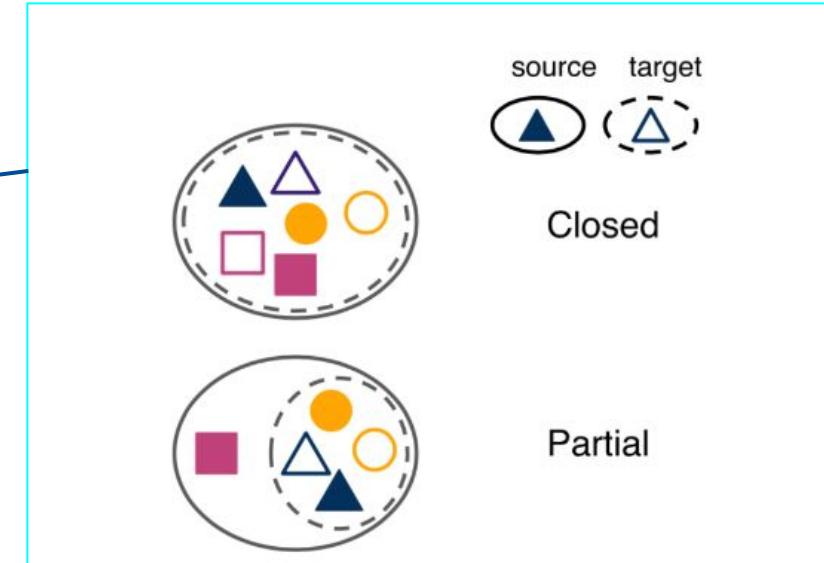
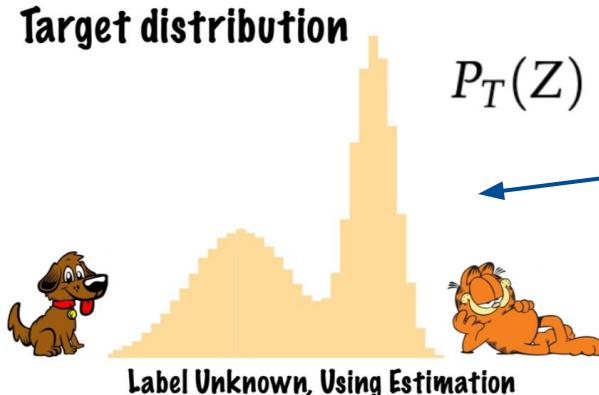


Limitations and things to be careful about

Types of domain overlap



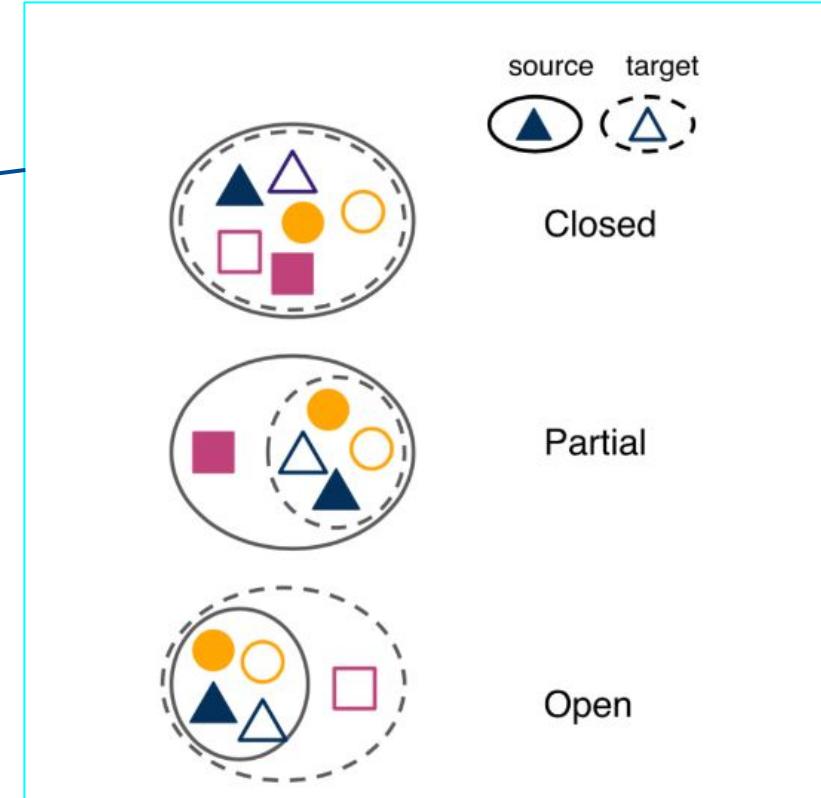
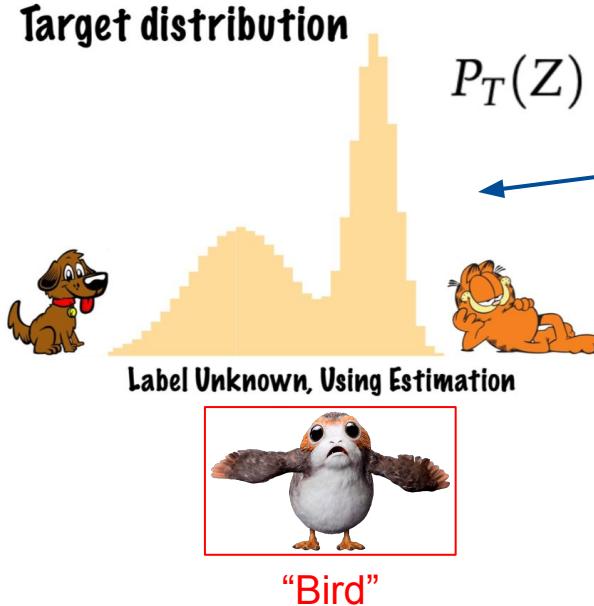
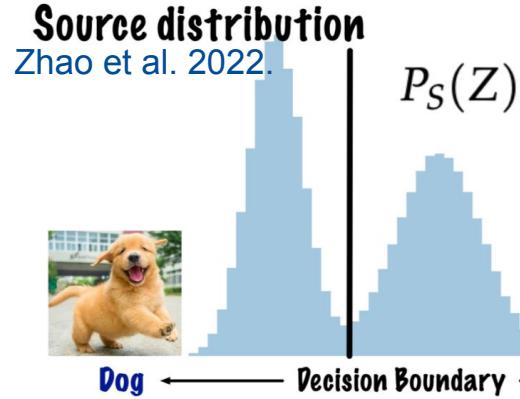
“Bird”





Limitations and things to be careful about

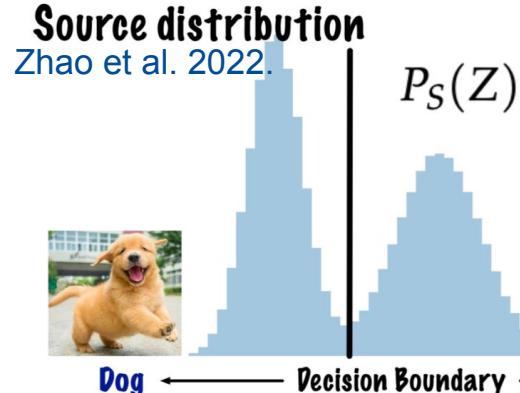
Types of domain overlap



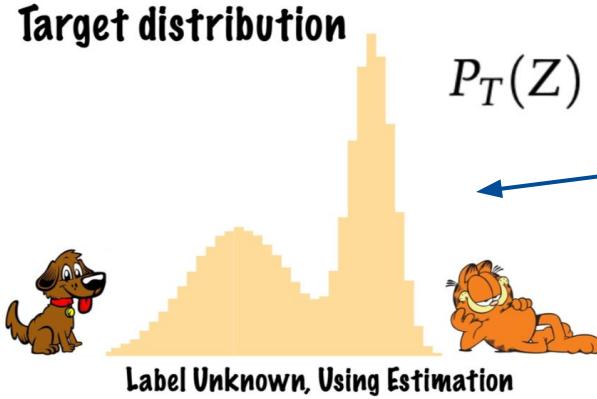


Limitations and things to be careful about

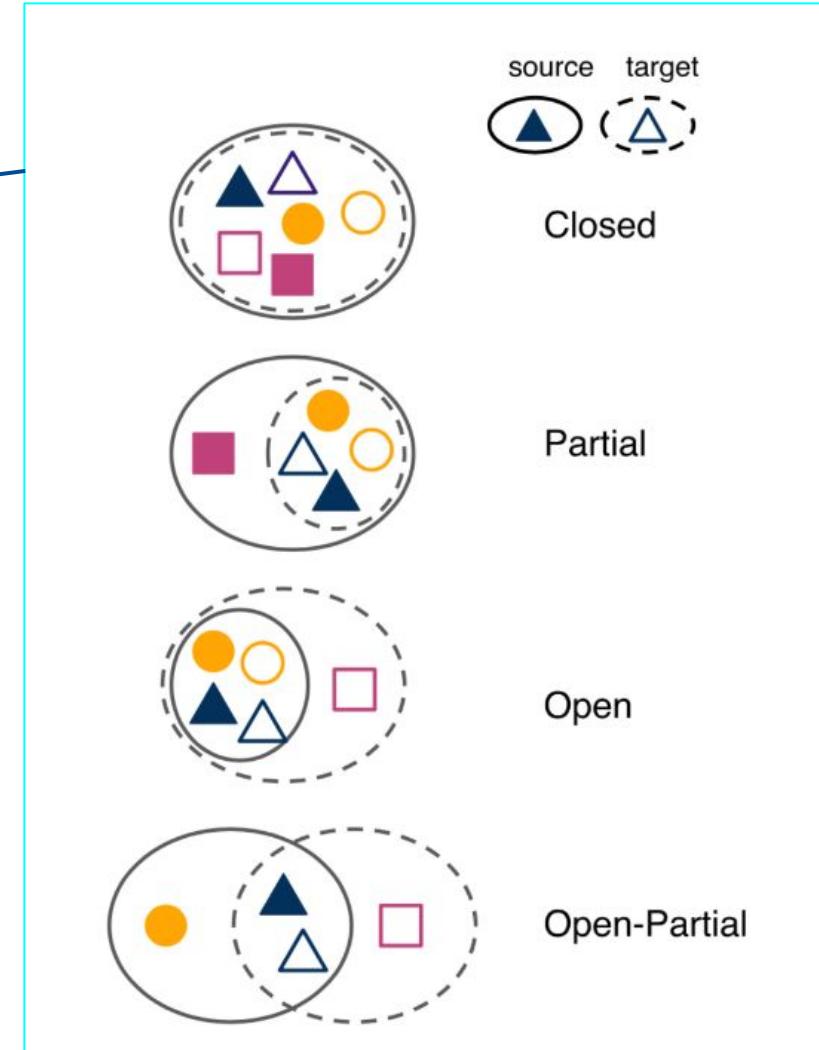
Types of domain overlap



“Bear”



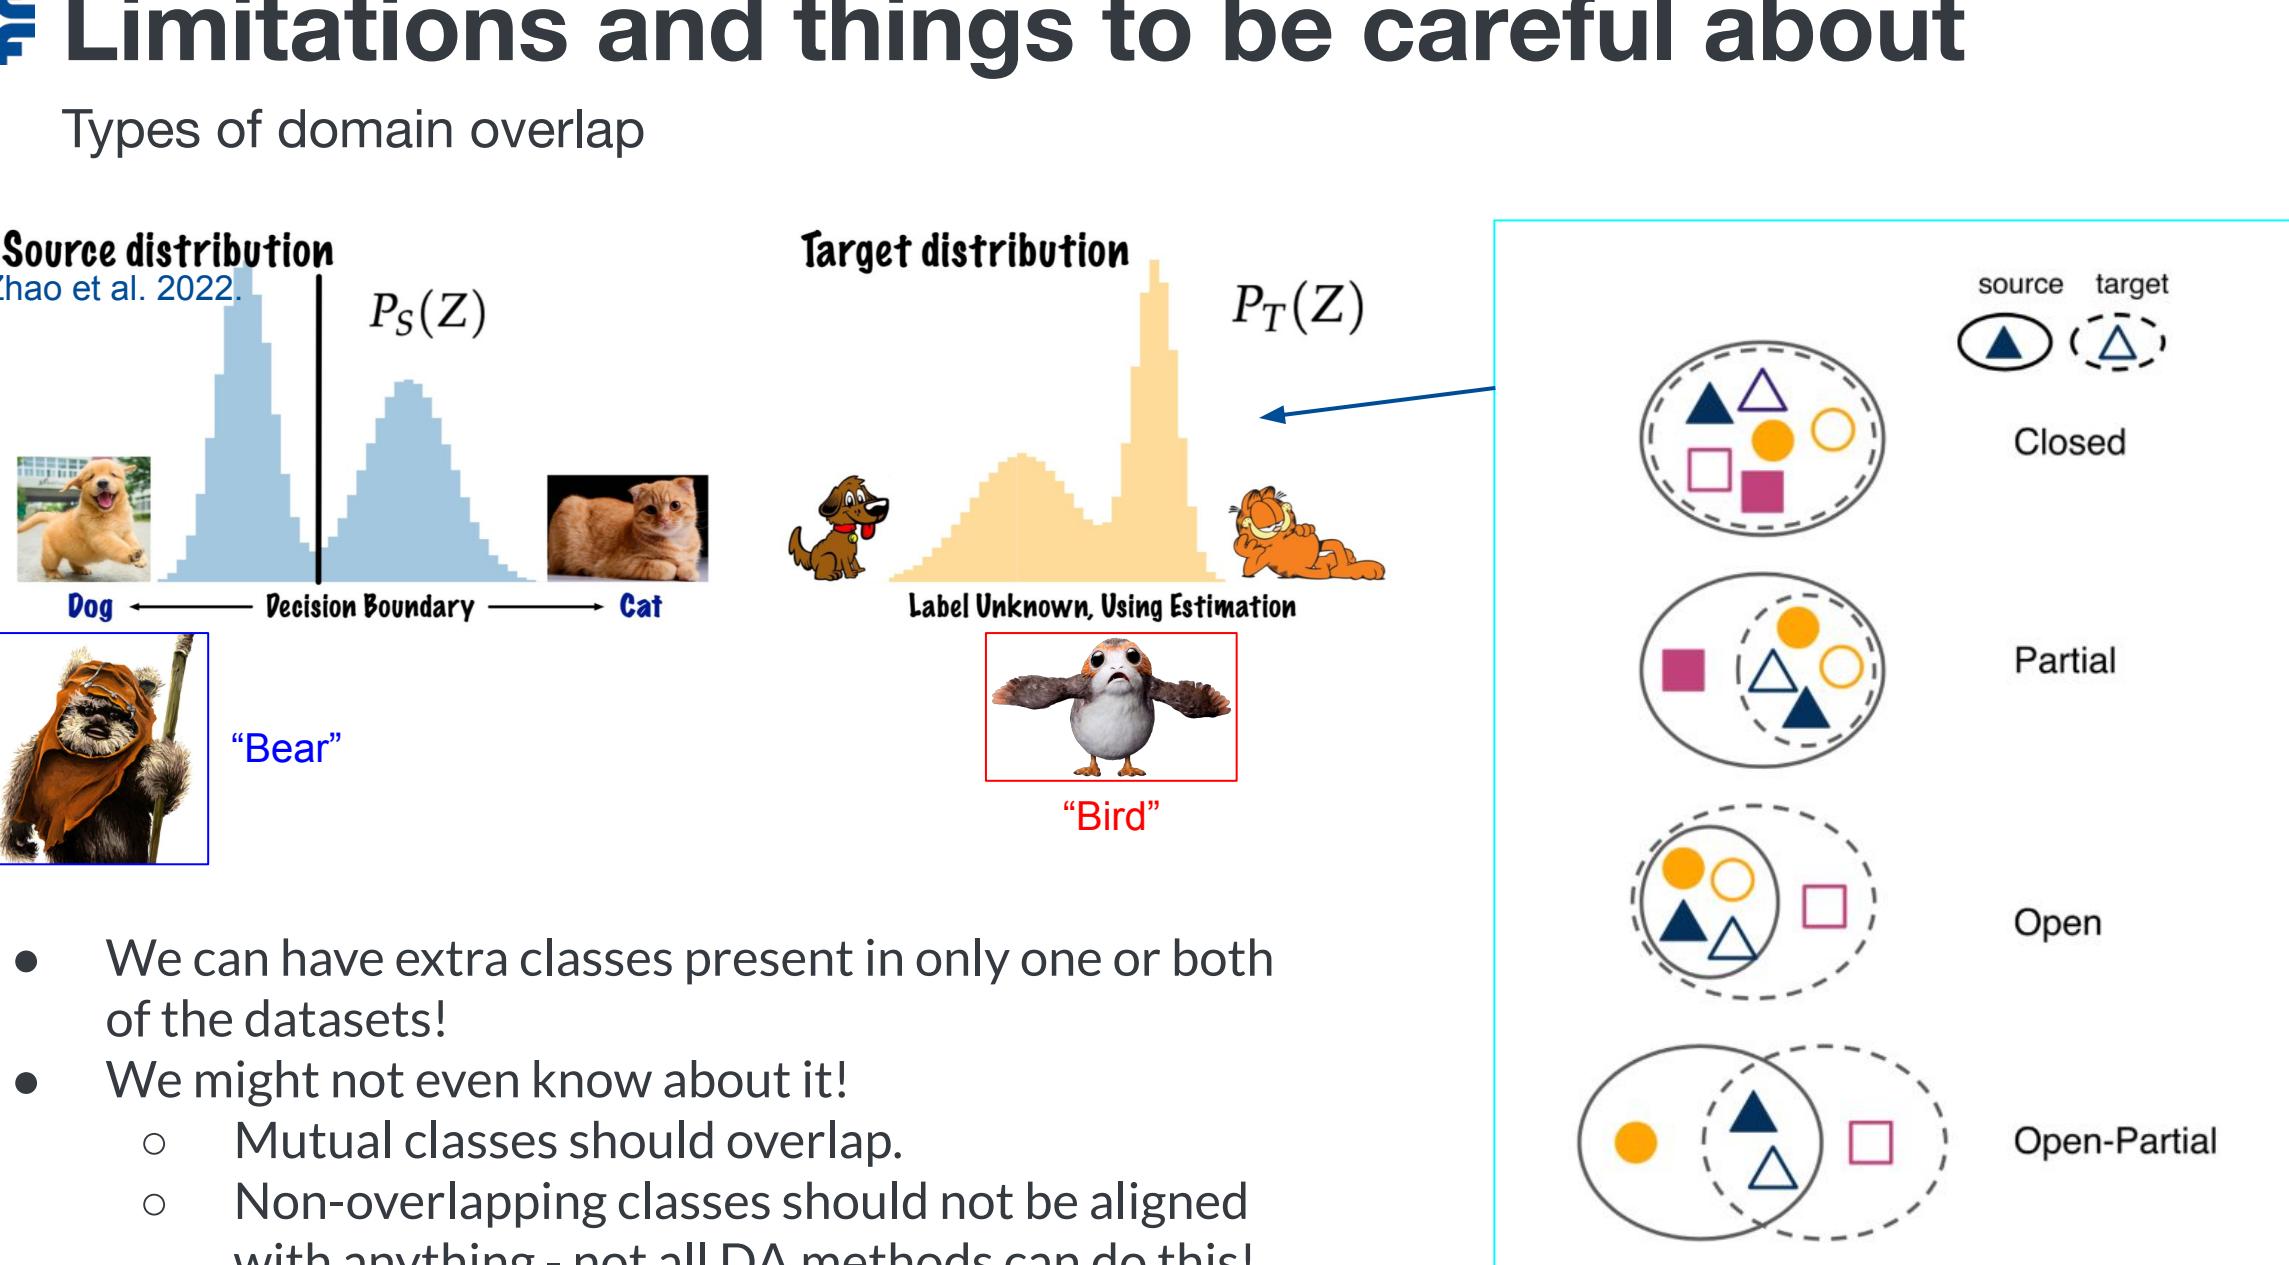
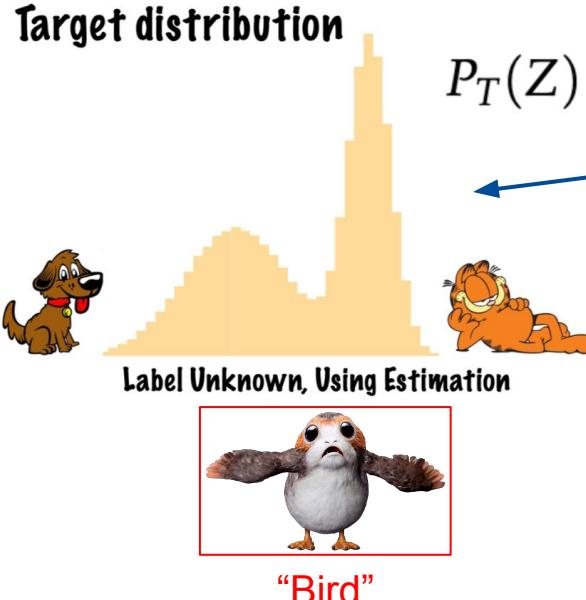
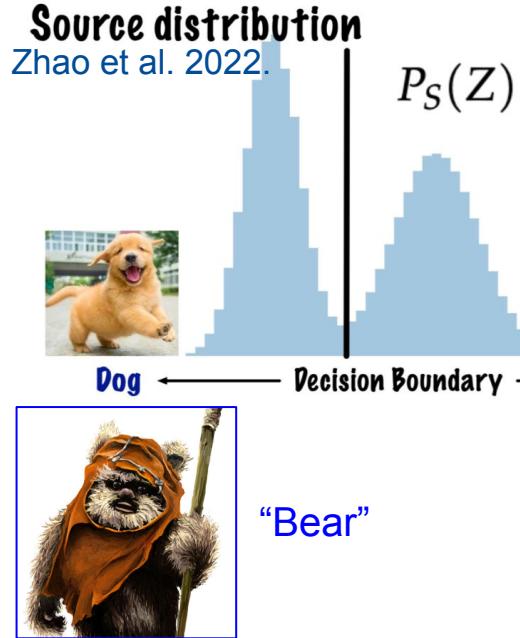
“Bird”





Limitations and things to be careful about

Types of domain overlap

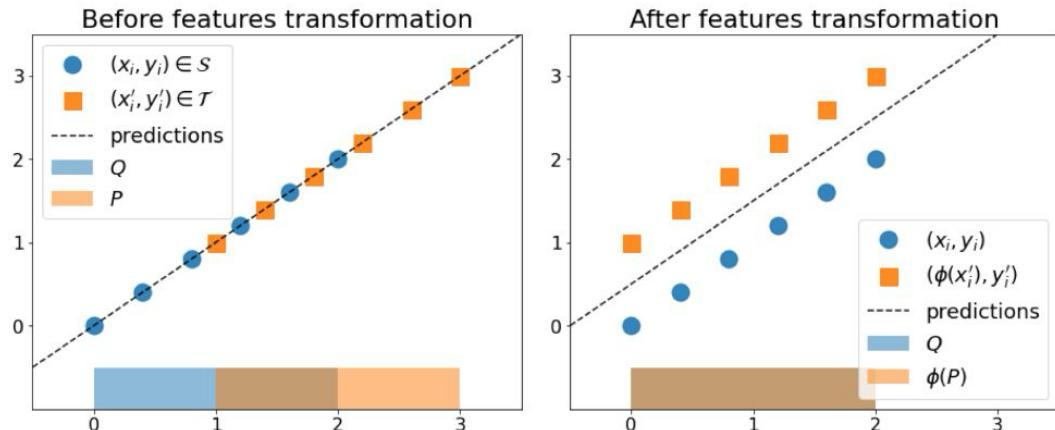


- We can have extra classes present in only one or both of the datasets!
- We might not even know about it!
 - Mutual classes should overlap.
 - Non-overlapping classes should not be aligned with anything - not all DA methods can do this!



Limitations for regression problems

Be careful if prior ranges of parameters being inferred do not match across domains!



de Mathelin et al. 2021.

Fig. 1. Negative transfer caused by feature-based methods. Matching Q and P is detrimental here, as the shift between the two distributions is informative with respect to the task.

Feature-based algorithms (distance or adversarial based) cannot be used if prior ranges don't match.

Only if differences between the domains is related to parameters not being inferred (noise, pixel scale etc).

Instance-based algorithms are necessary
- reweighting importance of labeled instances in the overlapping region.



Domain Adaptation Problems Quiz !

Scenario 1:

You apply a domain adaptation technique to align galaxy images from two telescopes. After adaptation, your model performs **worse** on the source data and loses its ability to distinguish spiral and elliptical galaxies. What likely went wrong?

Learning rate too small

Overfitting to source data

Too aggressive domain alignment



Answer:

problem?



Domain Adaptation Problems Quiz !

Scenario 1:

You apply a domain adaptation technique to align galaxy images from two telescopes. After adaptation, your model performs **worse** on the source data and loses its ability to distinguish spiral and elliptical galaxies. What likely went wrong?

Learning rate too small

Overfitting to source data

Too aggressive domain alignment



Answer:

This is a classic **over-alignment** issue: when features are overly forced to match across domains, the model can discard **discriminative features** that are important for the task. This is especially risky in scientific domains where small differences matter, or domain overlap is small. Careful balance between alignment and task supervision is crucial.

problem?



Domain Adaptation Problems Quiz !

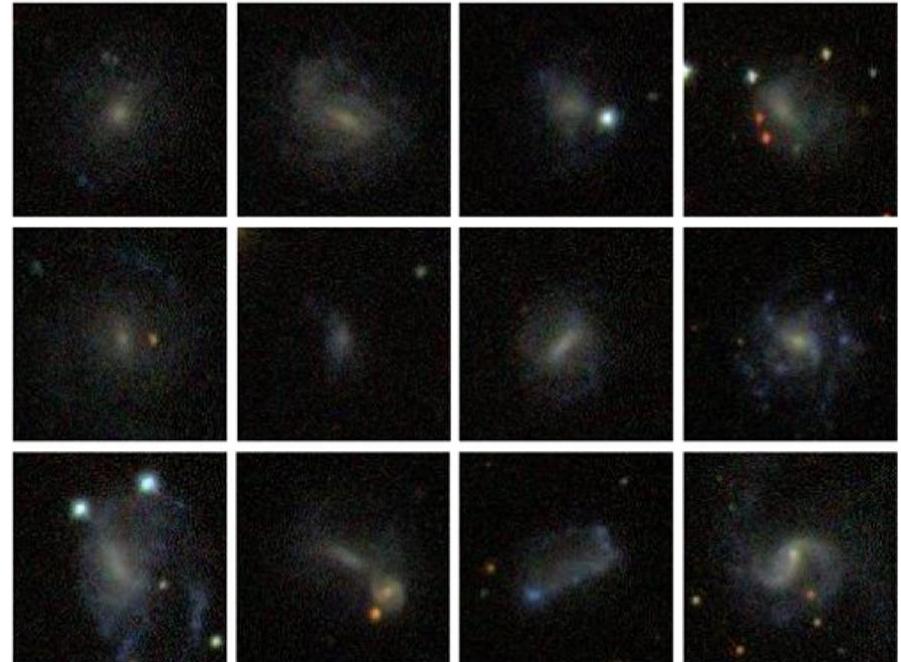
Scenario 2:

You trained a classifier on old survey data where low-surface brightness galaxies are rare, but in the new LSST survey, they dominate. Your adapted model performs poorly on the target domain, despite good feature alignment. What is the most likely cause?

We didn't account for label distribution shift

Feature extractor wasn't pre-trained

Training didn't converge



Answer:

Domain Adaptation Problems Quiz !

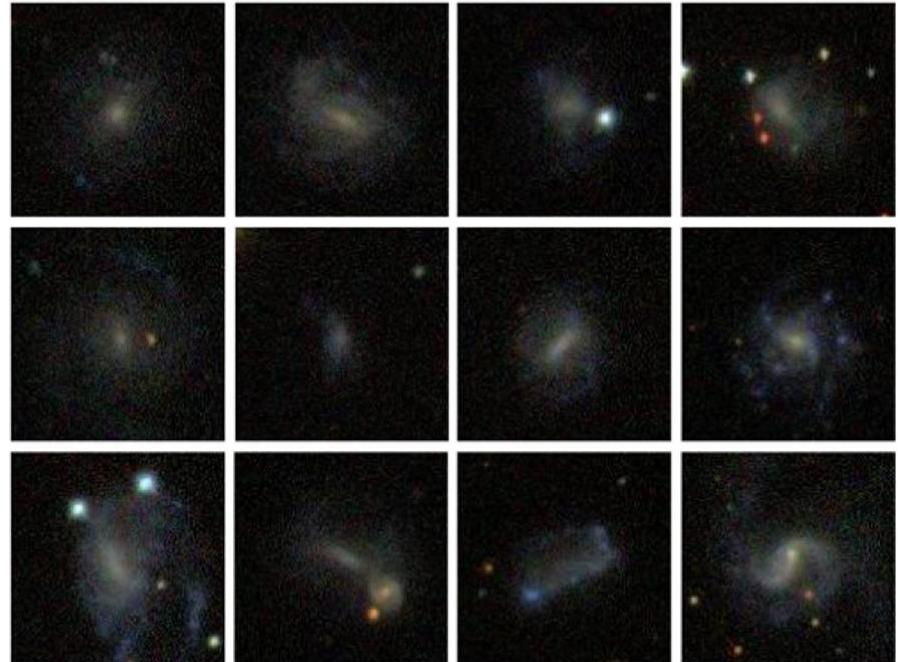
Scenario 2:

You trained a classifier on old survey data where low-surface brightness galaxies are rare, but in the new LSST survey, they dominate. Your adapted model performs poorly on the target domain, despite good feature alignment. What is the most likely cause?

We didn't account for label distribution shift

Feature extractor wasn't pre-trained

Training didn't converge



Answer:

Most domain adaptation methods assume that $P(Y|X)$ stays the same (covariate shift). If instead $P(Y)$ or $P(Y|X)$ changes, models trained on one distribution may misclassify dominant classes in the new domain. Solutions like target-prior correction, self-supervision or instance reweighting are needed.

Domain Adaptation Problems Quiz !

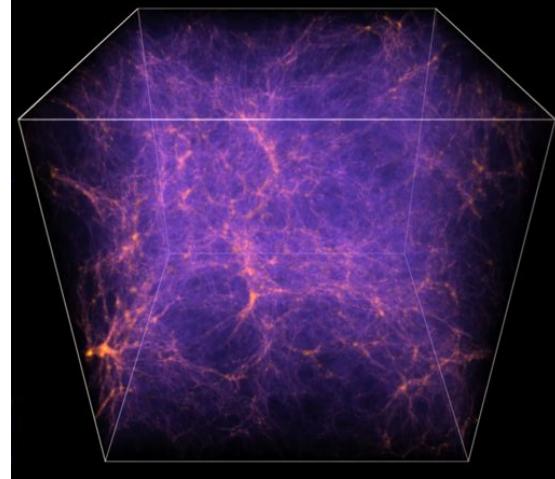
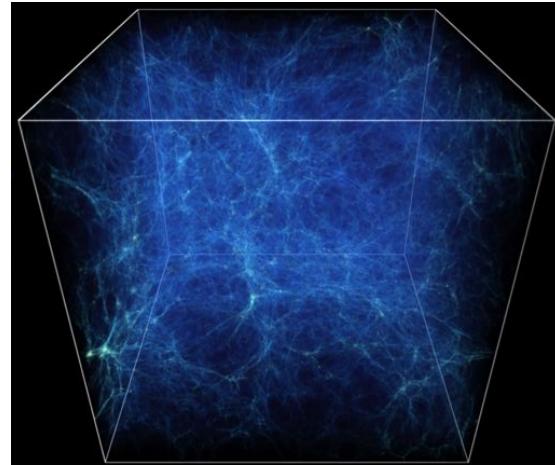
Scenario 3:

You apply a domain adaptation method assuming two datasets come from distinct domains. However, both were collected under nearly identical conditions, just with different random seeds. Model performance drops after adaptation. What likely went wrong?

You overestimated
domain shift

Domain alignment was
overdone

Overfitting



Answer:

Domain Adaptation Problems Quiz !

Scenario 3:

You apply a domain adaptation method assuming two datasets come from distinct domains. However, both were collected under nearly identical conditions, just with different random seeds. Model performance drops after adaptation. What likely went wrong?

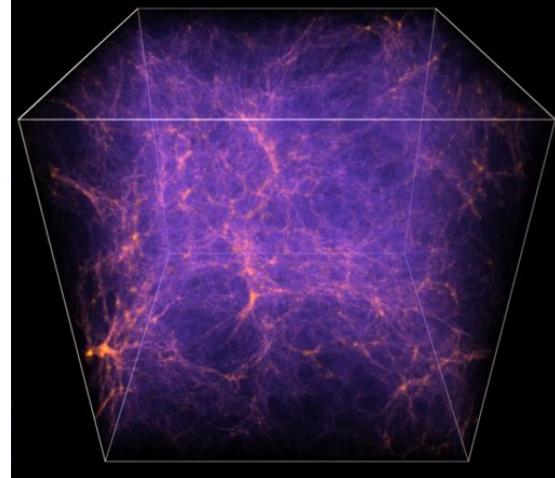
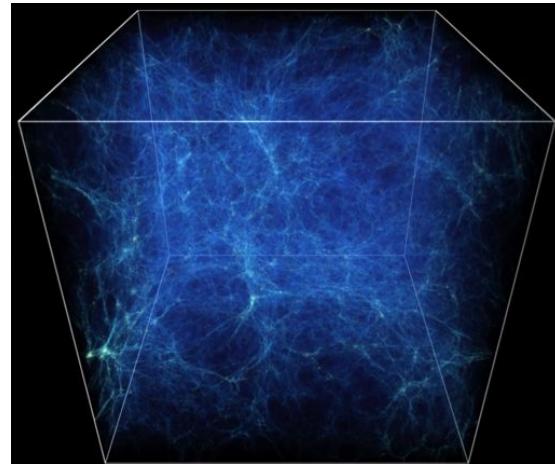
You overestimated domain shift

Domain alignment was overdone

Overfitting

Answer:

Applying domain adaptation when **no meaningful domain gap exists** can **inject noise** or unnecessary transformations into the model. **False domain shift** can degrade performance because of **spurious alignment pressure**. You end up learning representations to solve a non-existent problem, which may reduce task performance and hurt generalization.



03

Examples

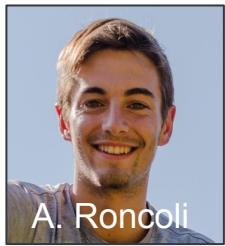
Useful Resources:

- [“SIDDA: Sinkhorn Dynamic Domain Adaptation for Image Classification with Equivariant Neural Networks” Pandya et al. 2025.](#)
- [“Domain-Adaptive Neural Posterior Estimation for Strong Gravitational Lens Analysis” Swierc et al. 2024.](#)
- [“Domain Adaptive Graph Neural Networks for Constraining Cosmological Parameters Across Multiple Data Sets” Roncoli et al. 2023.](#)



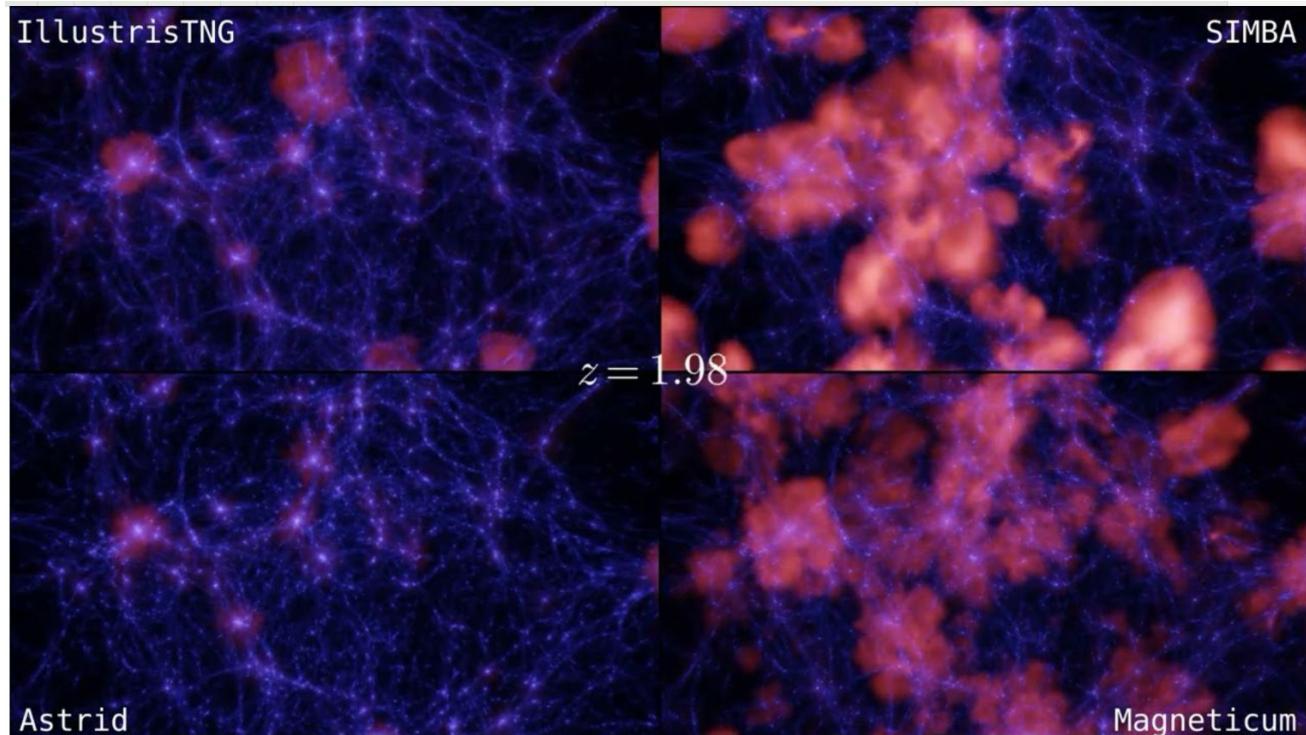
Cosmology with Graphs

NeurIPS 2023.
Roncoli et al. 2023.



A. Roncoli

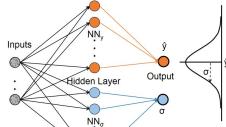
Can we correctly predict cosmology across different cosmological simulations?





Cosmology with Graphs

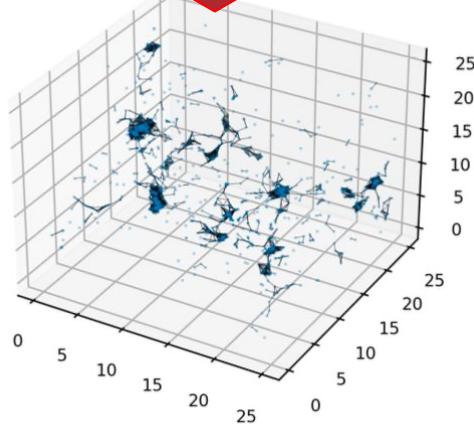
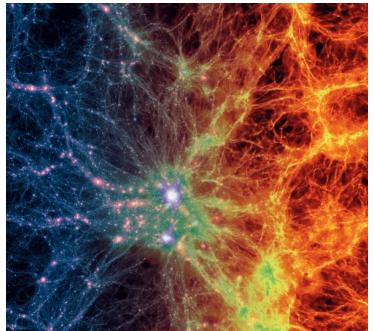
MVE: Infer matter density cosmological parameter



NeurIPS 2023.
Roncoli et al. 2023.



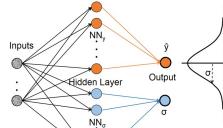
Graph Neural Networks:
ideal for sparse
galaxy catalogs!





Cosmology with Graphs

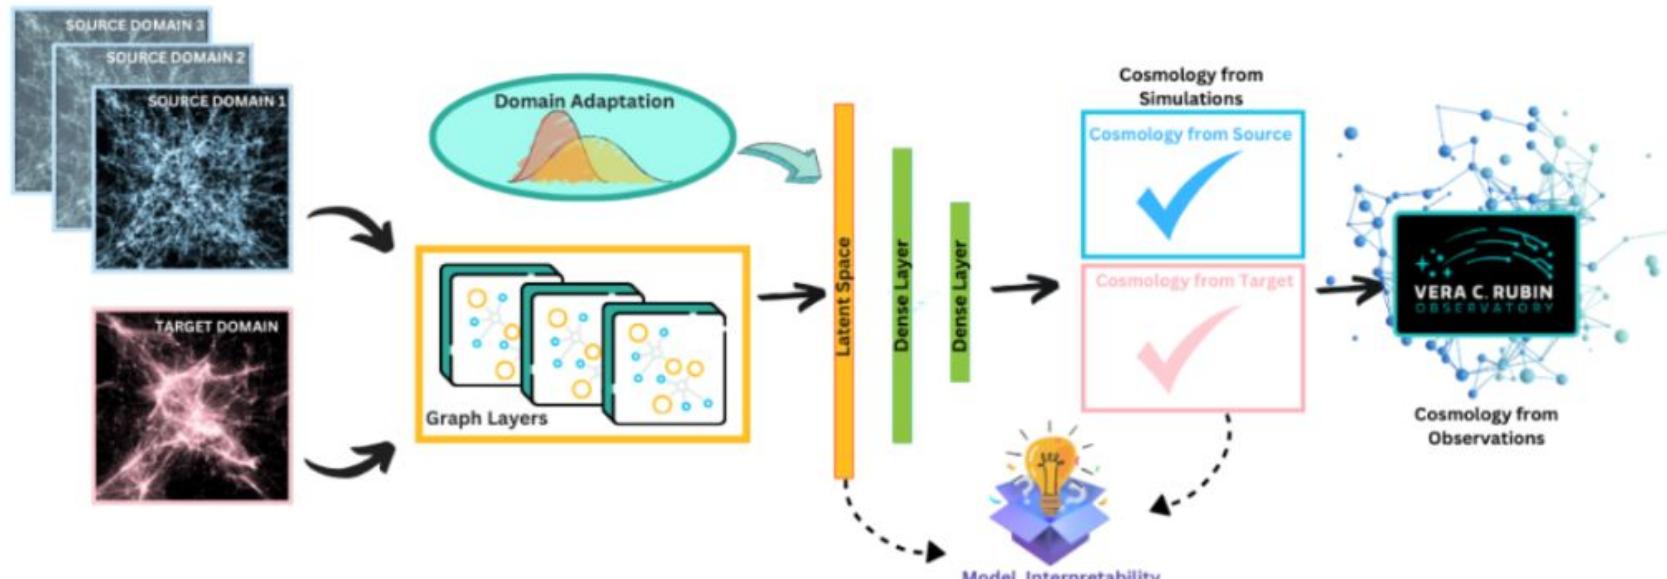
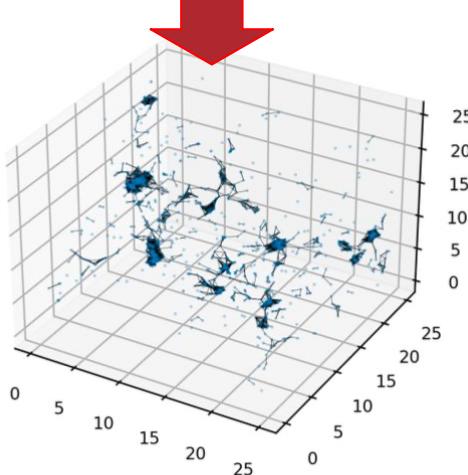
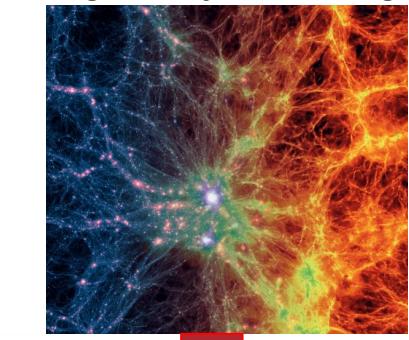
MVE: Infer matter density cosmological parameter



NeurIPS 2023.
Roncoli et al. 2023.



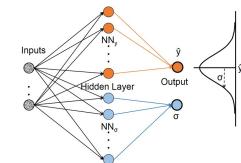
Graph Neural Networks:
ideal for sparse
galaxy catalogs!





Cosmology with Graphs

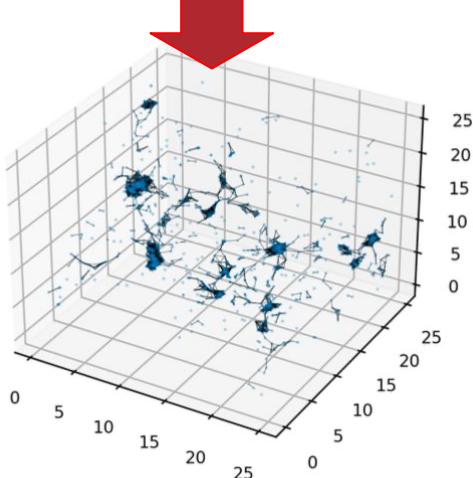
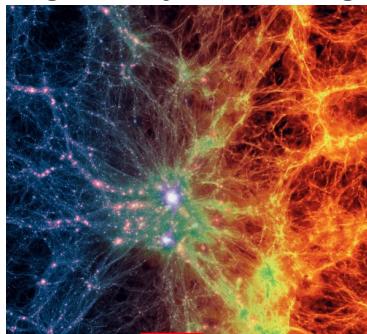
MVE: Infer matter density cosmological parameter



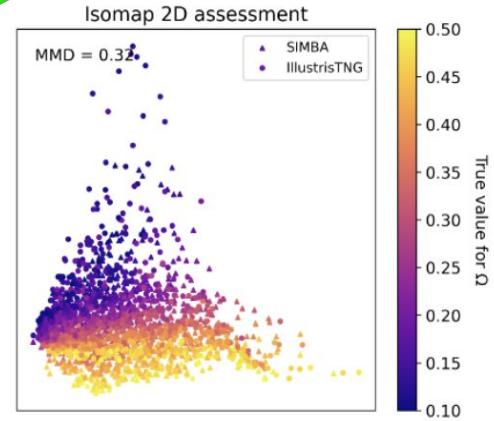
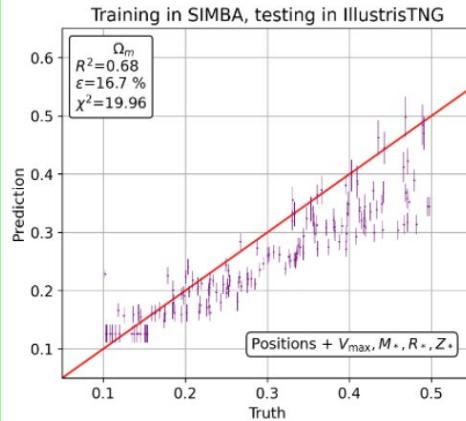
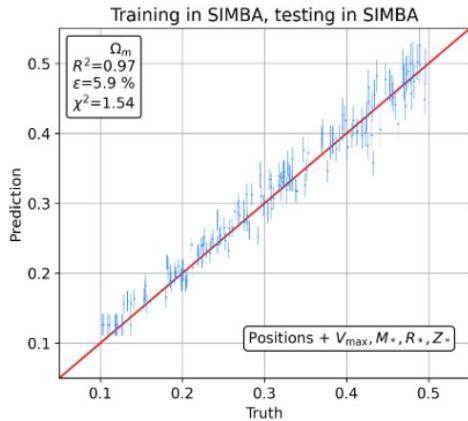
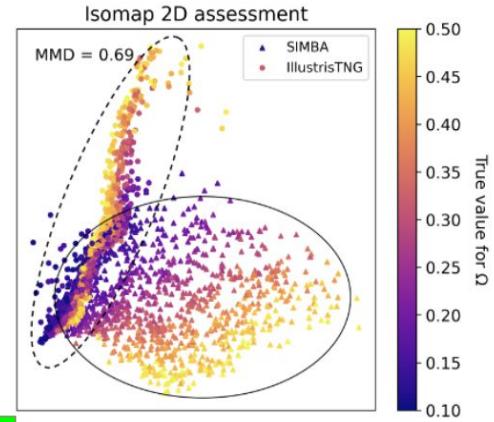
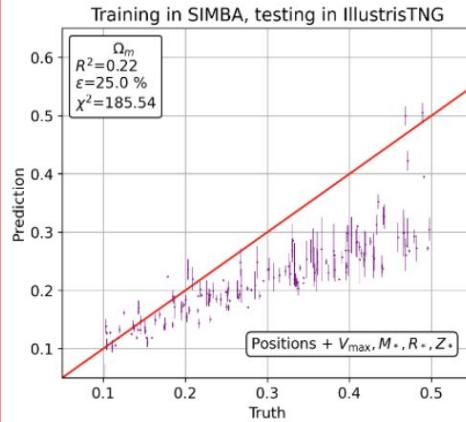
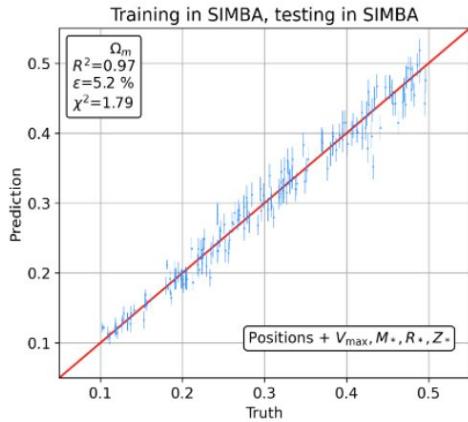
NeurIPS 2023.
Roncoli et al. 2023.



Graph Neural Networks:
ideal for sparse
galaxy catalogs!



$z=0$ 1000 simulations each



28%
better
relative
error

Order
of mag.
better
 χ^2



Strong gravitational lensing



S. Agarwal

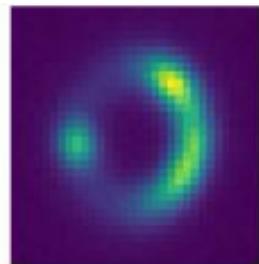


P. Swierc

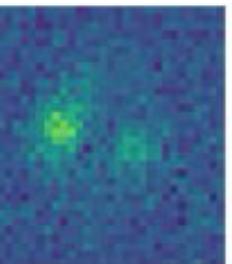
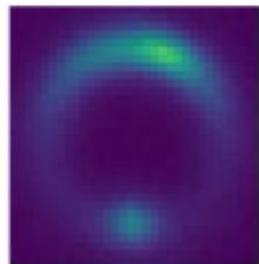


M. Tamargo

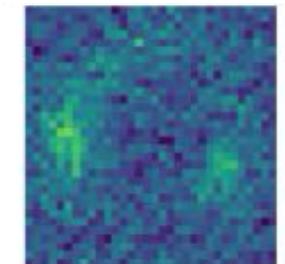
Can we infer strong lensing parameters (Einstein radius, ellipticity, position/offset) robustly in both simulated and real data?



Source - Clean



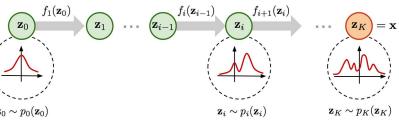
Target - DES



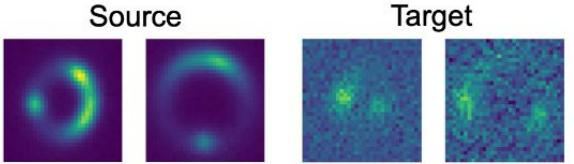


Strong gravitational lensing

SBI - infer Einstein radius, ellipticity, position/offset



(a)



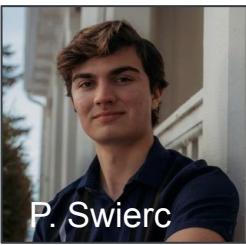
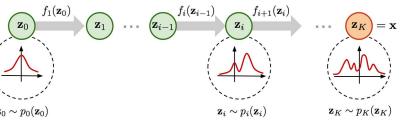
- a. 100k images (1-filter) with and without noise
- b. **Latent distributions are correctly aligned.**
- c. Differences between true values and predicted posteriors are small. **Accuracy improves up to 2 orders of magnitude.**
- d. Posterior are well calibrated and no longer overconfident.

Swierc et al. 2024.

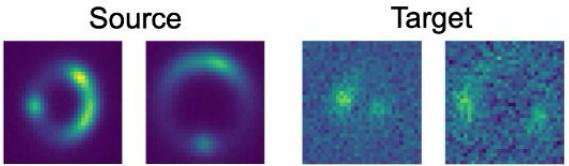


Strong gravitational lensing

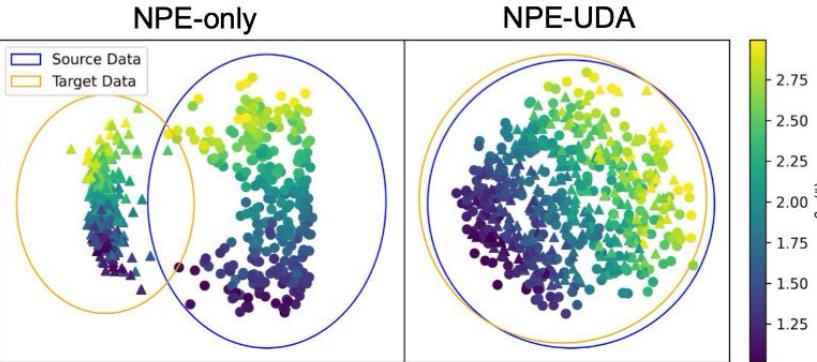
SBI - infer Einstein radius, ellipticity, position/offset



(a)



(b)



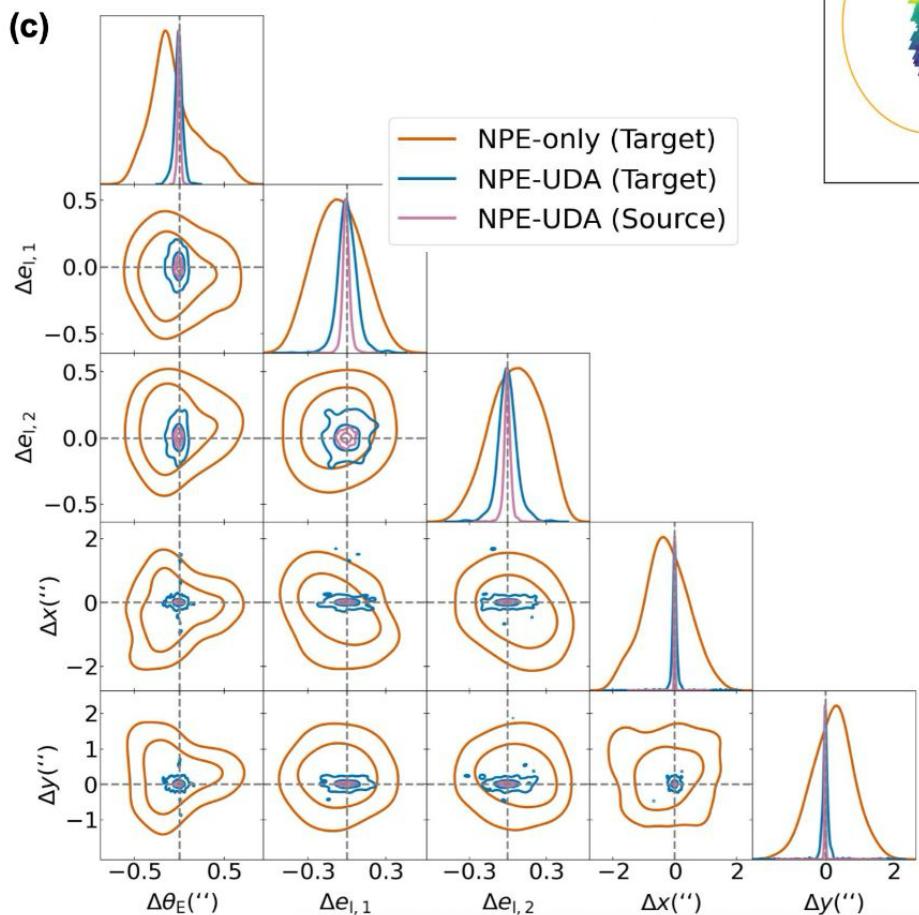
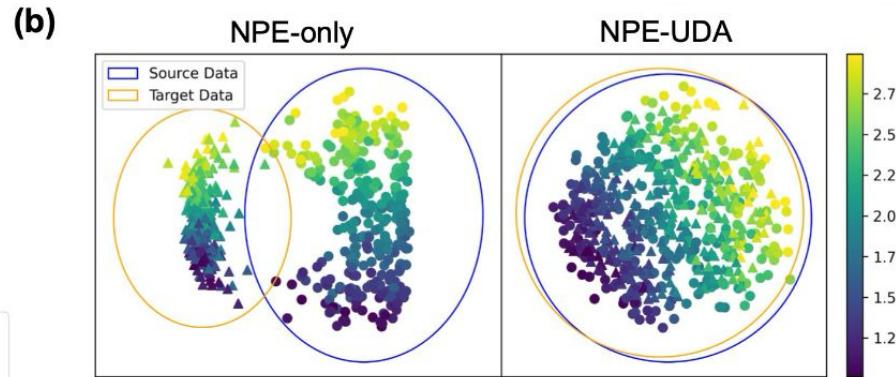
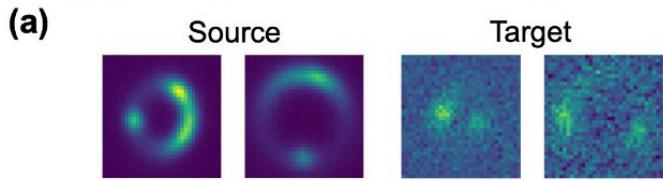
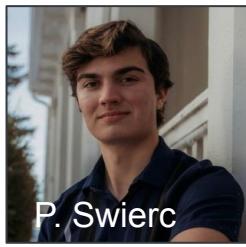
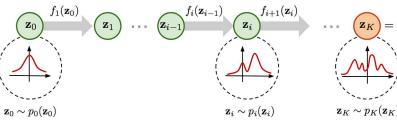
- a. 100k images (1-filter) with and without noise
- b. **Latent distributions are correctly aligned.**
- c. Differences between true values and predicted posteriors are small. **Accuracy improves up to 2 orders of magnitude.**
- d. Posterior are well calibrated and no longer overconfident.

Swierc et al. 2024.



Strong gravitational lensing

SBI - infer Einstein radius, ellipticity, position/offset



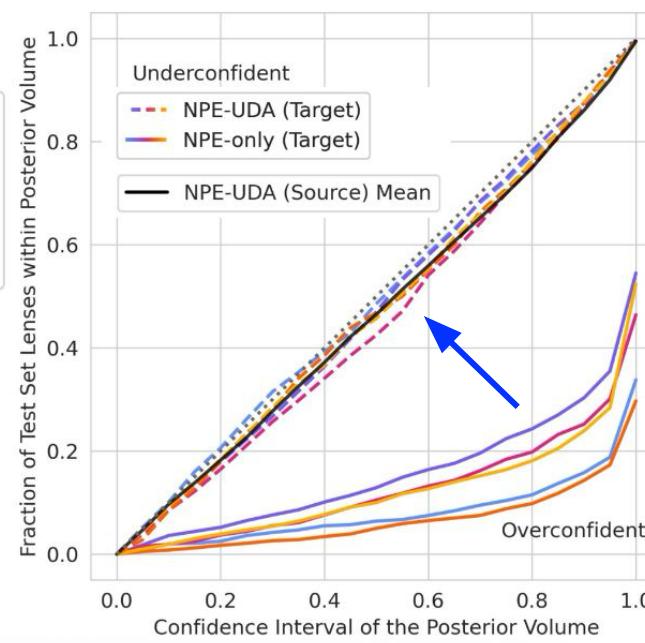
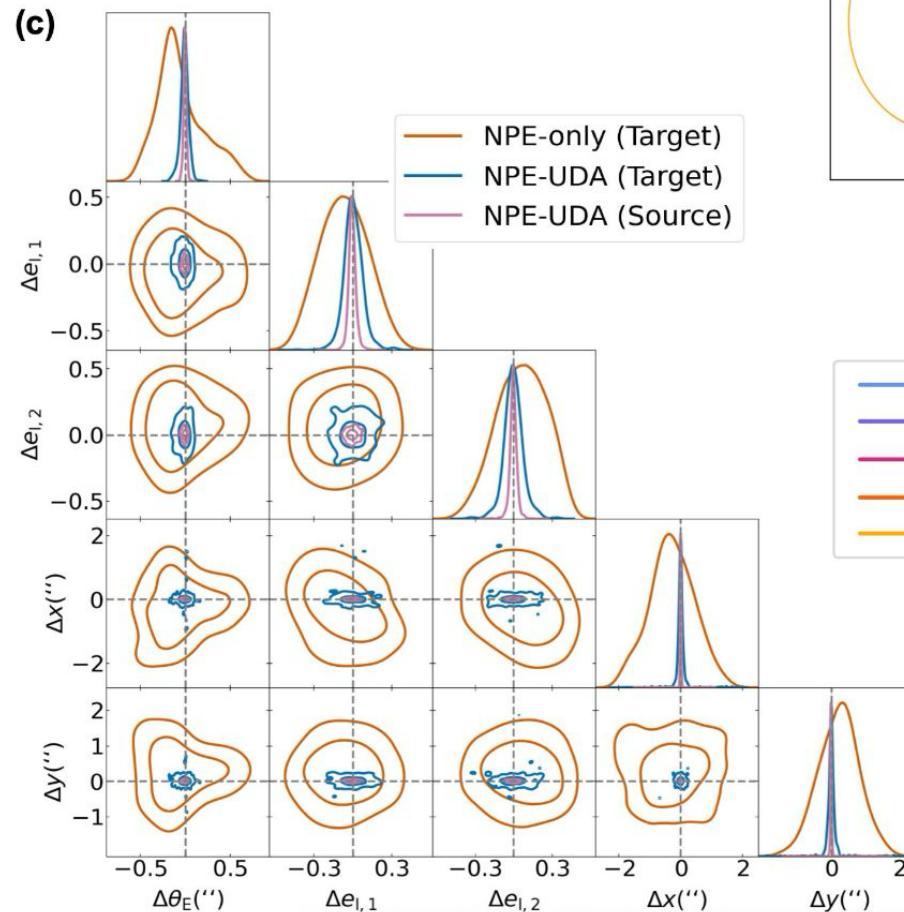
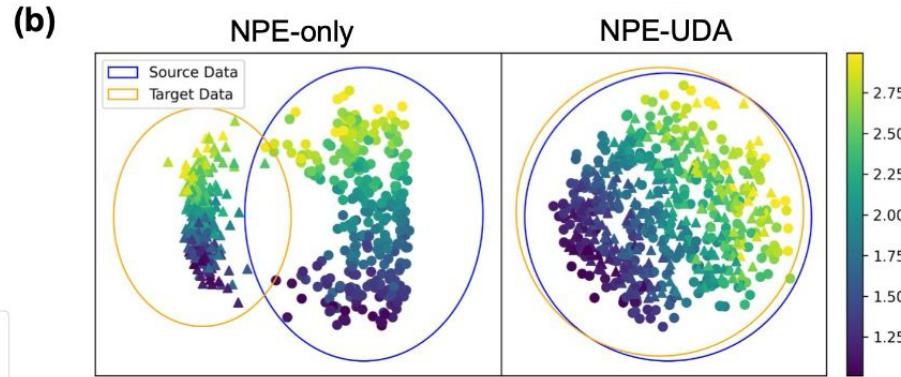
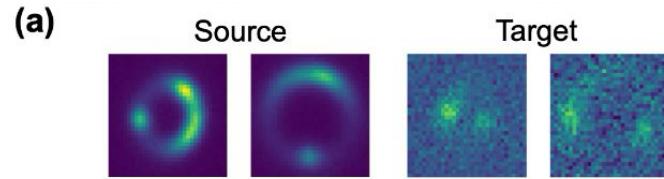
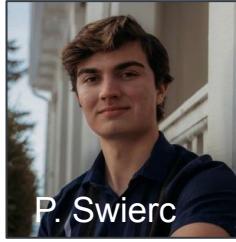
- a. 100k images (1-filter) with and without noise
- b. **Latent distributions are correctly aligned.**
- c. Differences between true values and predicted posteriors are small. **Accuracy improves up to 2 orders of magnitude.**
- d. Posterior are well calibrated and no longer overconfident.

Swierc et al. 2024.



Strong gravitational lensing

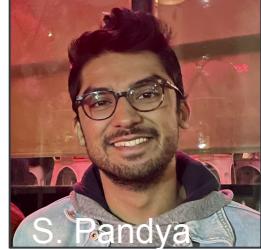
SBI - infer Einstein radius, ellipticity, position/offset



- 100k images (1-filter) with and without noise
- Latent distributions are correctly aligned.**
- Differences between true values and predicted posteriors are small. **Accuracy improves up to 2 orders of magnitude.**
- Posterior are well calibrated and no longer overconfident.



SIDDA: SInkhorn Dynamic Domain Adaptation



S. Pandya

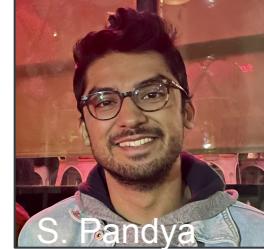
Can we perform automated domain alignment and avoid time-consuming hyperparameter tuning?





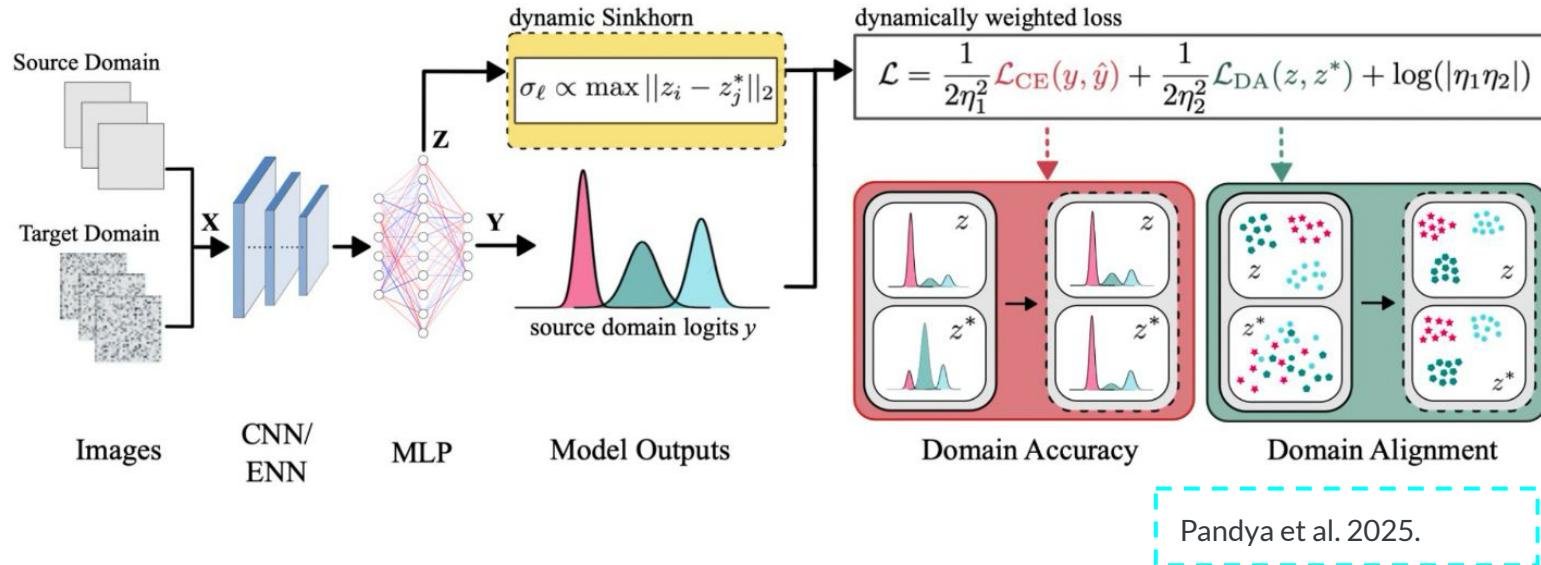
SIDDA: Sinkhorn Dynamic Domain Adaptation

Automated domain alignment for “any” application



S. Pandya

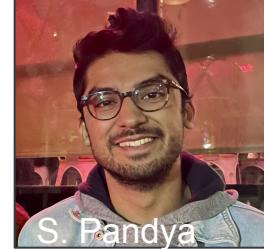
- **Trainable scaling** of the main task loss and DA loss.
- **Trainable Sinkhorn plan** i.e., how detailed should the distance measure be.
- **No tuning needed!**
- Up to **40% better accuracy** on unlabeled data.





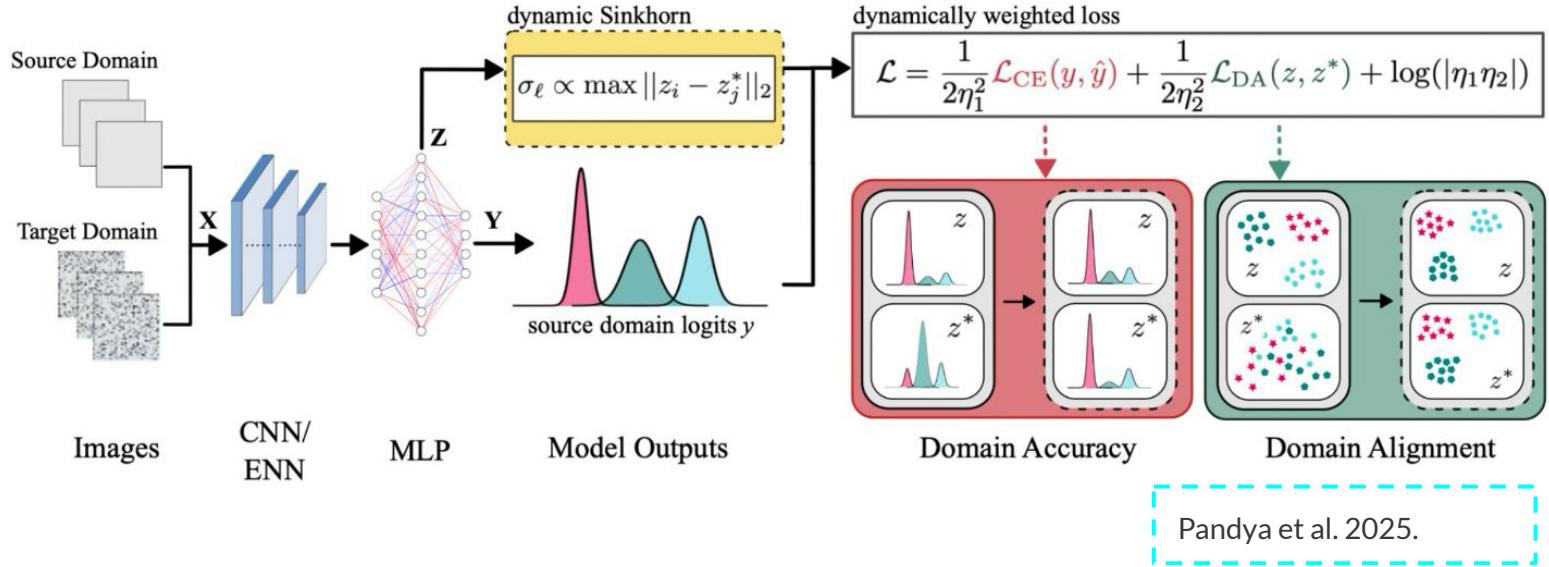
SIDDA: Sinkhorn Dynamic Domain Adaptation

Automated domain alignment for “any” application

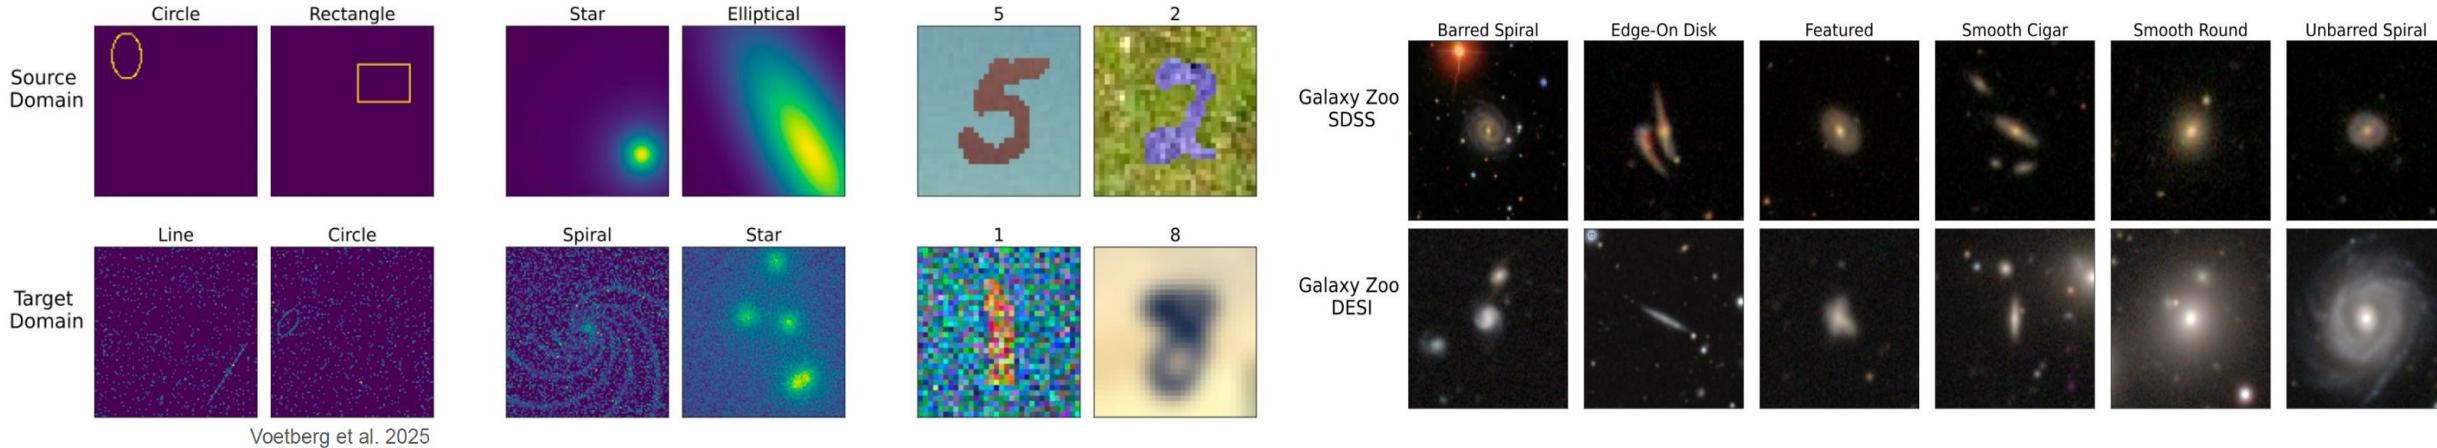


S. Pandya

- **Trainable scaling** of the main task loss and DA loss.
- **Trainable Sinkhorn plan** i.e., how detailed should the distance measure be.
- **No tuning needed!**
- Up to **40% better accuracy** on unlabeled data.



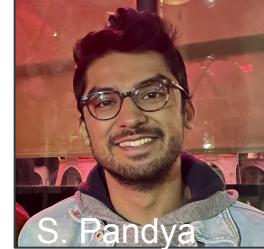
Pandya et al. 2025.



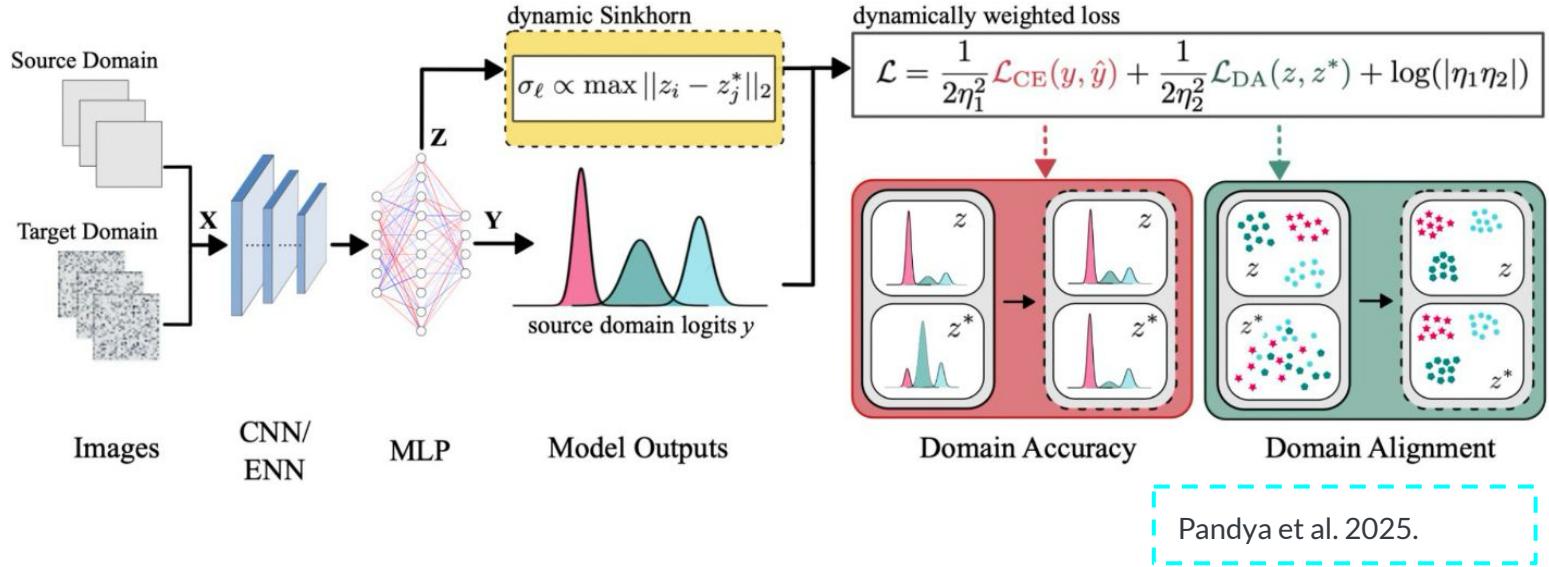


SIDDA: Sinkhorn Dynamic Domain Adaptation

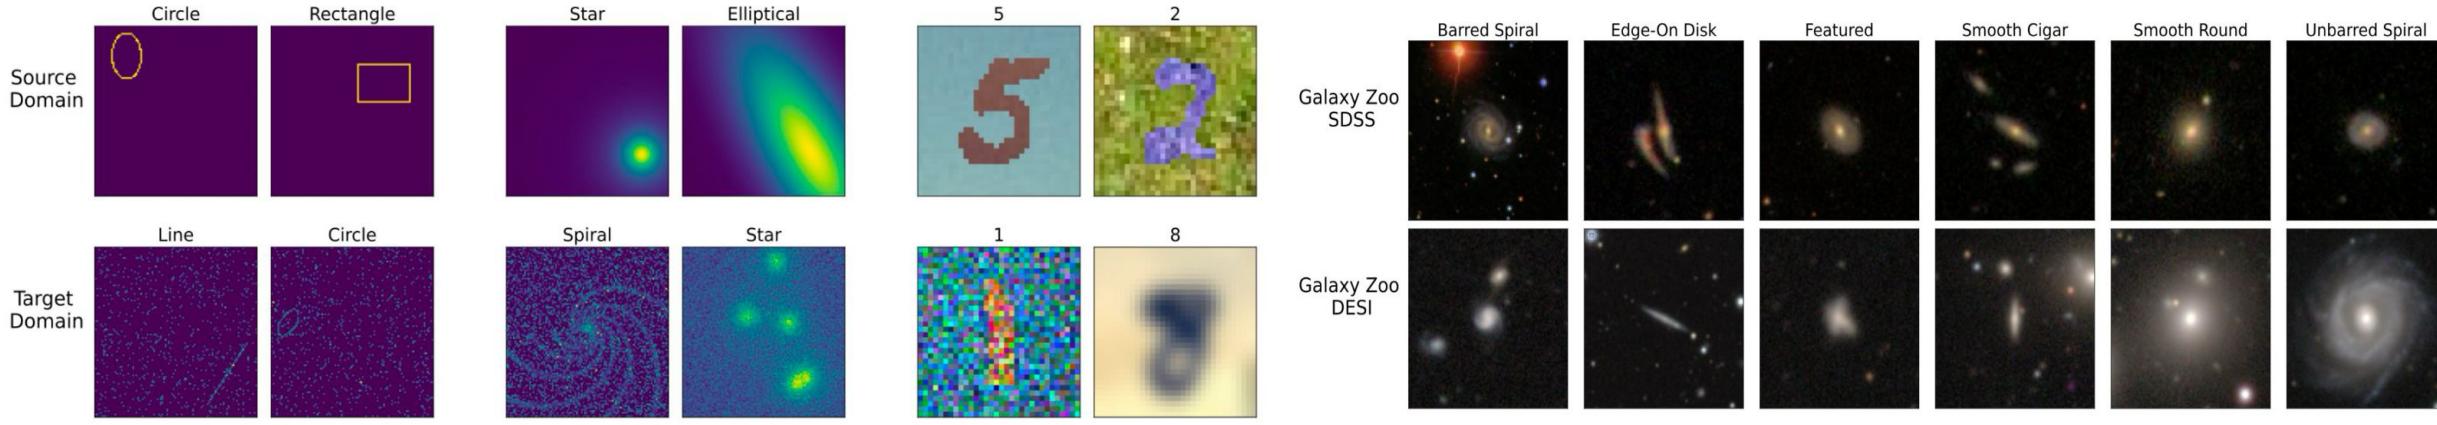
Automated domain alignment for “any” application



- **Trainable scaling** of the main task loss and DA loss.
- **Trainable Sinkhorn plan** i.e., how detailed should the distance measure be.
- **No tuning needed!**
- Up to **40% better accuracy** on unlabeled data.



Pandya et al. 2025.



Voetberg et al. 2025

Stay tuned for the hands-on session to talk to Sneh!

04

Diving Deeper Into Distance-Based Methods

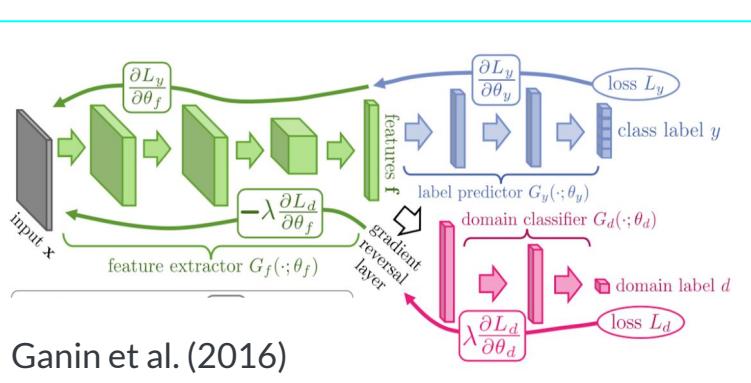
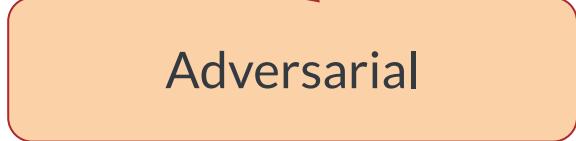
Useful Resources:

- “A Kernel Two-Sample Test” Gretton et al. 2012., [Video 1](#), [Video 2](#)
- “Domain Adversarial Training of Neural Networks” Ganin et al. 2016.
- “Deep Reconstruction-Classification Networks for Unsupervised Domain Adaptation” Ghifary et al. 2016.
- “Interpolating between Optimal Transport and MMD using Sinkhorn Divergences” Feydy et al. 2018

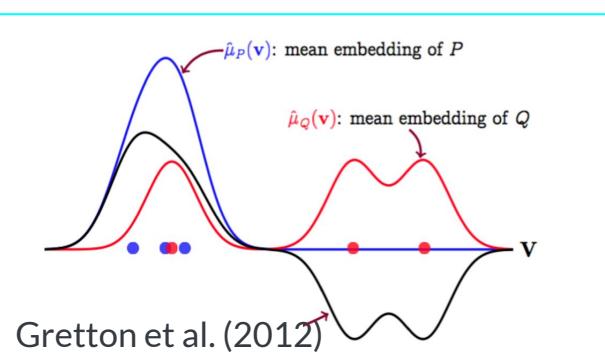
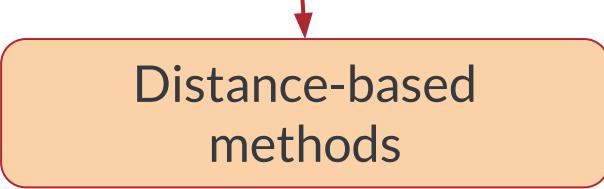
Combining Datasets

DOMAIN ADAPTATION

Align data distributions in the latent space of the network by forcing the network to **find more robust domain-invariant features**.



Using domain discriminators to encourage domain confusion through an adversarial objective.



Minimize a distance metric between source and target data.



Ghifary et al. (2016)

Data reconstruction as an auxiliary task to ensure feature invariance.



Common types of distance measures

φ -divergences

Czisar 1963.

- Use convex function φ to measure the ratio of the two distributions.
- Likelihood based models and information theory.

$$D_\varphi(\alpha\|\beta) = \int \varphi\left(\frac{d\alpha}{d\beta}\right) d\beta$$

- Examples:

Kullback-Liebler (KL) divergence ($\varphi=x \log x$)
Jensen-Shannon (JS) distance
Total Variation (TV) distance
 \square^2 distance

Kernel distances

Gretton 2006. ("king of kernels")

- Kernel-based metric between distributions.
- Measures distance in Reproducing Kernel Hilbert Space (RKHS).
- Used in generative modeling and domain adaptation.

$$\begin{aligned} \text{MMD}^2(\alpha, \beta) &= \mathbb{E}_{x,x' \sim \alpha}[k(x, x')] + \mathbb{E}_{y,y' \sim \beta}[k(y, y')] \\ &\quad - 2\mathbb{E}_{x \sim \alpha, y \sim \beta}[k(x, y)] \end{aligned}$$

Optimal transport

Monge 1781, Kantorovich 1942

- Finding the transport plan π that minimized the total cost of moving all the mass (c - the cost of moving a unit mass) from α to \square ?

$$\text{OT}(\alpha, \beta) = \inf_{\gamma \in \Pi(\alpha, \beta)} \int c(x, y) d\gamma(x, y)$$

$$c(x, y) = \|x - y\|^p$$

- Examples metrics:

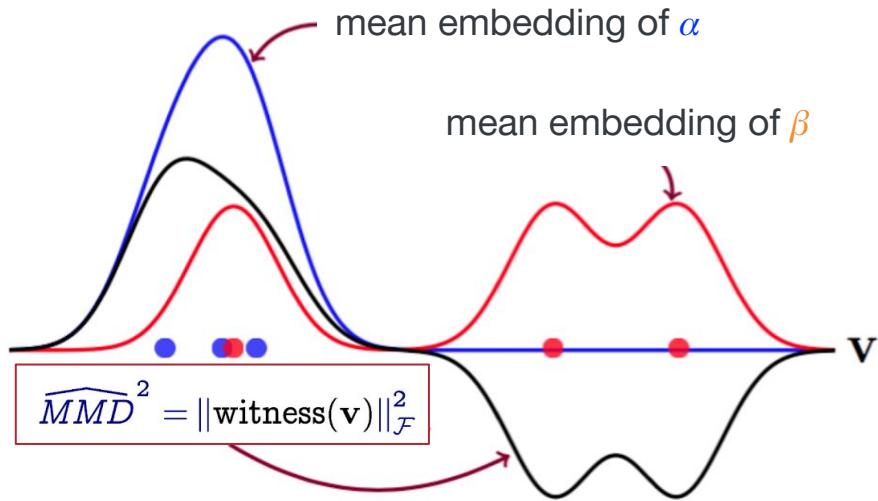
Wasserstein distance ($p=2$)
Earth movers distance ($p=1$)
Sinkhorn divergence (regularized OT)



Maximum Mean Discrepancy (MMD)

Gaussian (RBF) kernel

$$k(x, y) = \exp\left(-\frac{\|x - y\|^2}{2\sigma^2}\right)$$



$$\text{MMD}^2(\alpha, \beta) = \mathbb{E}_{x, x' \sim \alpha}[k(x, x')] + \mathbb{E}_{y, y' \sim \beta}[k(y, y')] - 2\mathbb{E}_{x \sim \alpha, y \sim \beta}[k(x, y)]$$

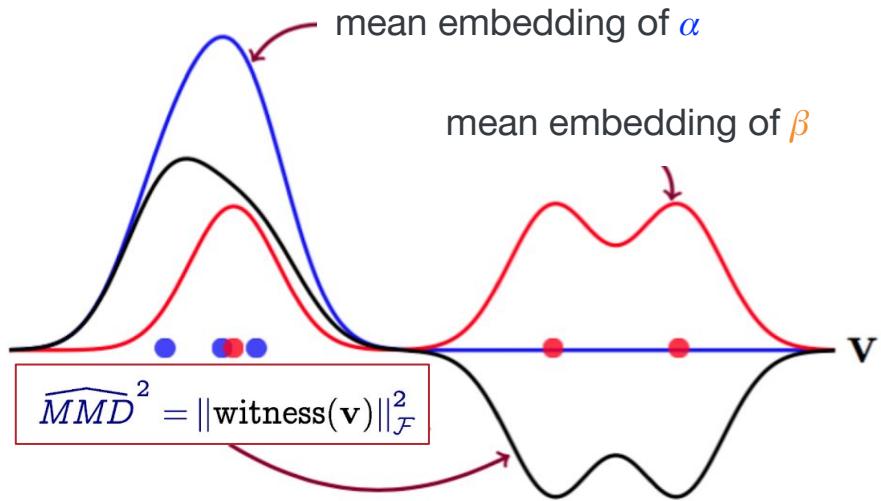


Maximum Mean Discrepancy (MMD)

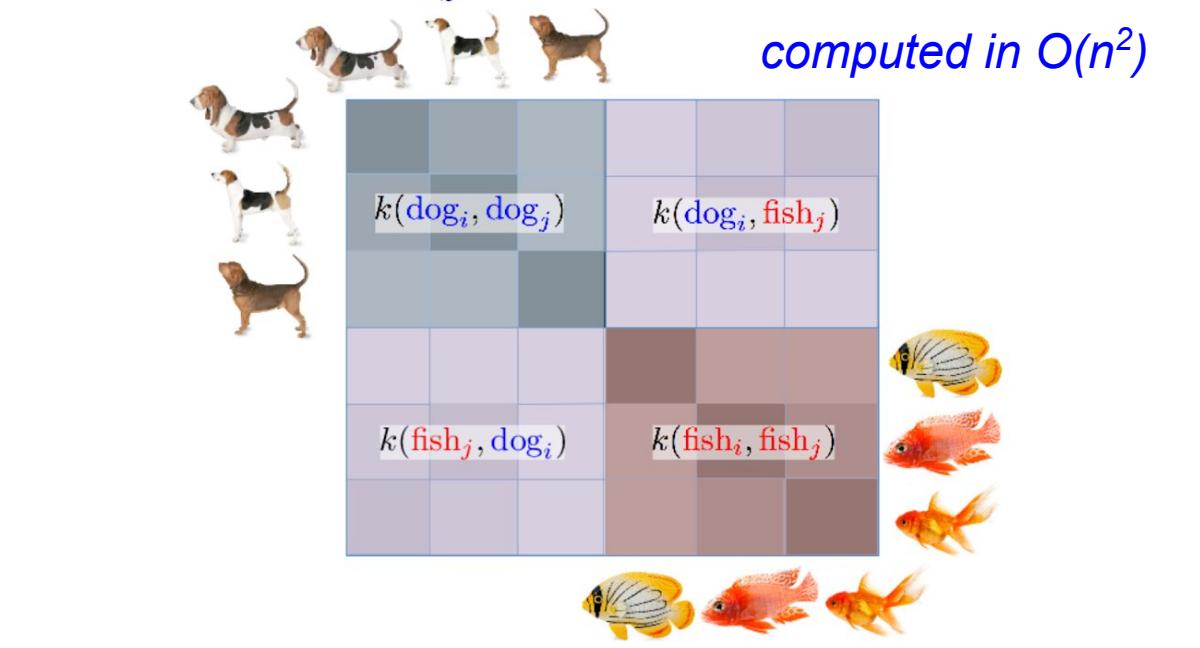
Smola et al. (2007)
Gretton et al. (2012)

Gaussian (RBF) kernel

$$k(x, y) = \exp\left(-\frac{\|x - y\|^2}{2\sigma^2}\right)$$



$$\begin{aligned} \widehat{MMD}^2 &= \frac{1}{n(n-1)} \sum_{i \neq j} k(\text{dog}_i, \text{dog}_j) + \frac{1}{n(n-1)} \sum_{i \neq j} k(\text{fish}_i, \text{fish}_j) \\ &\quad - \frac{2}{n^2} \sum_{i,j} k(\text{dog}_i, \text{fish}_j) \end{aligned}$$



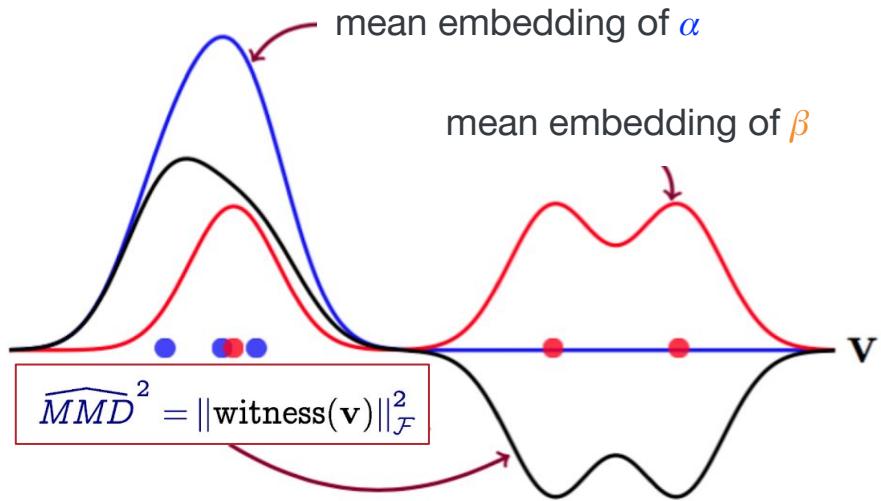
$$MMD^2(\alpha, \beta) = \mathbb{E}_{x,x' \sim \alpha}[k(x, x')] + \mathbb{E}_{y,y' \sim \beta}[k(y, y')] - 2\mathbb{E}_{x \sim \alpha, y \sim \beta}[k(x, y)]$$

Maximum Mean Discrepancy (MMD)

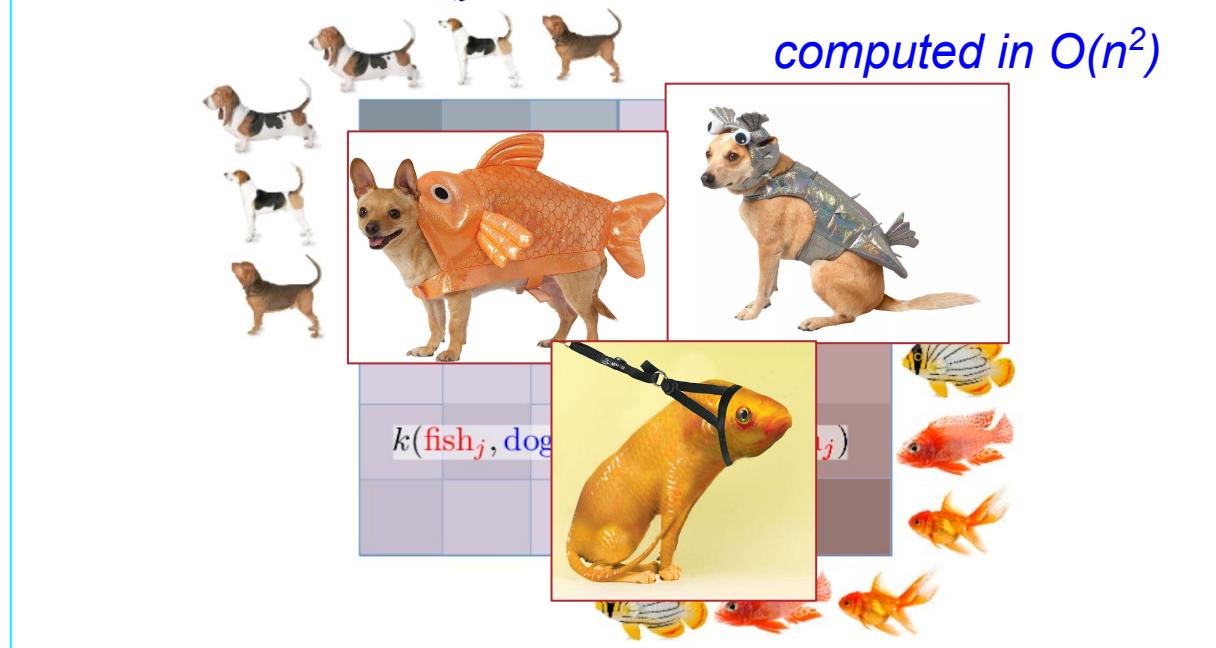
Smola et al. (2007)
Gretton et al. (2012)

Gaussian (RBF) kernel

$$k(x, y) = \exp\left(-\frac{\|x - y\|^2}{2\sigma^2}\right)$$



$$\begin{aligned} \widehat{MMD}^2 &= \frac{1}{n(n-1)} \sum_{i \neq j} k(\text{dog}_i, \text{dog}_j) + \frac{1}{n(n-1)} \sum_{i \neq j} k(\text{fish}_i, \text{fish}_j) \\ &\quad - \frac{2}{n^2} \sum_{i,j} k(\text{dog}_i, \text{fish}_j) \end{aligned}$$

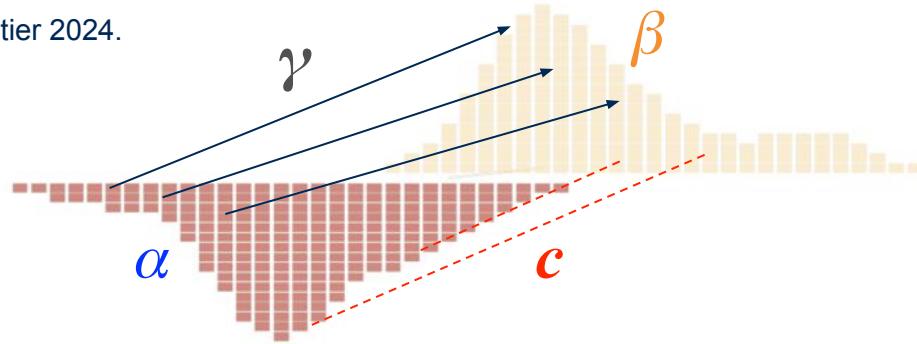


$$\text{MMD}^2(\alpha, \beta) = \mathbb{E}_{x,x' \sim \alpha}[k(x, x')] + \mathbb{E}_{y,y' \sim \beta}[k(y, y')] - 2\mathbb{E}_{x \sim \alpha, y \sim \beta}[k(x, y)]$$



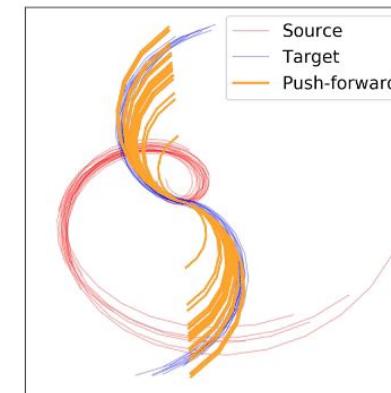
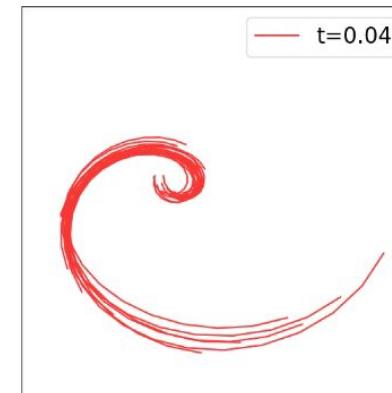
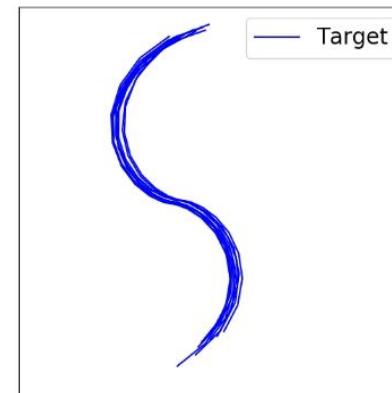
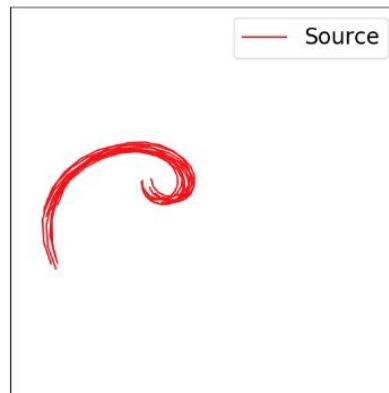
Optimal Transport

Source: Charpentier 2024.



OPTIMAL TRANSPORT:

Finding the “most efficient” transportation plan γ that minimizes the total transportation cost c of moving the entire mass from probability α distribution to β .



Source: Zhu et al. 2023.

Optimal Transport

Finding the “most efficient” transportation plan γ that minimizes the total transportation cost c of moving the entire mass from probability α distribution to β .

Optimal Transport

$$\text{OT}(\alpha, \beta) = \inf_{\gamma \in \Pi(\alpha, \beta)} \int c(x, y) d\gamma(x, y)$$

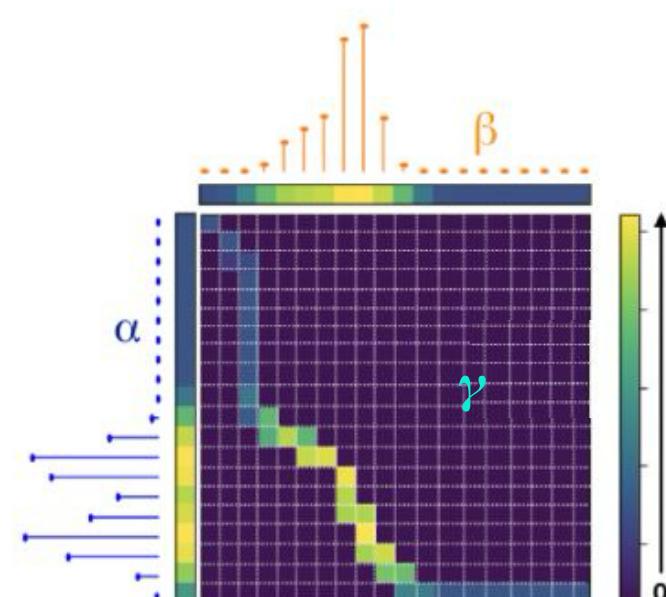
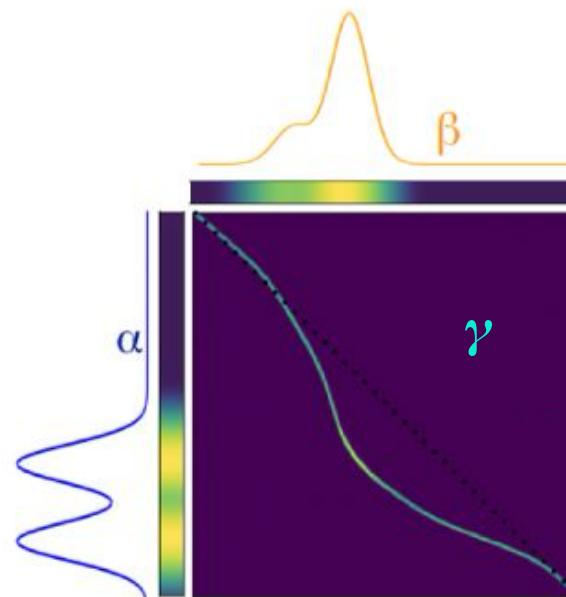
Solving OT exactly becomes impossible for large datasets as it scales as $O(n^3 \log n)$!

Wasserstein distance
 $p=2$

$$W_p(\alpha, \beta) = \left(\inf_{\gamma \in \Pi(\alpha, \beta)} \int \|x - y\|^p d\gamma(x, y) \right)^{1/p}$$

$c(x, y)$ cost of moving a unit of mass from x to y

$\gamma(x, y)$ transport plan (coupling) - how much mass moves from x to y



Genevay 2019.



Optimal Transport

Finding the “most efficient” transportation plan γ that minimizes the total transportation cost c of moving the entire mass from probability α distribution to β .

Optimal Transport

$$\text{OT}(\alpha, \beta) = \inf_{\gamma \in \Pi(\alpha, \beta)} \int c(x, y) d\gamma(x, y)$$

Solving OT exactly becomes impossible for large datasets as it scales as $O(n^3 \log n)$!

Wasserstein distance
 $p=2$

$$W_p(\alpha, \beta) = \left(\inf_{\gamma \in \Pi(\alpha, \beta)} \int \|x - y\|^p d\gamma(x, y) \right)^{1/p}$$

Entropy Regularization

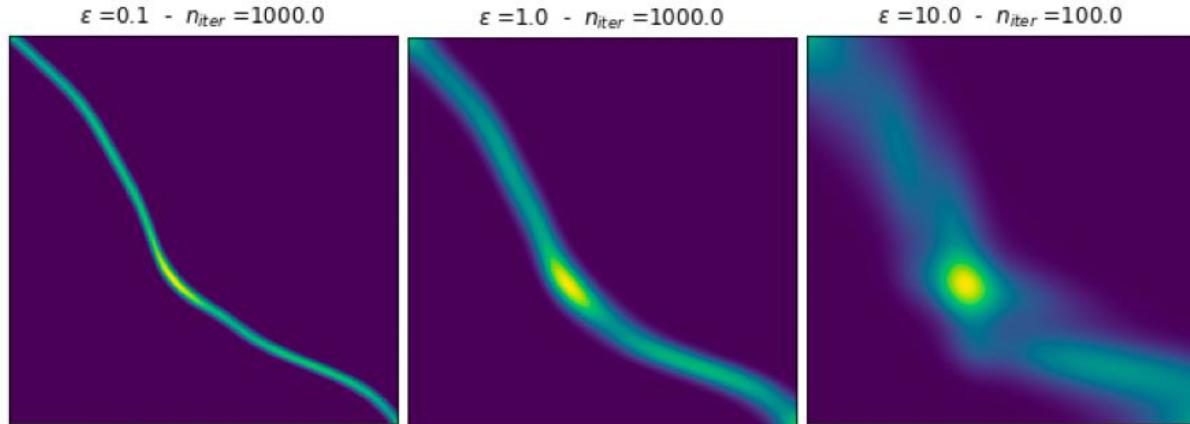
$$OT_{\varepsilon, p}(\alpha, \beta) = \inf_{\gamma \in \Pi(\alpha, \beta)} \int_{\mathcal{X} \times \mathcal{Y}} \|x - y\|^p d\gamma(x, y) - \varepsilon \int_{\mathcal{X} \times \mathcal{Y}} \log \left(\frac{d\gamma}{dx dy} \right) d\gamma(x, y)$$

The **entropy term** smooths the optimization and ensures faster convergence with better numerical properties. ε controls the **strength of the entropy regularization**.

Optimal Transport

Finding the “most efficient” transportation plan γ that minimizes the total transportation cost c of moving the entire mass from probability α distribution to β .

ε smooths the problem to avoid overfitting.



Genevay 2019.

Entropy Regularization

$$OT_{\varepsilon,p}(\alpha, \beta) = \inf_{\gamma \in \Pi(\alpha, \beta)} \int_{\mathcal{X} \times \mathcal{Y}} \|x - y\|^p d\gamma(x, y) - \varepsilon \int_{\mathcal{X} \times \mathcal{Y}} \log \left(\frac{d\gamma}{dxdy} \right) d\gamma(x, y)$$

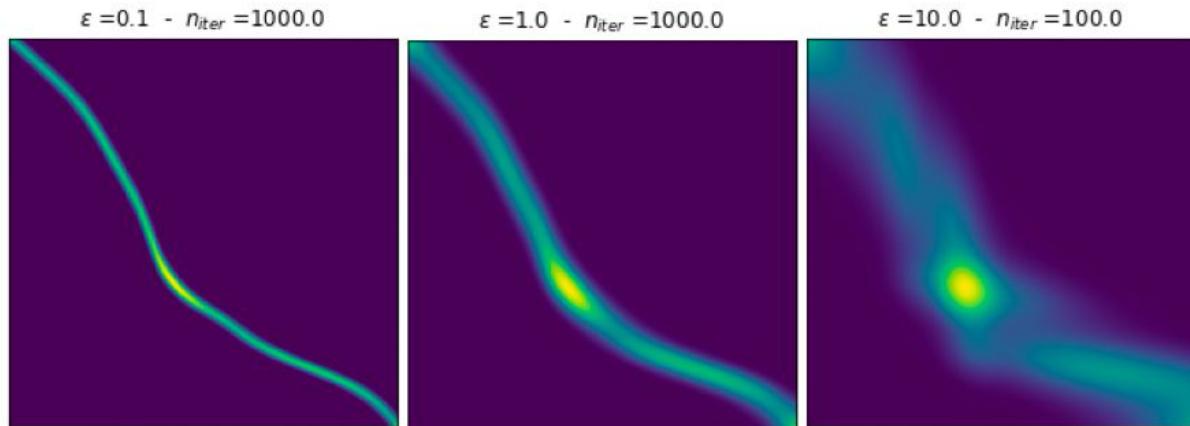
The **entropy term** smooths the optimization and ensures faster convergence with better numerical properties. ε controls the **strength of the entropy regularization**.

Optimal Transport

Finding the “most efficient” transportation plan γ that minimizes the total transportation cost c of moving the entire mass from probability α distribution to β .

ε smooths the problem to avoid overfitting.

But **Entropy Regularization introduces a bias**. $OT_{\varepsilon,p}(\alpha, \alpha) \neq 0$



Entropy Regularization

$$OT_{\varepsilon,p}(\alpha, \beta) = \inf_{\gamma \in \Pi(\alpha, \beta)} \int_{\mathcal{X} \times \mathcal{Y}} \|x - y\|^p d\gamma(x, y) - \varepsilon \int_{\mathcal{X} \times \mathcal{Y}} \log \left(\frac{d\gamma}{dx dy} \right) d\gamma(x, y)$$

The **entropy term** smooths the optimization and ensures faster convergence with better numerical properties. ε controls the **strength of the entropy regularization**.

Sinkhorn Algorithm

Debiased and efficient OT

$$\text{Sink}_\varepsilon(\alpha, \beta) = OT_\varepsilon(\alpha, \beta) - \frac{1}{2} [OT_\varepsilon(\alpha, \alpha) + OT_\varepsilon(\beta, \beta)]$$

Entropic OT between source
and target distributions.

Self-distances of source and target
(should be zero in classical OT, but
not in entropic OT).

Asymptotic behaviours:

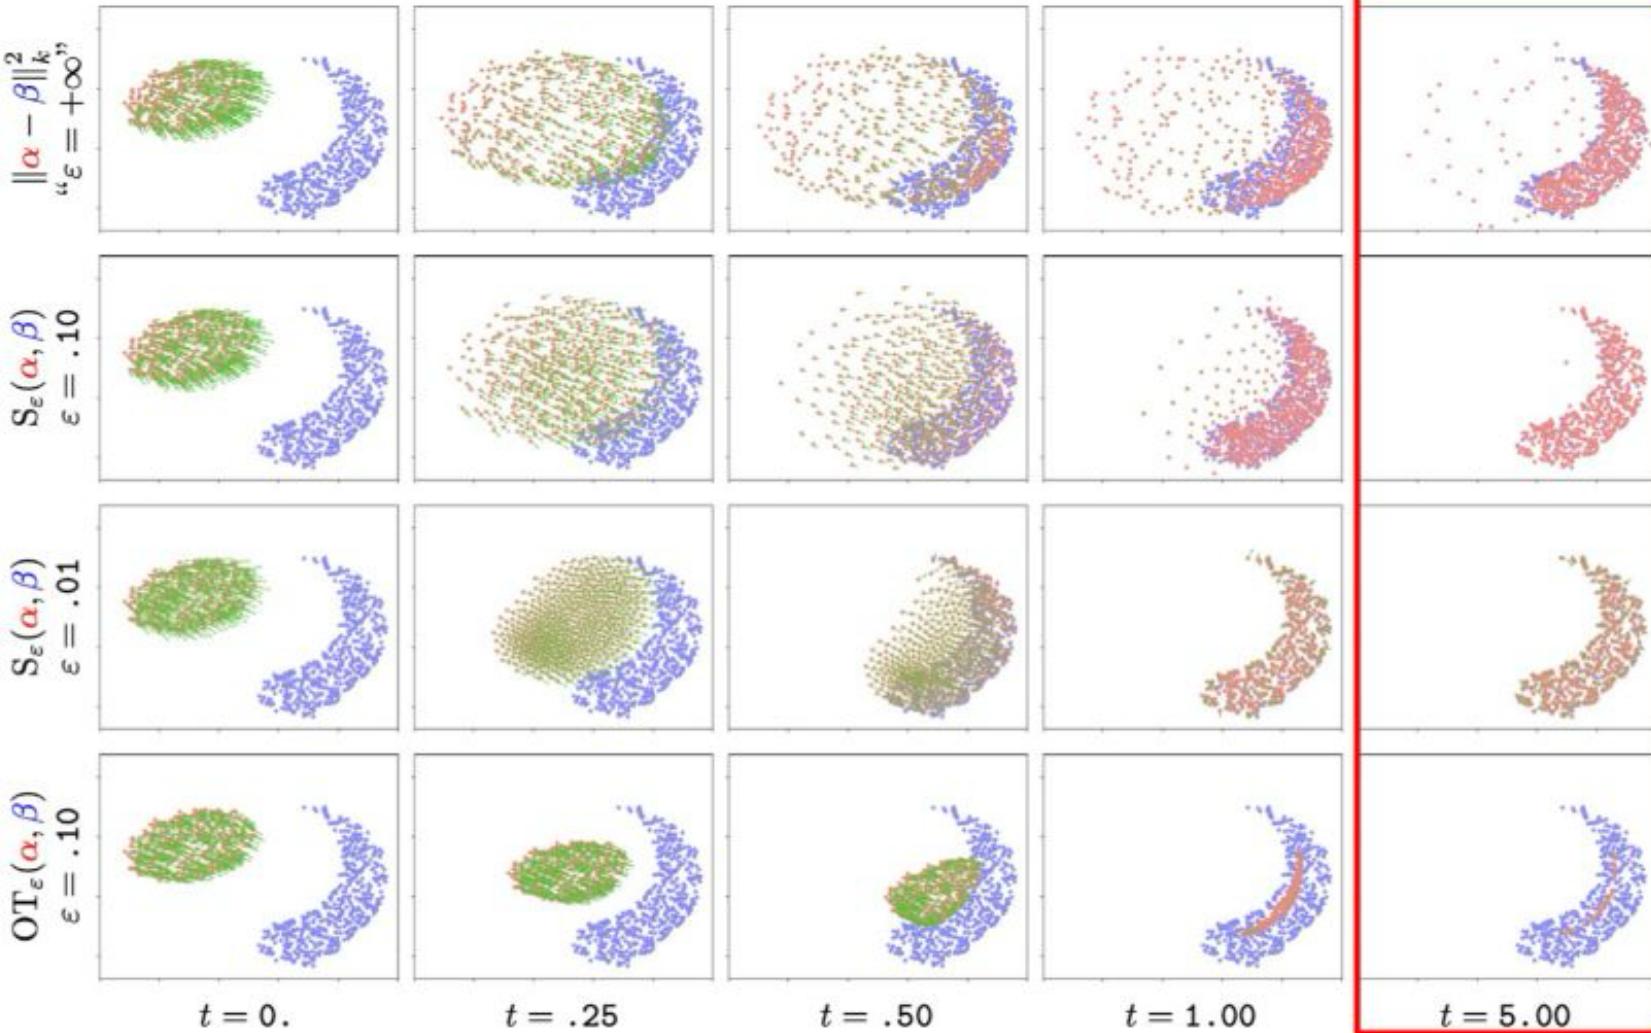
$$\text{Sink}_\varepsilon(\alpha, \beta) \rightarrow \begin{cases} \text{Wasserstein} & \text{as } \varepsilon \rightarrow 0 \\ \text{MMD-like mean-matching} & \text{as } \varepsilon \rightarrow \infty \end{cases}$$

Feydy et al. 2018.

Sinkhorn - best of both worlds

Interpolating between MMD and Optimal Transport (OT)

MMD



Sinkhorn

OT

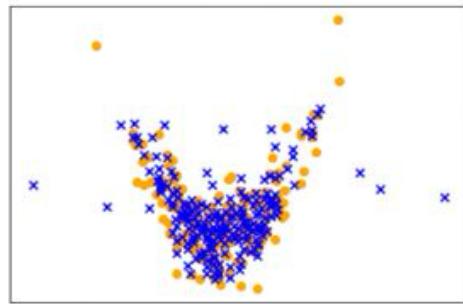
MMD and Optimal Transport (OT)

MMD PROS:

- **Simple and fast** to compute - $O(n^2)$.
- **Differentiable**
- Works well in **high dimensions** (if kernel is well-chosen).
- Captures subtle distributional differences.

CONS:

- **Sensitive to kernel choice.**
- Bad in low-density environments.
- Fails to capture **geometric structure**
- Less interpretable.



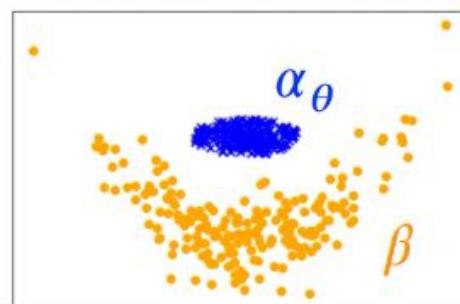
$$MMD_k - k = - \|\cdot\|_2^{1.5}$$

OT PROS:

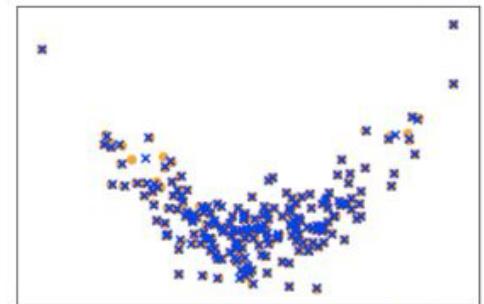
- Captures **geometry of distributions**.
- More **interpretable**: measures "how much mass needs to move".
- Good for structured data (images, point clouds, time series).

CONS:

- **Computationally expensive** $O(n^3 \log(n))$.
- Curse of dimensionality.
- **Non-differentiable** (needs entropic regularization, but this requires tuning). **Sinkhorn to the rescue!**



Initial Setting



$$W_c - c = \|\cdot\|_2^{1.5}$$

MMD and Optimal Transport (OT)

MMD PROS:

- **Simple and fast** to compute - $O(n^2)$.
- **Differentiable**
- Works well in **high dimensions** (if kernel is well-chosen).
- Captures subtle distributional differences.

CONS:

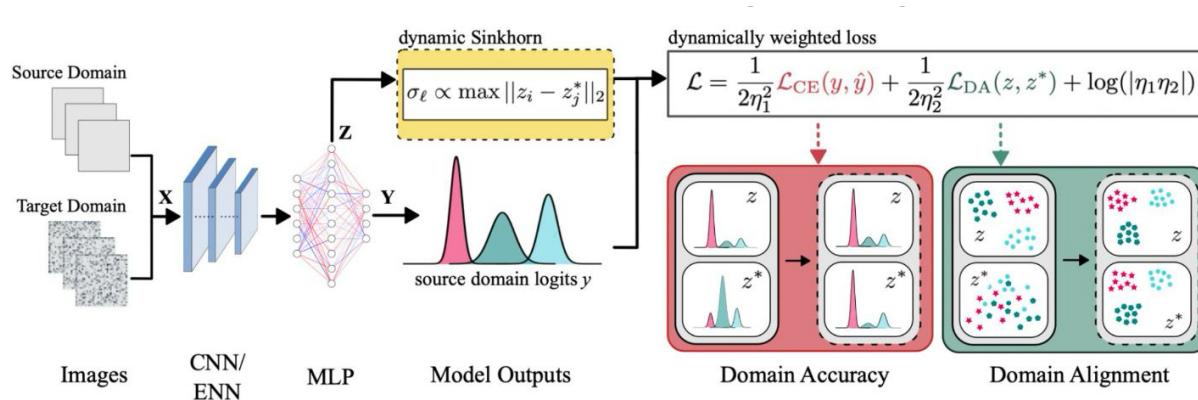
- **Sensitive to kernel choice.**
- Bad in low-density environments.
- Fails to capture **geometric structure**
- Less interpretable.

OT PROS:

- Captures **geometry of distributions**.
- More **interpretable**: measures "how much mass needs to move".
- Good for structured data (images, point clouds, time series).

CONS:

- **Computationally expensive** $O(n^3 \log(n))$.
- Curse of dimensionality.
- **Non-differentiable** (needs entropic regularization, but this requires tuning). **Sinkhorn to the rescue!**



05

Uncertainty Quantification

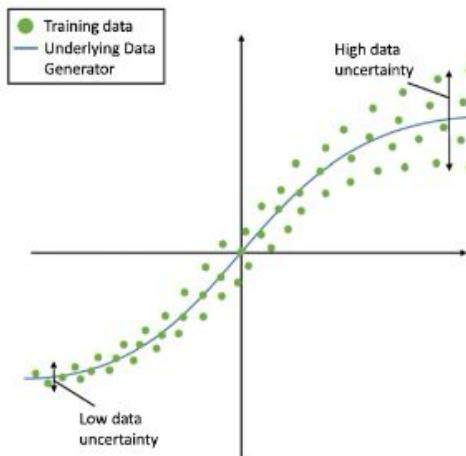
Useful Resources:

- ["A survey of uncertainty in deep neural networks"](#) Gawlikowski et al. 2023.
- ["On the Need to Align Intent and Implementation in Uncertainty Quantification for Machine Learning"](#) Trivedi & Nord 2025.

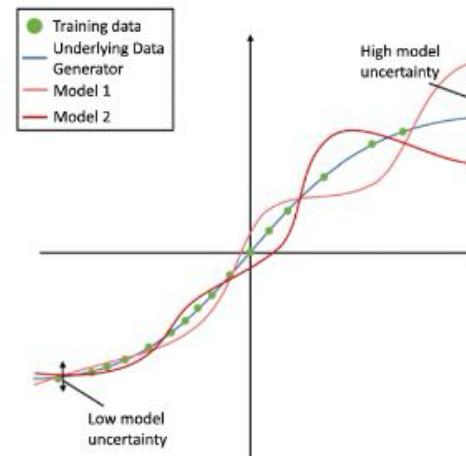


Sources of uncertainties

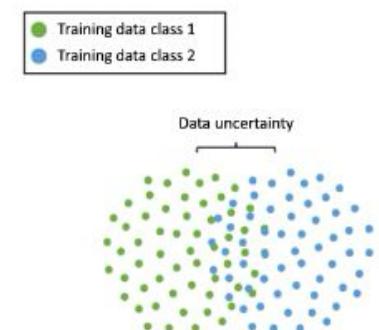
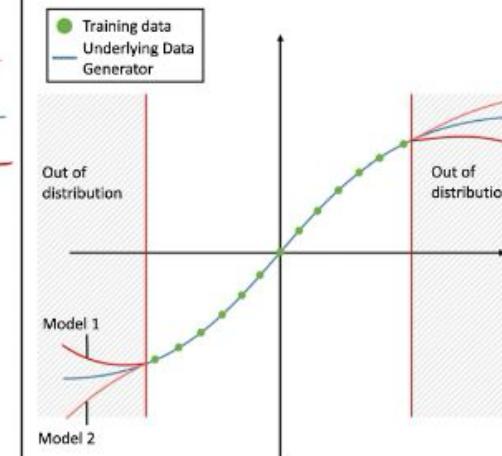
Data



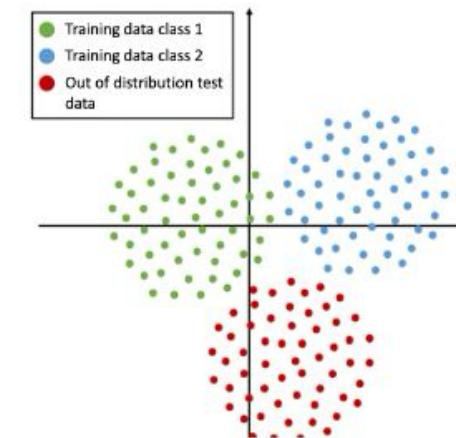
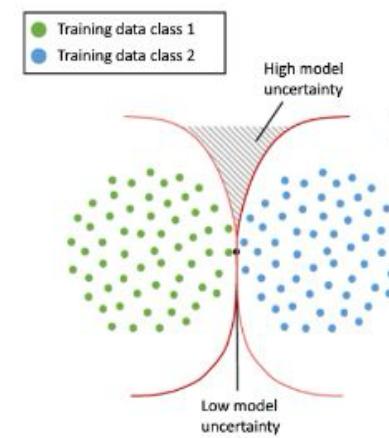
Model



Distributional



Gawlikowski et al. 2023.





Sources of uncertainties

All of this will influence model parameters and model outputs!

Systematic

persistent bias from instruments, calibration, or imperfect theory

Statistical

variation due to limited data

Epistemic

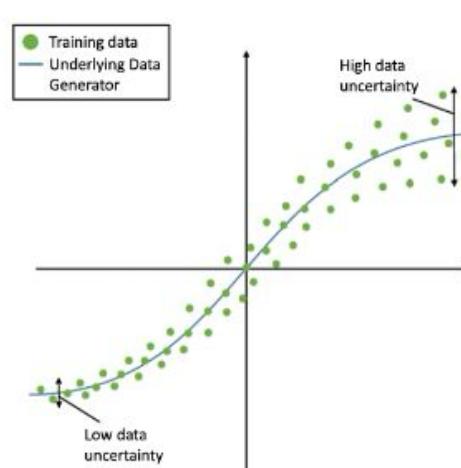
model or data-driven uncertainty

Aleatoric

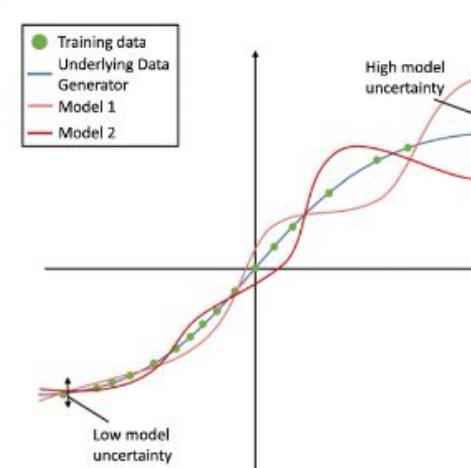
irreducible noise in the data

These vocabularies are not interchangeable!

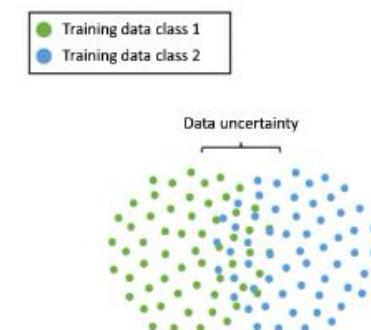
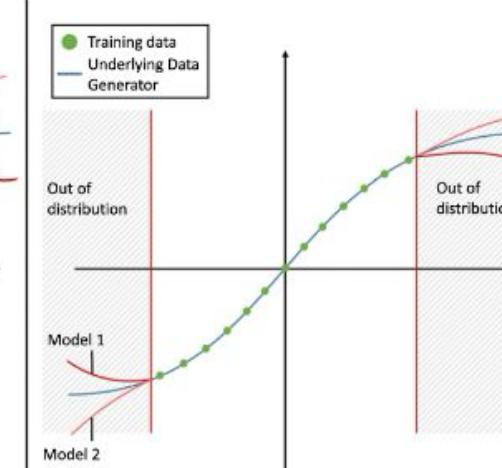
Data



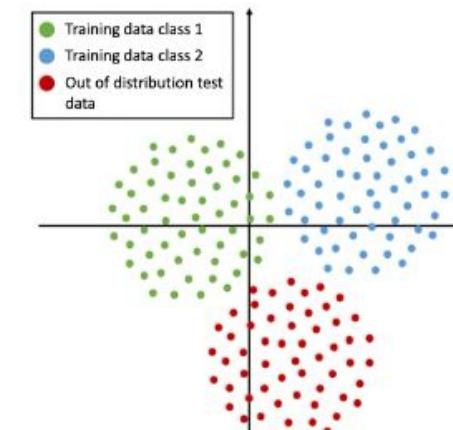
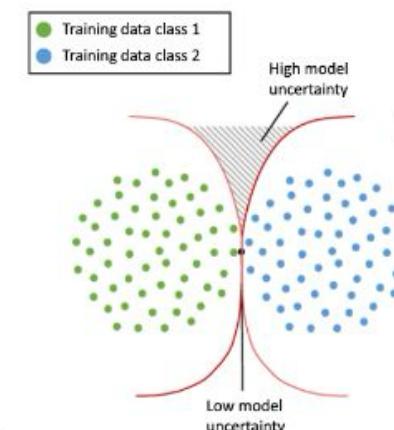
Model



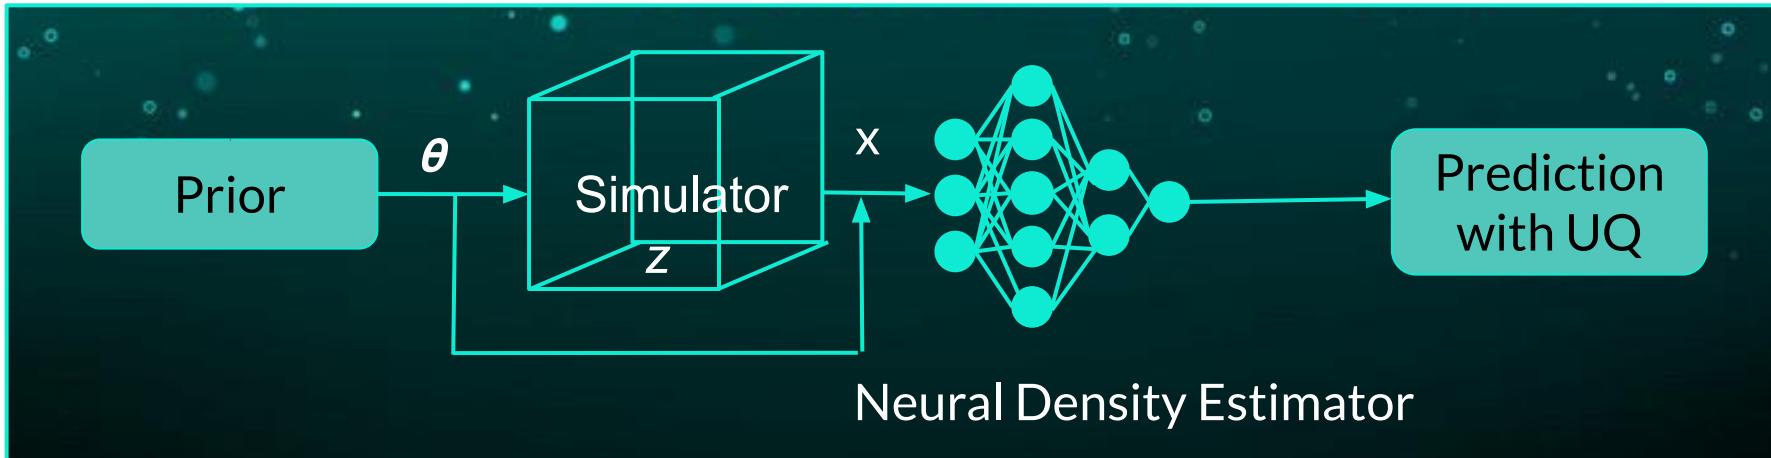
Distributional



Gawlikowski et al. 2023.

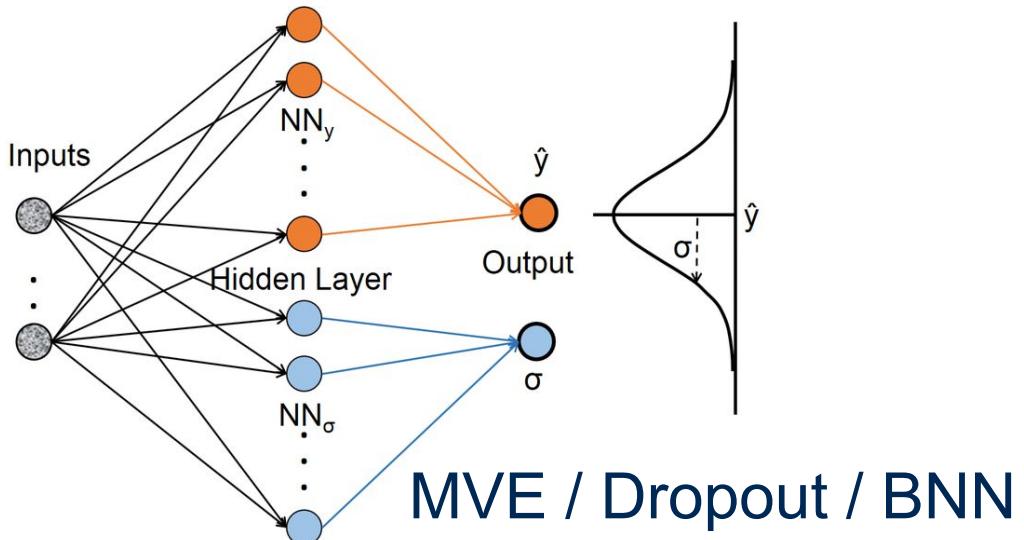
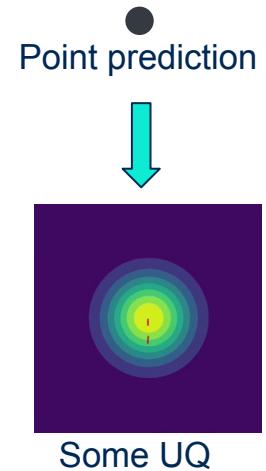
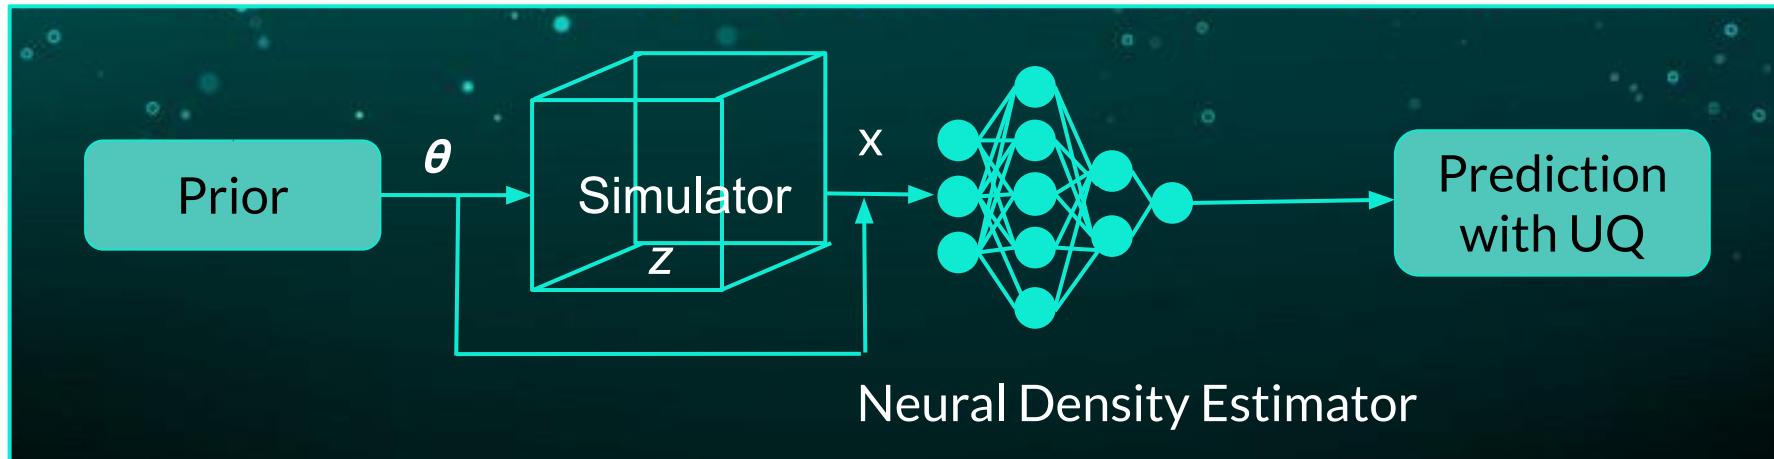


Inference with UQ

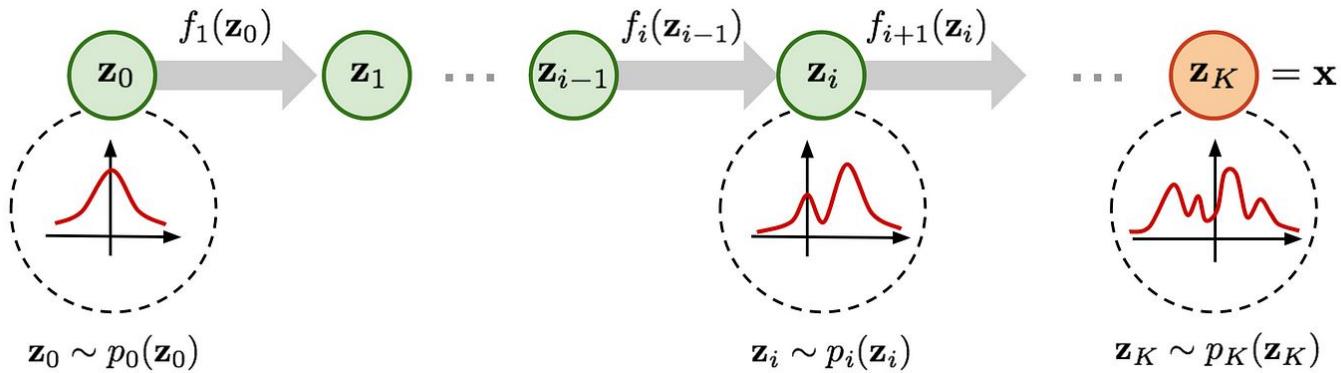
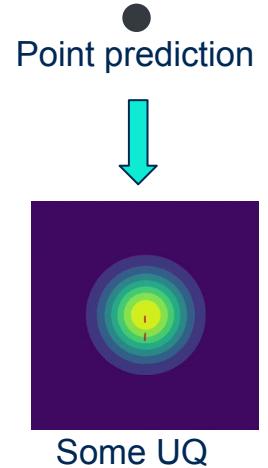
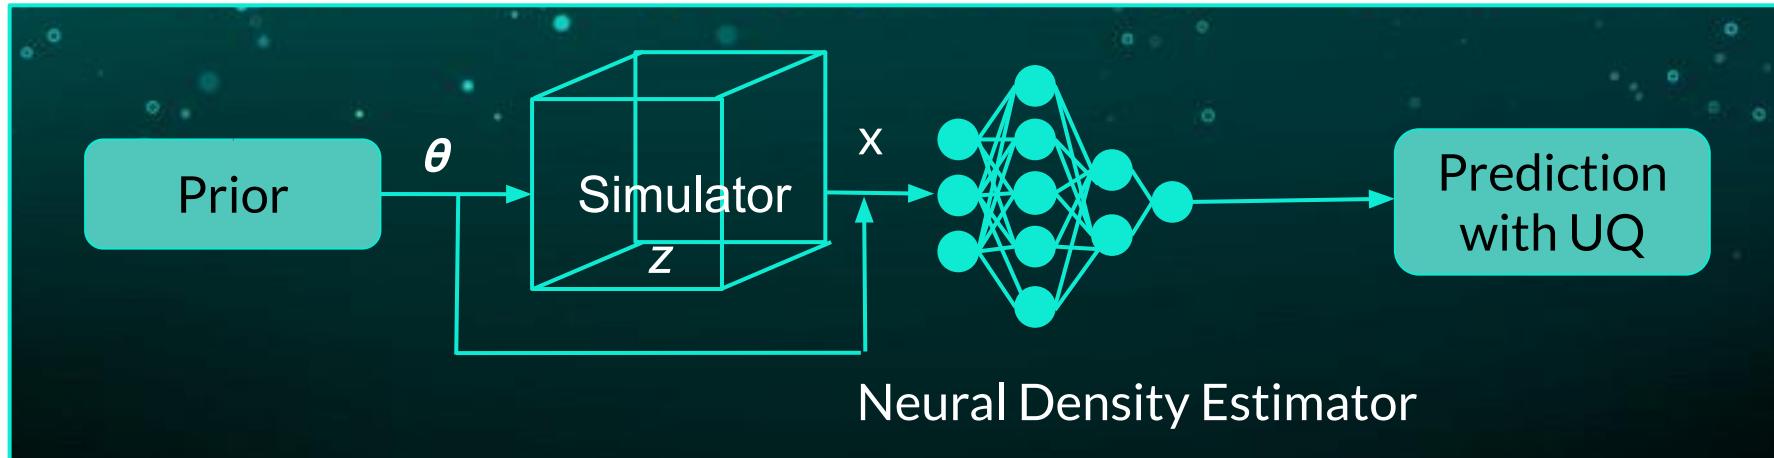


Point prediction

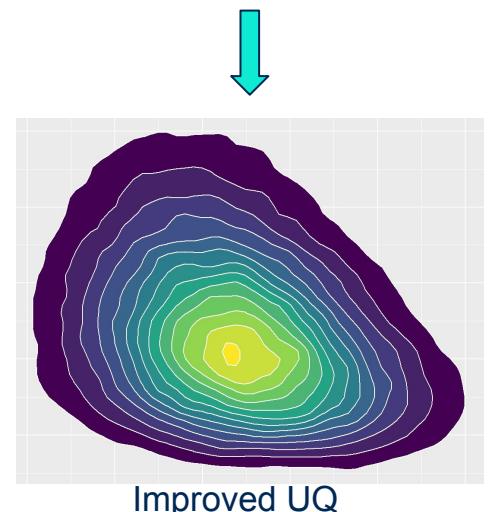
Inference with UQ



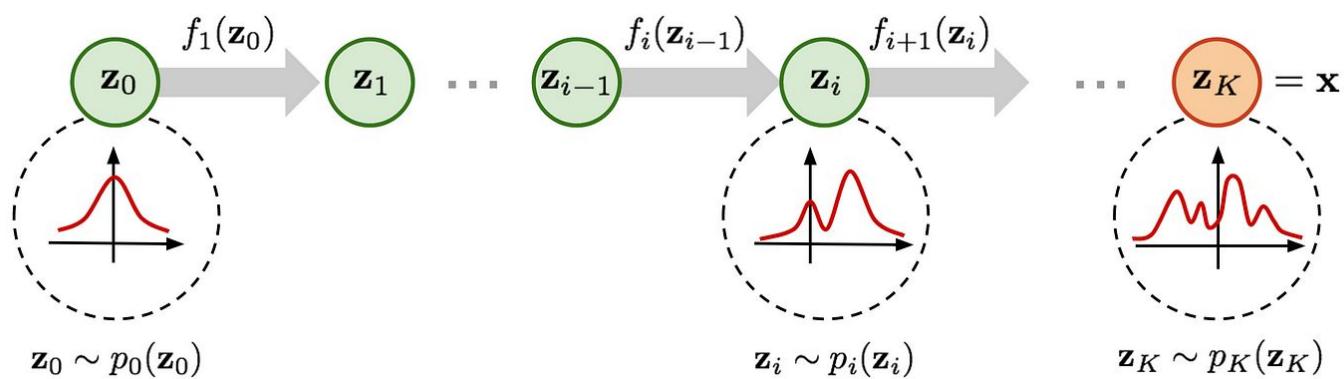
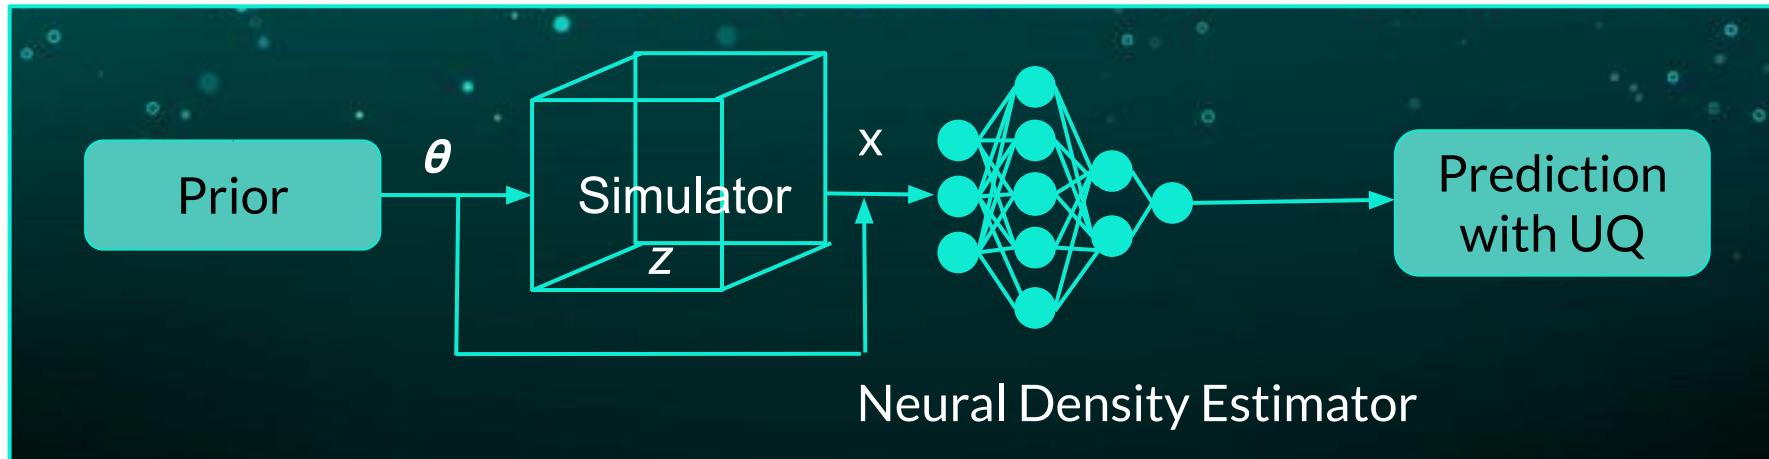
Inference with UQ



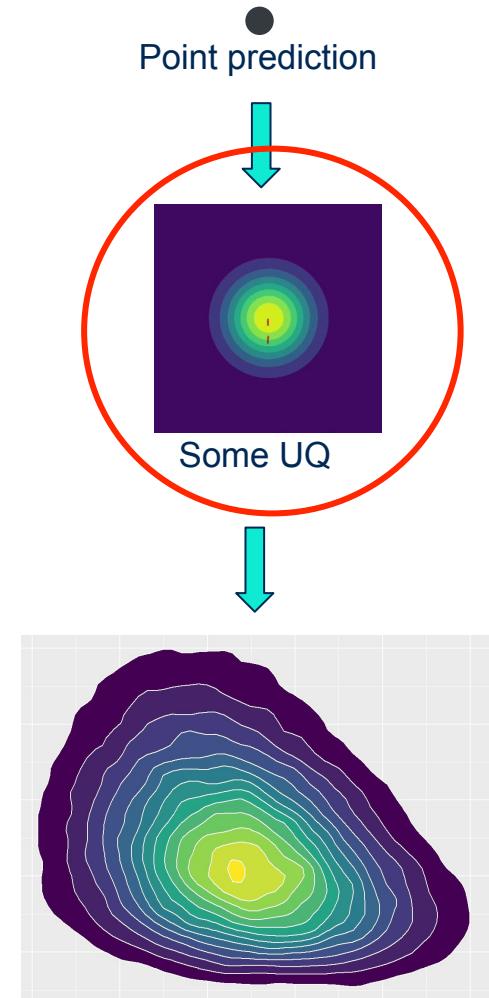
Estimate the full posterior (Normalizing Flows)



Inference with UQ



Estimate the full posterior (Normalizing Flows)

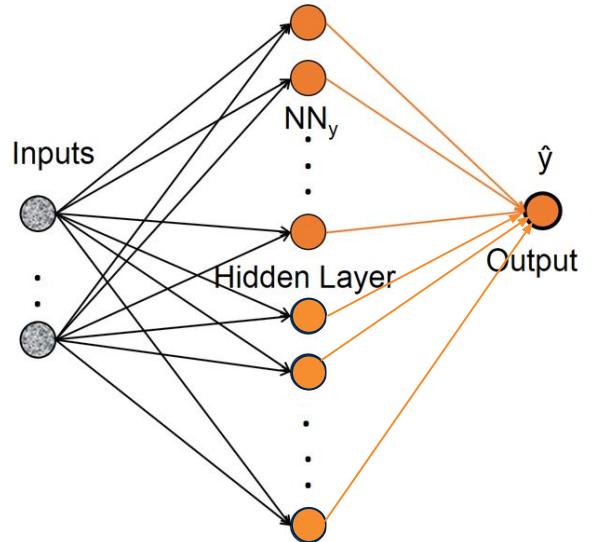




UQ for different types of models

Epistemic, aleatoric or both?

Deterministic Model



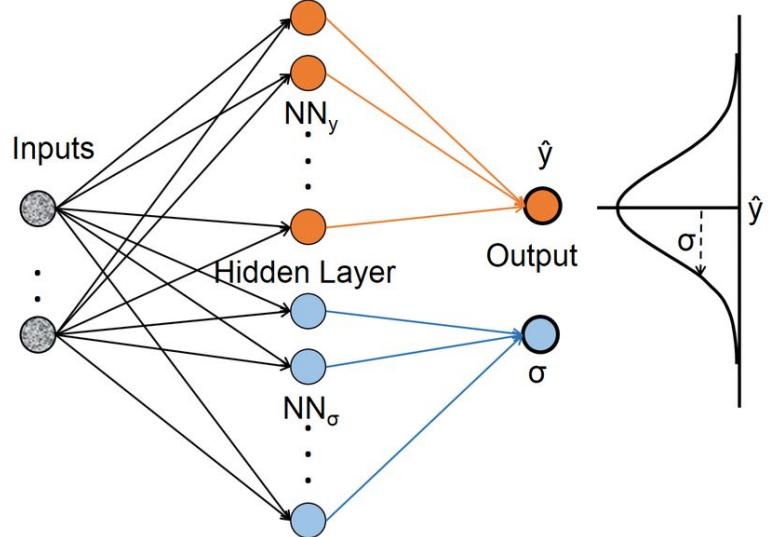
No Errors



UQ for different types of models

Epistemic, aleatoric or both?

Deterministic Model - MVE



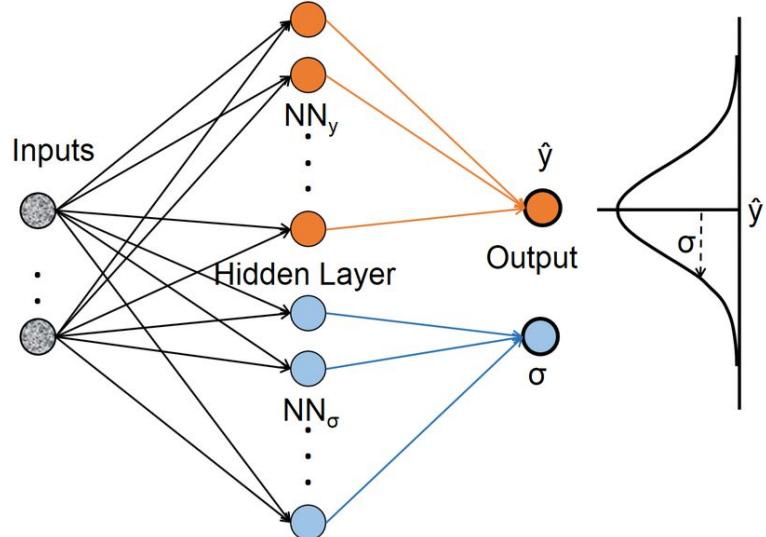
Data (Aleatoric)



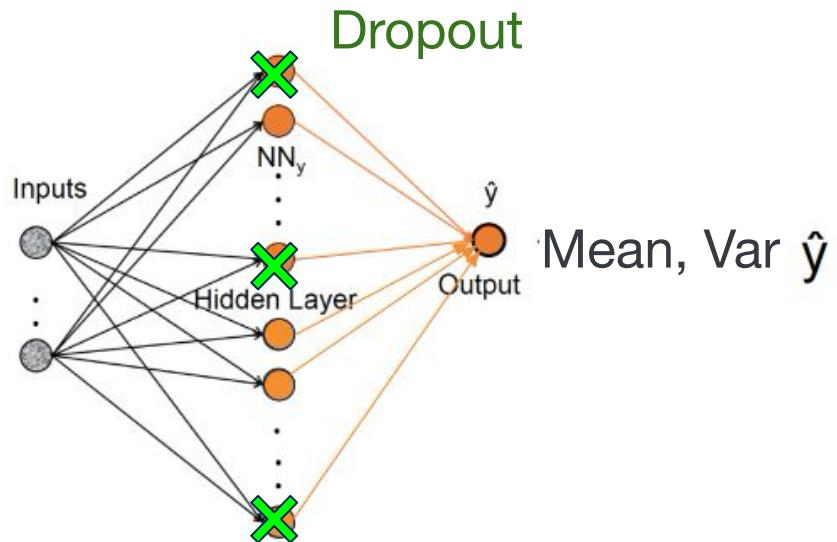
UQ for different types of models

Epistemic, aleatoric or both?

Deterministic Model - MVE



Ensembles

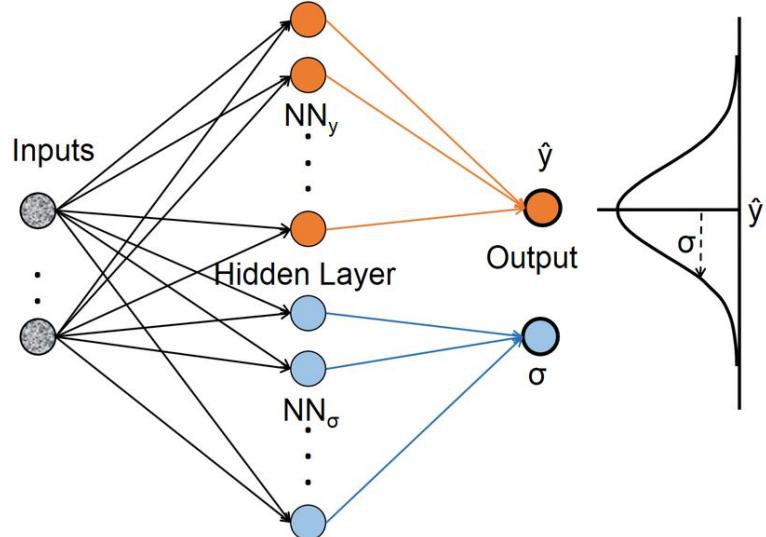




UQ for different types of models

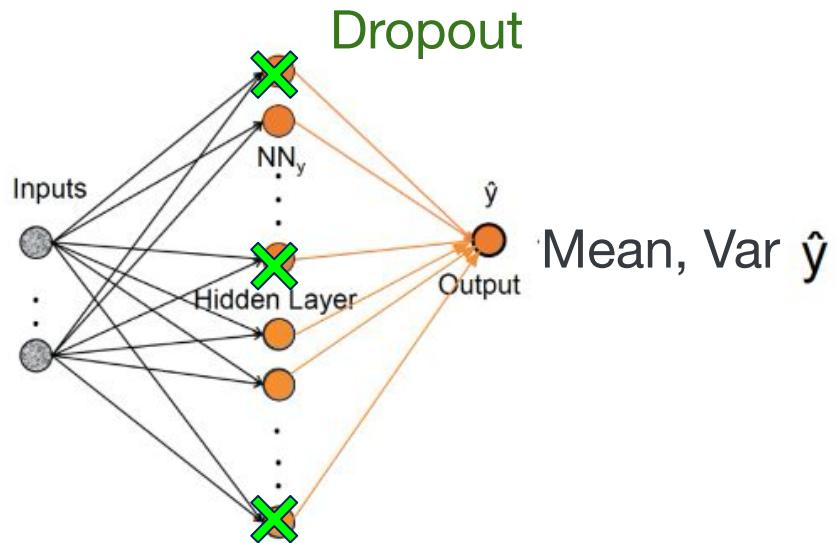
Epistemic, aleatoric or both?

Deterministic Model - MVE



Data (Aleatoric)

Ensembles



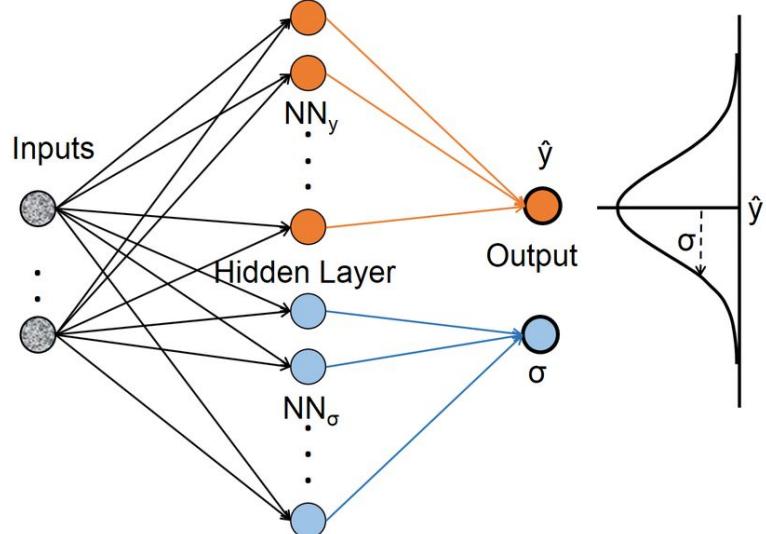
Model (Epistemic)



UQ for different types of models

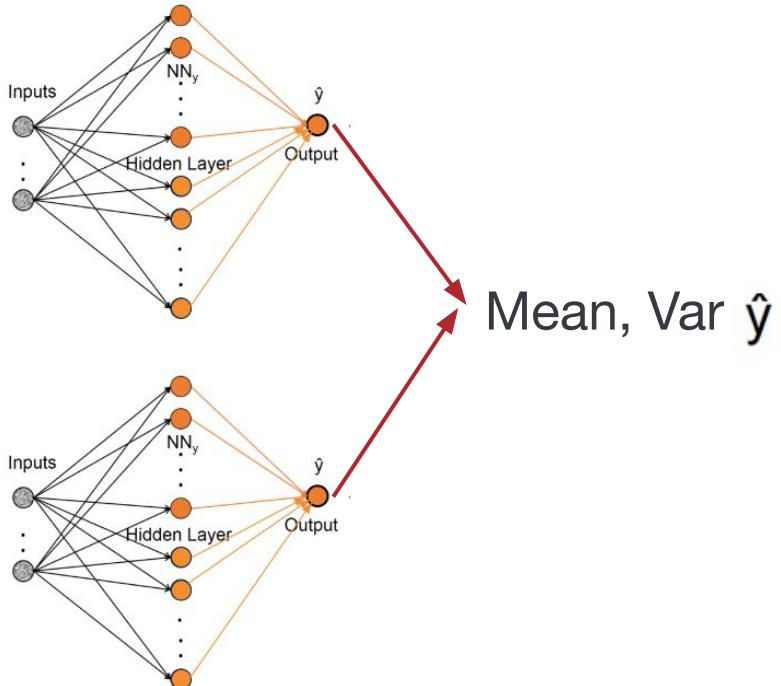
Epistemic, aleatoric or both?

Deterministic Model - MVE



Data (Aleatoric)

Ensembles



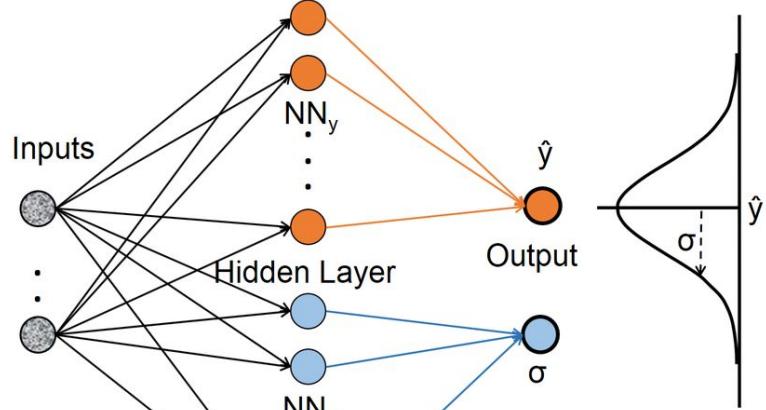
Model (Epistemic)



UQ for different types of models

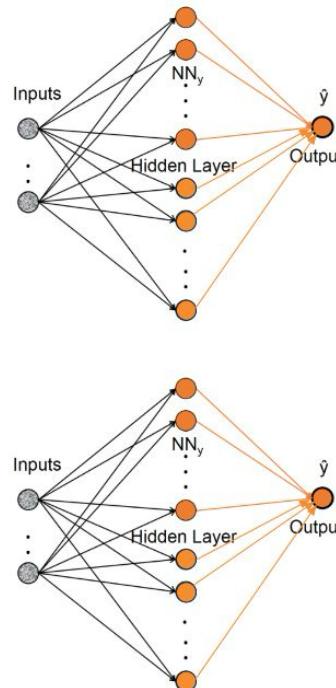
Epistemic, aleatoric or both?

Deterministic Model - MVE



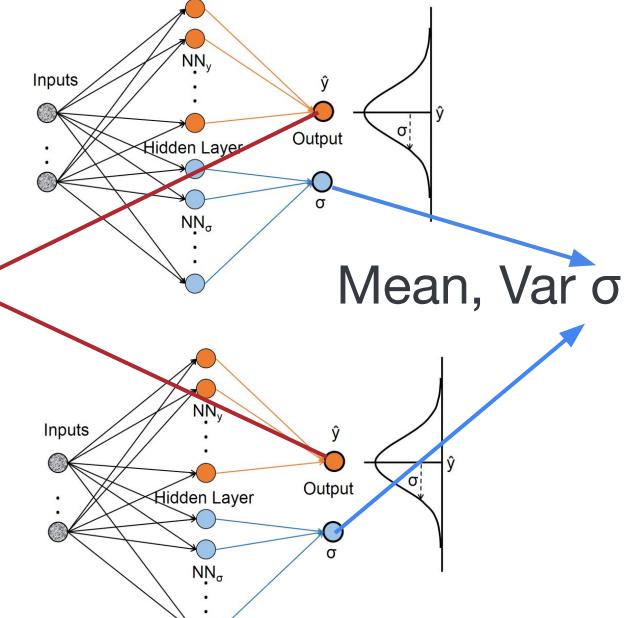
Data (Aleatoric)

Ensembles



Model (Epistemic)

Mean, Var \hat{y}



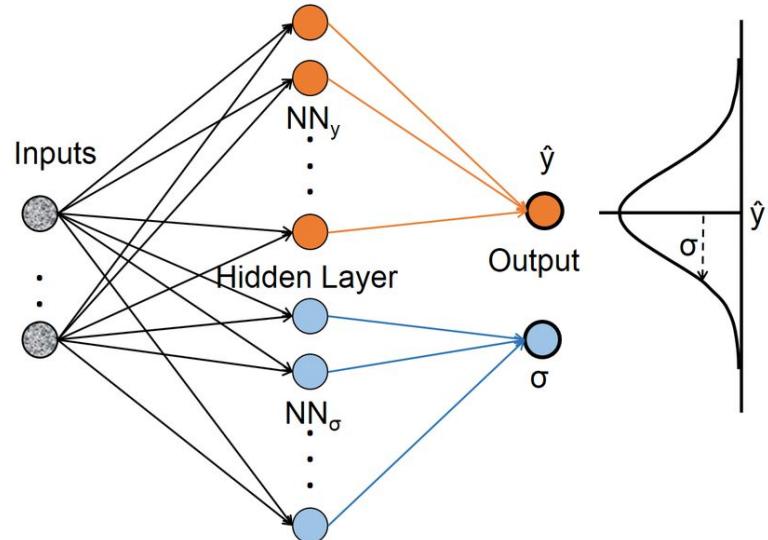
Data (Aleatoric)



UQ for different types of models

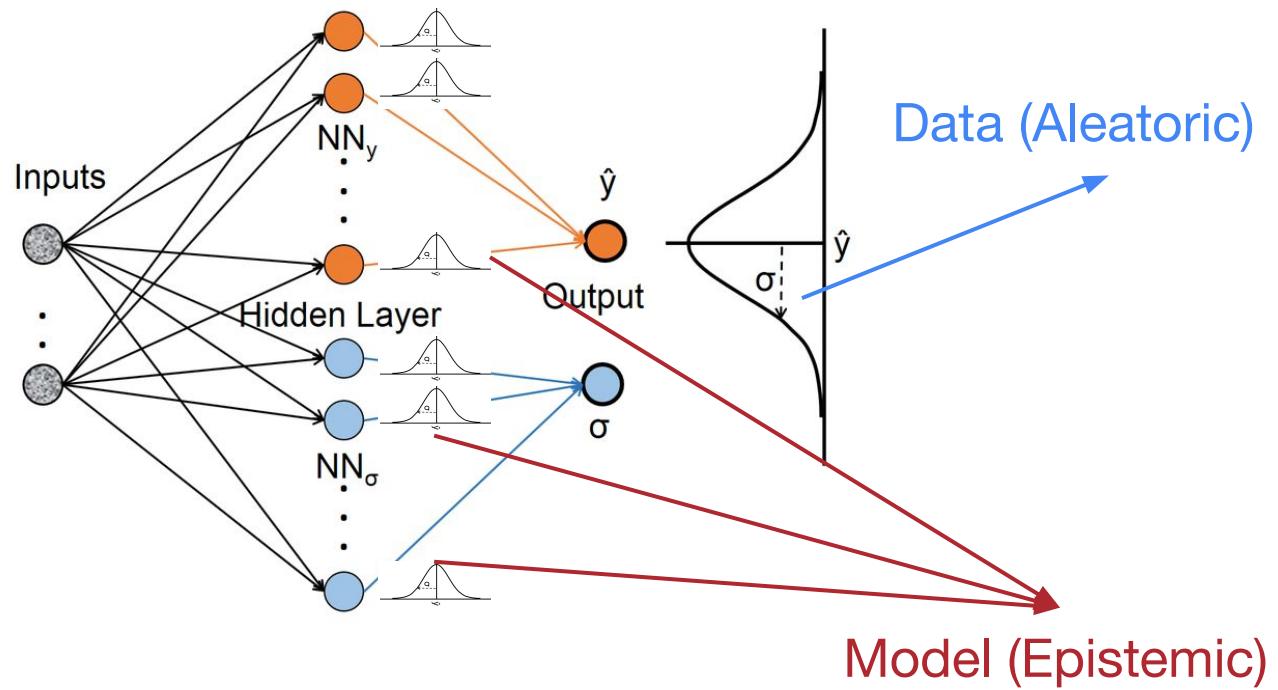
Epistemic, aleatoric or both?

Deterministic Model - MVE



Data (Aleatoric)

Bayesian Neural Networks

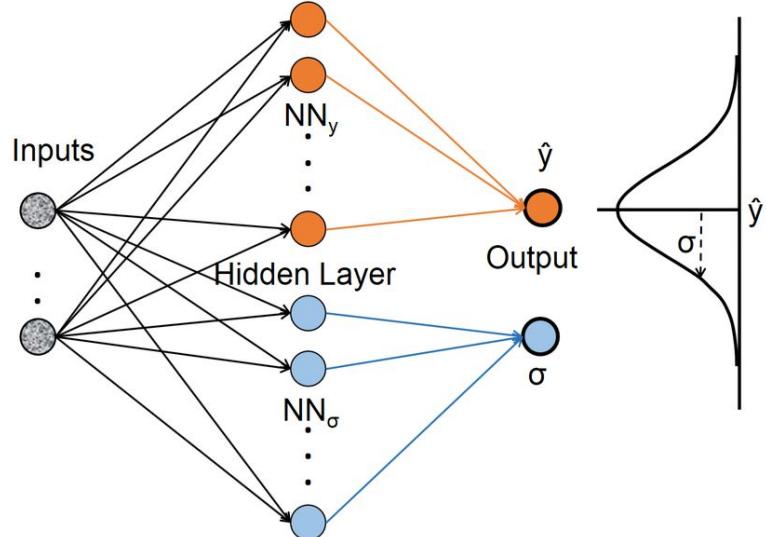




UQ for different types of models

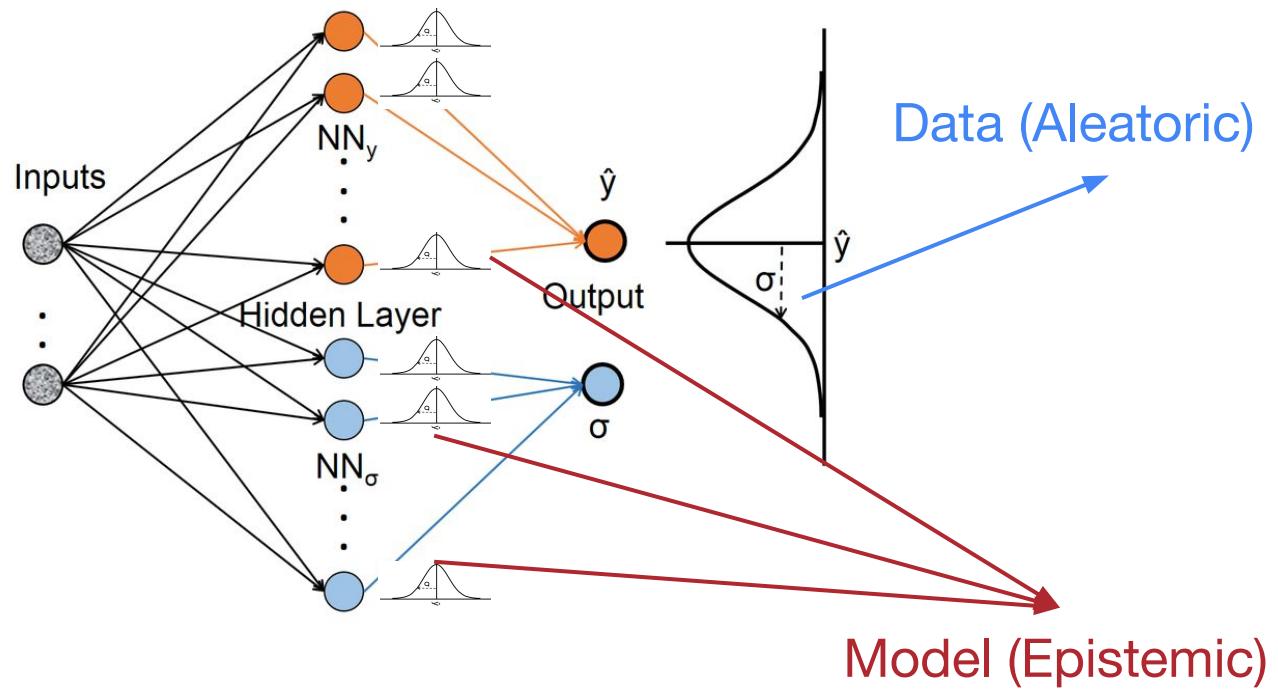
Epistemic, aleatoric or both?

Deterministic Model - MVE



Data (Aleatoric)

Bayesian Neural Networks



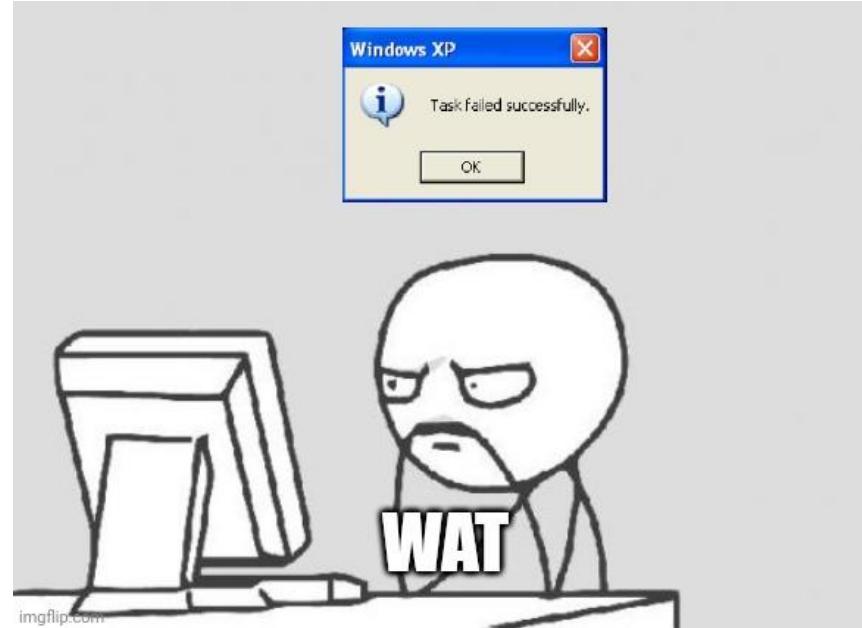
Model (Epistemic)

When domain shift is added to this picture interpretation of errors becomes even more difficult!

UQ and Domain Shift

No UQ method is robust to domain shift by default.

- Standard UQ methods often assume **i.i.d. data** → can lead to **misleading uncertainty estimates** when this is violated.
- For scientific or safety-critical applications, **overconfident but wrong predictions** under domain shift are a serious risk.
 - **Key challenges:**
 - Overconfidence - highly certain (small errors) even on out of domain data.
 - Shifted epistemic/aleatoric balance - epistemic error dominates.
 - Miscalibrated models - confidence ≠ accuracy.
 - Lack of target labels - we cannot evaluate UQ.



06

Open Questions & Recent Trends



1. Unsupervised and Few-shot DA

- Moving beyond labeled source-heavy settings.
- Combining **few-shot learning** and **DA** to adapt with just a few target labels.

2. Test-time Adaptation (TTA) & Source-free Domain Adaptation

- Models adapt **on-the-fly** during inference using only target data (entropy minimization, pseudo-labeling).
- Adapting a pretrained model using **only a trained source model**, not the data.
- No access to source data required; useful in privacy-constrained or streaming settings.

3. Multi-source and Multi-target DA

- Leveraging multiple source domains or generalizing to multiple target domains.
- Key challenge: dealing with **conflicting domain shifts** and source reliability.



4. Out-of-Distribution Generalization

- Blurring lines between DA and OOD robustness.
- Focus on **domain generalization**: training on multiple domains and generalizing to unseen ones without adaptation.

5. Domain Adaptation for Regression / Scientific Applications

- DA research has mostly focused on classification; regression problems (e.g., physics, medical data) are getting more attention.
- Requires new alignment methods and uncertainty-aware losses.

6. Vision-Language Models & Multimodal DA

- Adapting **multimodal representations** (e.g., text+image) to new domains.
- Requires alignment across **both modalities and domains**.



Unsolved Problems



1. Mismatch of Label Distributions

- Most methods assume $P(Y|X)$ remains the same, but in real-world settings it may change → **label shift**, which is under-addressed.

2. Lack of Theoretical Guarantees

- Few practical methods come with **provable generalization bounds** under domain shift.
- Need better **theory for real-world, high-dimensional settings**.

3. Interpretability and Trust

- Black-box adaptations make it hard to understand **what features are being aligned**.
- Especially important in **scientific and safety-critical domains**.



Unsolved Problems



4. Measuring Adaptation Success

- Evaluating whether a model has "adapted well" remains vague.
- Need better **metrics and diagnostics** (e.g., generalization error decomposition, target risk estimation).

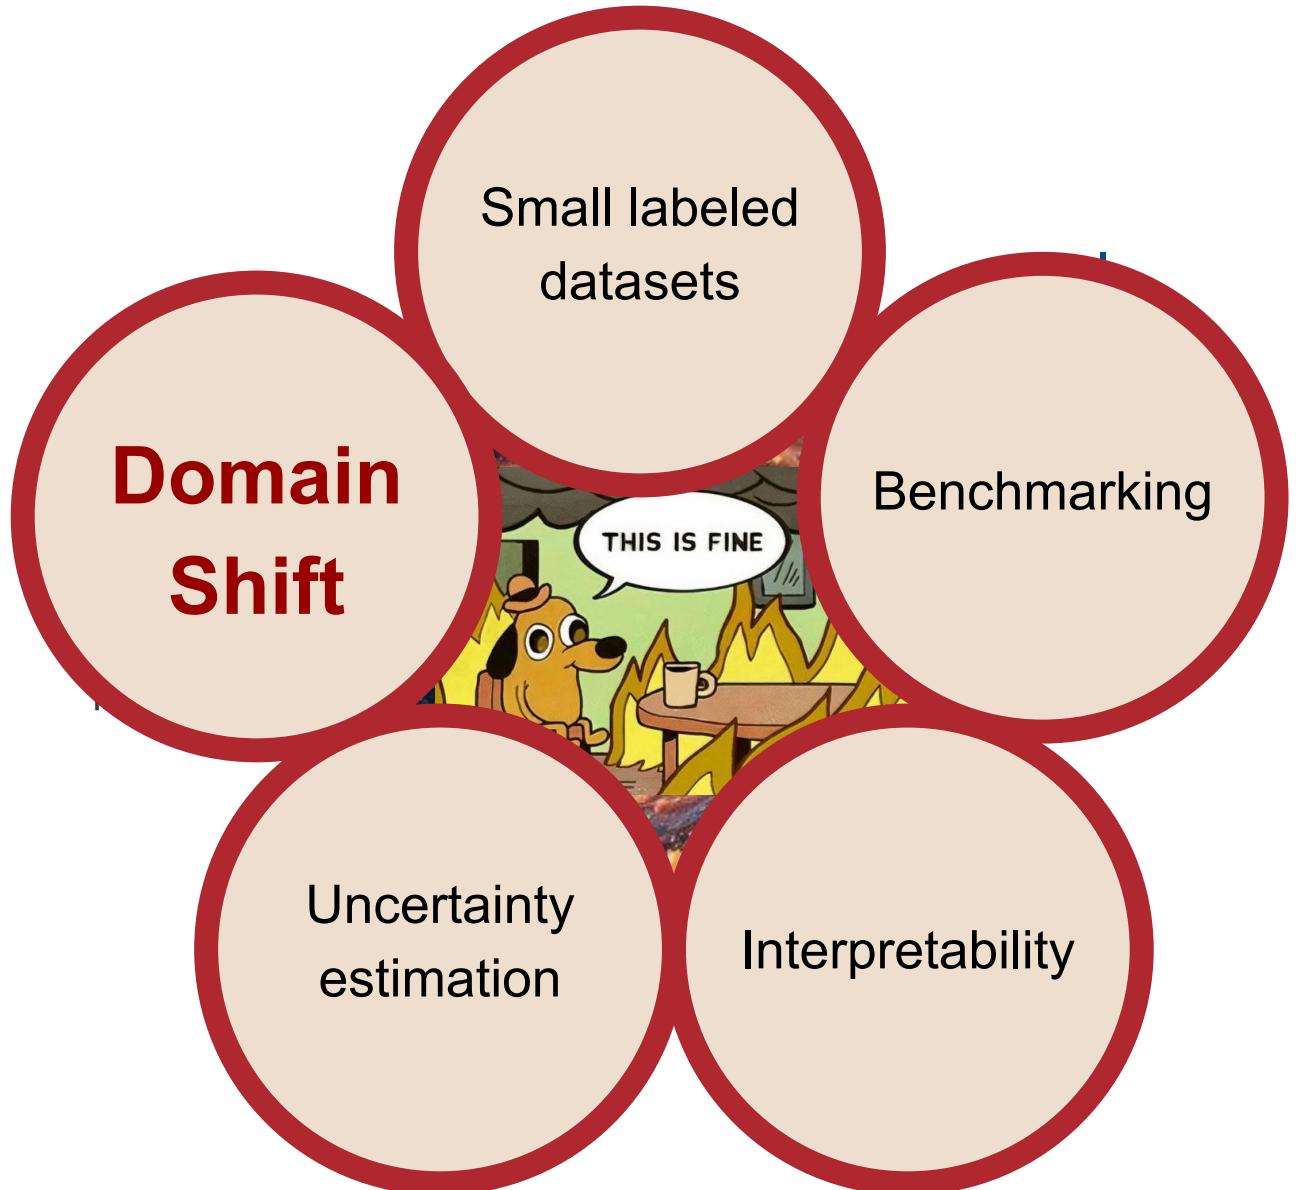
5. Continual and Lifelong Adaptation

- Adapting to a **sequence of domains** without forgetting past ones.
- Combines domain adaptation with **continual learning**.

6. Efficient and Scalable Adaptation

- Many DA algorithms are compute-heavy or memory-intensive.
- Need **efficient, lightweight** solutions for deployment in edge devices or real-time systems.

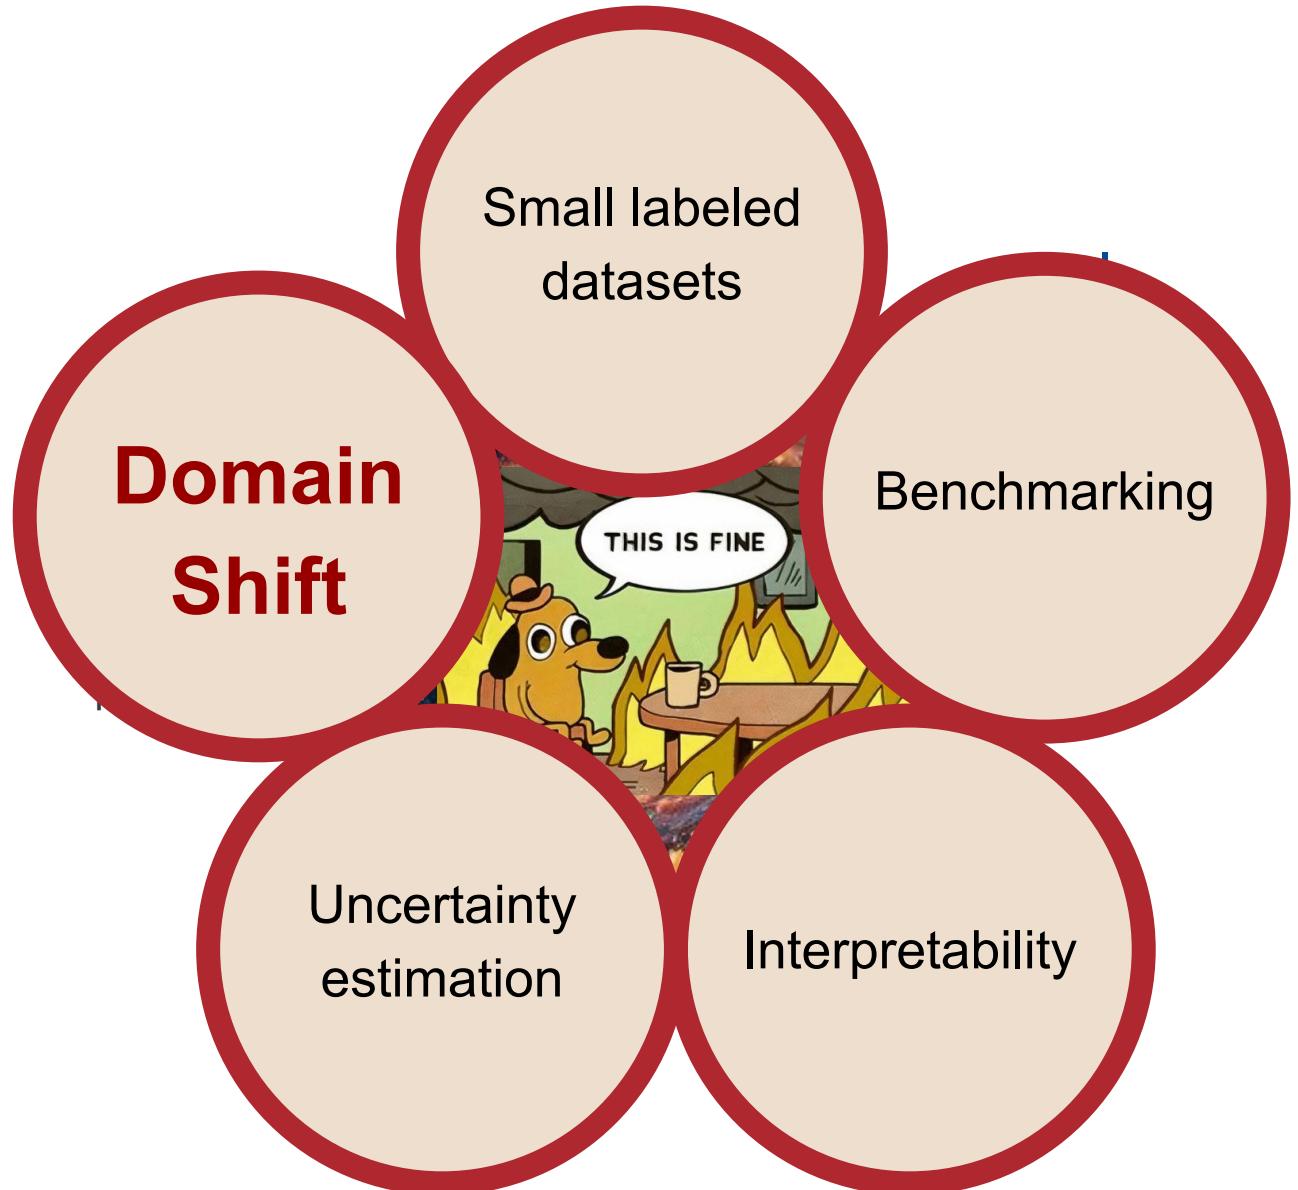
Summary



- Era of big astro surveys - **large amount of data but not all of it labeled** (i.e. suitable for easy AI applications).
- Simulations and old data are different - **domain shift problem!**
- **Domain Adaptation** can help but it cannot solve all problems.
- Hands on session:
 - Learn about different DA methods & how to tune them.



Summary



THANK YOU!

aleksand@fnal.gov

- Era of big astro surveys - **large amount of data but not all of it labeled** (i.e. suitable for easy AI applications).
- Simulations and old data are different - **domain shift problem!**
- **Domain Adaptation** can help but it cannot solve all problems.
- Hands on session:
 - Learn about different DA methods & how to tune them.



Fermi**FORWARD**



U.S. DEPARTMENT
of ENERGY